

Udemy – SQL

Tao Lin

Section 1: Intro to databases, SQL, and MySQL

1. Field-column, record - whole rows,
Relational algebra

Horizontal entity / entity instance – column

Vertical entity – row

Database entity/object -whole table

Section 2: SQL Theory

2. declarative lang: result matters but not how
3. Data Definition Language (DDL)
4. Create, alter

The screenshot shows a SQL editor interface. On the left, there's a toolbar with icons for file operations. The main area contains SQL code and a table definition. The code is:

```
CREATE TABLE object_name (column_name data_type);  
CREATE TABLE sales (purchase_number INT);
```

Below the code, a table named 'sales' is shown with one column 'purchase_number'. The table has a single row.

(1)

```
ALTER TABLE sales  
ADD COLUMN date_of_purchase DATE;
```

(2)

```
DROP TABLE customers
```

(3)

```
RENAME TABLE customer TO customer_data
```

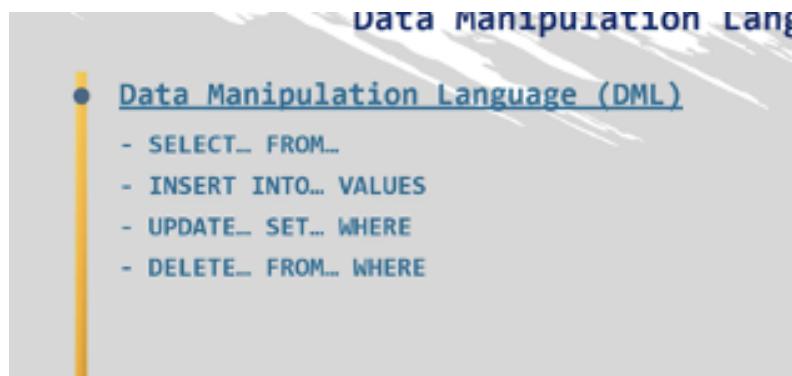
(4) delete the record

```
TRUNCATE TABLE customers
```

5. SQL key words

No confusion with variables

6. Data manipulation language (DML)



SELECT * FROM sales; #delete everything

INSERT INTO sales (purchase_number, date_of_purchase) VALUES (1, '2017-10-11');

Same as:

INSERT INTO sales VALUES (1, '2017-10-11');

UPDATE sales

SET data_of_purchase = '2017-12-12'

WHERE purchase_number = 1;

DELETE (specify precisely what to remove)

DELETE FROM sales; #DELETE is the same as TRUNCATE

DELETE FROM sales

WHERE

Purchase_number = 1;

7. Data Control Language (DCL) database administrator

(1)GRANT

GRANT type_of_permission ON database_name.table_name TO 'username'@ 'localhost';

#localhost example: IP 127.0.0.1

```
CREATE USER 'frank'@'localhost' IDENTIFIED BY 'pass';
GRANT SELECT ON sales.customers TO 'frank'@'LOCALHOST';
GRANT ALL ON Sales.* TO 'frank'@'LOCALHOST'; #grant all editing access to Frank on any
Sales table
```

(2)REVOKE: reverse process of GRANT

```
REVOKE SELECT ON sales.customers FROM 'frank'@'LOCALHOST';
```

8.Transaction Control Language

Not every change you make to a database is saved automatically

COMMIT: related to INSERT, DELETE, UPDATE

so that other users can have access to the modified database

ROLLBACK: save any changes made by making a step back

Example:

```
#individual change so access to other ppl can be provided
```

```
UPDATE customers
```

```
SET last_name = 'Johnson'
```

```
WHERE customer_id = 4
```

```
COMMIT;
```

```
ROLLBACK;
```

Reverts to the non committed changes

- DDL - Data Definition Language
creation of data
- DML - Data Manipulation Language
manipulation of data
- DCL - Data Control Language
assignment and removal of permissions to use this data
- TCL - Transaction Control Language
saving and restoring changes to a database

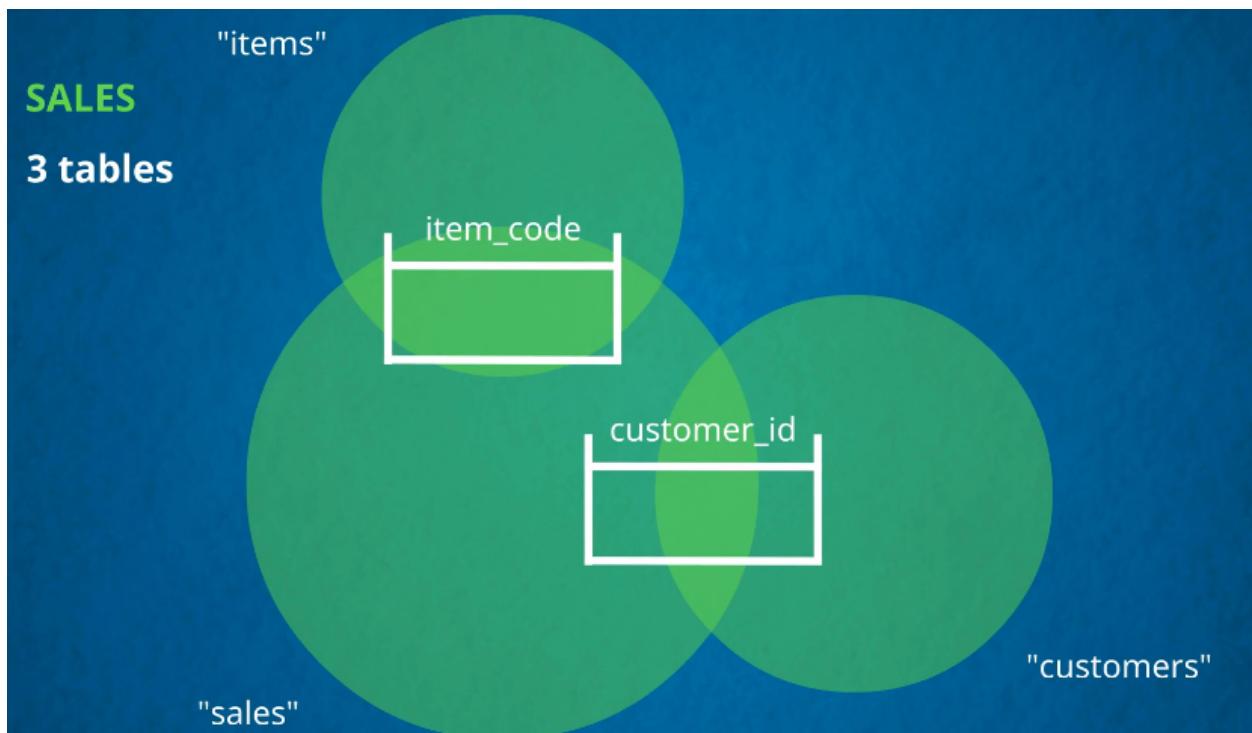
Section 3: database terminology

8. Relational database essentials

This course focuses on relational database

Understanding non-relational DB requires more complex systems and math + programming

Main goal: organize huge amounts of data that can be quickly retrieved
compact, well-structured, and efficient → relational algebra



We can use the overlapping in different tables to increase the efficiency

Tables are also called relations, which are the smallest units in the entire system that can carry integral logical meaning.

Relational database management system → RDBMS

9. Database and spreadsheets

(1) Spreadsheets

An electronic ledger, an electronic version of paper accounting worksheets

(2) Similarity with relational databases

- Can contain a large amount of tabular data
- Can use existing data to make calculations
- Are used by many users

(3) Differences

- In spreadsheets there are different values and formats; in databases, you need to pre-set the type of data contained in a certain field
- In spreadsheets, data are stored in a cell; in databases, data are stored in a record of table
- In spreadsheets, cells can contain calculations (functions and formulas); in DBs, all calculations and operations are done after data **retrieval**
- Excel only handles 1 million rows max; no limit for DBs
- Google docs cannot find who changed the deleted information incorrectly; DBs can create tables about permissions

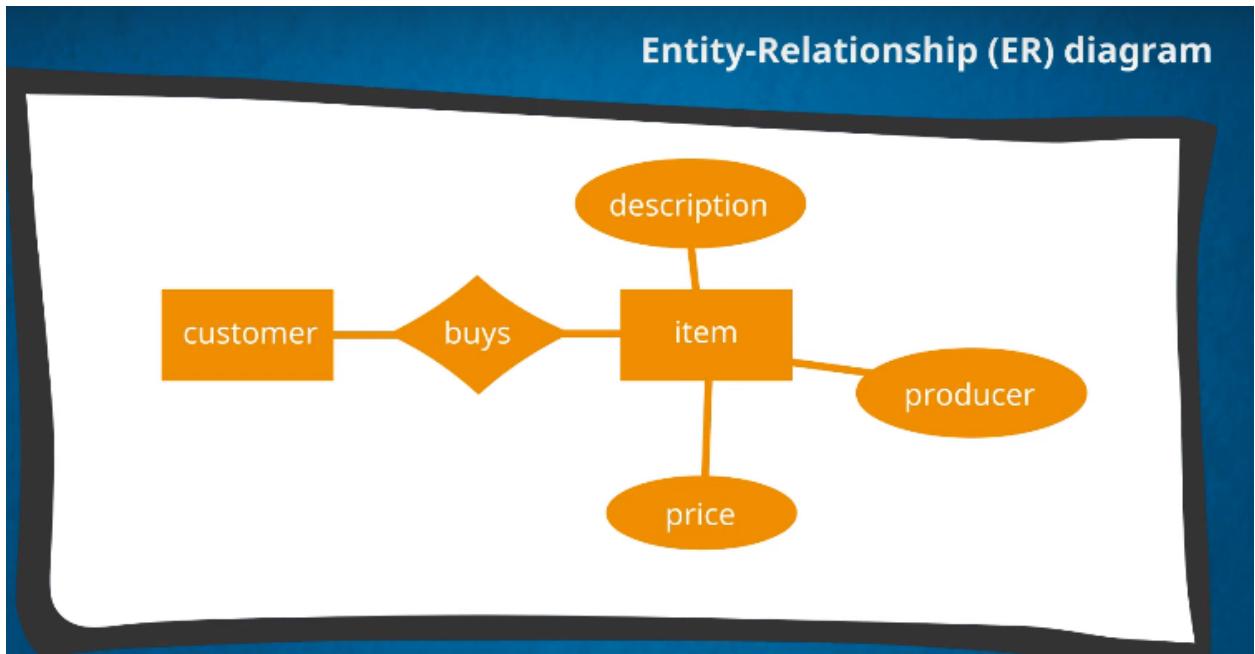
DBs have data consistency and integrity → so no duplicates → save space



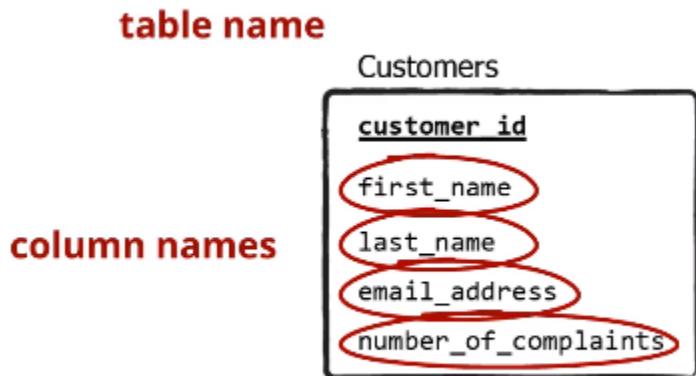
10. Database terminology

Database design: plot the entire database system on a canvas using a visualization tool
 Two main ways

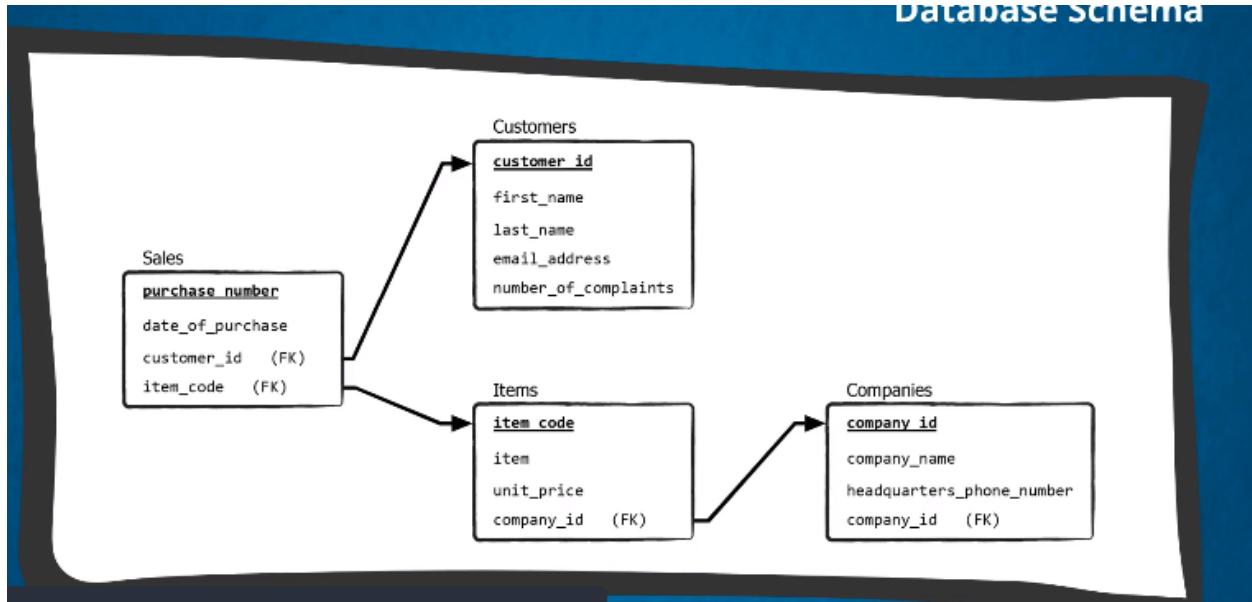
- (1) Entity relationship (ER) diagram



(2) Relational schema - how the database is organized



Database Schema



Arrows show relationships

Database creation: use SQL to set up database physically

Data manipulation: use dataset to extract business insights

DB management = design + creation + manipulation

Database administration

11. Relational schemas: a primary key

Example: four tabulars - sales, customers, items, and companies

Customers					
customer_id	first_name	last_name	email_address	number_of_complaints	
1	John	McKinley	john.mackinley@365careers.com	0	
2	Elizabeth	McFarlane	e.mcfarlane@365careers.com	2	
3	Kevin	Lawrence	kevin.lawrence@365careers.com	1	
4	Catherine	Winnfield	c.winnfield@365careers.com	0	

Sales					
purchase_number	date_of_purchase	customer_id	item_code		
1	9/3/2016	1	A_1		
2	12/2/2016	2	C_1		
3	4/15/2017	3	D_1		
4	5/24/2017	1	B_2		
5	5/25/2017	4	B_2		
6	6/6/2017	2	B_1		
7	6/10/2017	4	A_2		
8	6/13/2017	3	C_1		
9	7/20/2017	1	A_1		
10	8/11/2017	2	B_1		

Items			
item_code	item	unit_price_usd	company_id
A_1	Lamp	20	1
A_2	Desk	250	1
B_1	Lamp	30	2
B_2	Desk	350	2
C_1	Chair	150	3
D_1	Loudspeakers	400	4

Companies		
company	Headquarters_Phone_number	company_id
Company A	+1 (202) 555-0196	1
Company B	+1 (202) 555-0152	2
Company C	+1 (229) 853-9913	3
Company D	+1 (618) 369-7392	4

Sales: date of purchase and customer ID can have the same values; purchase number is unique

Primary key → Column(s) whose value exists and is unique for every record in a table is called a primary key

Features:

- each table can have one and only one primary key
- Usually single item primary key (or unique identifier). Sometimes the key can be composed of a set of primary keys. For example, purchase number + date of purchase, like (2, 9/3/16)
- Primary key does not allow nulls/blanks

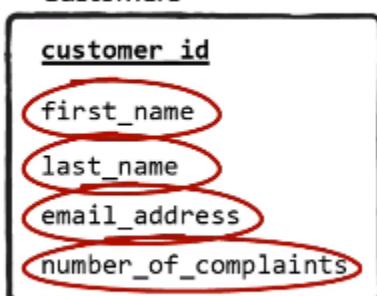
Sales			
<u>purchase_number</u>	<u>date_of_purchase</u>	<u>customer_id</u>	<u>item_code</u>
1	9/3/2016	1	A_1
2	12/2/2016	2	C_1
3	4/15/2017	3	D_1
	5/24/2017	1	B_2
5	5/25/2017	4	B_2
6	6/6/2017	2	B_1
	6/10/2017	4	A_2
8	6/10/2017	3	C_1
9	7/20/2017	1	A_1
10	8/11/2017	2	B_1

Primary key is always underlined

table name

Customers

column names



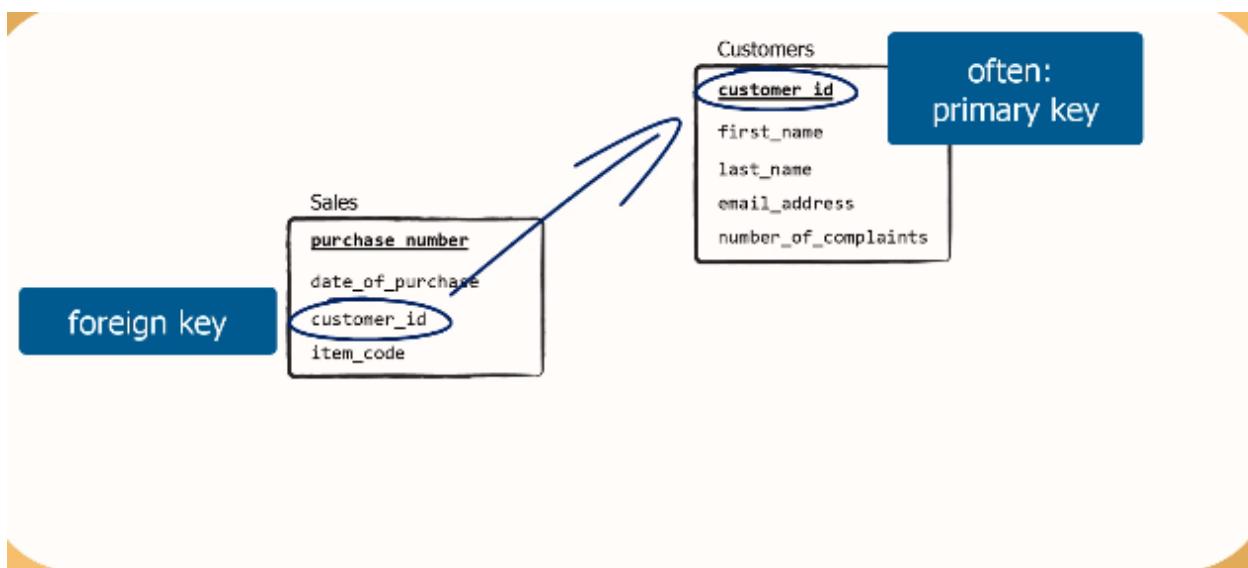
- Not all the tables have a primary key

12. Foreign key

Customers					
<u>customer_id</u>	<u>first_name</u>	<u>last_name</u>	<u>email_address</u>	<u>number_of_complaints</u>	
1	John	McKinley	john.mackinley@365careers.com	0	
2	Elizabeth	McFarlane	e.mcfarlane@365careers.com	2	
3	Kevin	Lawrence	kevin.lawrence@365careers.com	1	
4	Catherine	Winnfield	c.winnfield@365careers.com	0	

Sales			
purchase_number	date_of_purchase	customer_id	item_code
1	9/3/2016	1	A_1
2	12/2/2016	2	C_1
3	4/15/2017	3	D_1
	5/24/2017	1	B_2
5	5/25/2017	4	B_2
6	6/6/2017	2	B_1
	6/10/2017	4	A_2
8	6/10/2017	3	C_1
9	7/20/2017	1	A_1
10	8/11/2017	2	B_1

In these two tables, customer_id is the primary key in Customers, but foreign key in Sales.



- Foreign keys often points to primary keys
- Always look for the foreign keys, as they show us where the relations are
- FKs can be repeated and missing (e.g., the same customers can buy items several times)
- **FK identifies the relationship between tables, not the tables themselves**

**no repeating and missing values
(unique values only)**

repeating and missing values

13. Relational schemas: unique key and null values

In the example, we need to create the Companies table

Unique key: used whenever you would like to specify that you don't want to see duplicate data in a given field

Comparing to primary keys:

	primary key	unique key
NULL VALUES	no	yes
NUMBER OF KEYS	1	0, 1, 2...
APPLICATION TO MULTIPLE COLUMNS	yes	yes

14. Relationships

Relationships tell you how much of the data from a foreign key field can be seen in the primary key column of the table the data is related to and vice versa

(1) One-to-many type of relationship

One value from the customer_id under customers table can be found many times in the customer_id column in the Sales table

Sales			
purchase_number	date_of_purchase	customer_id	item_code
1	9/3/2016	1	A_1
2	12/2/2016	2	C_1
3	4/15/2017	3	D_1
4	5/24/2017	1	B_2
5	5/25/2017	4	B_2
6	6/6/2017	2	B_1
7	6/10/2017	4	A_2
8	6/13/2017	3	C_1
9	7/20/2017	1	A_1
10	8/11/2017	2	B_1

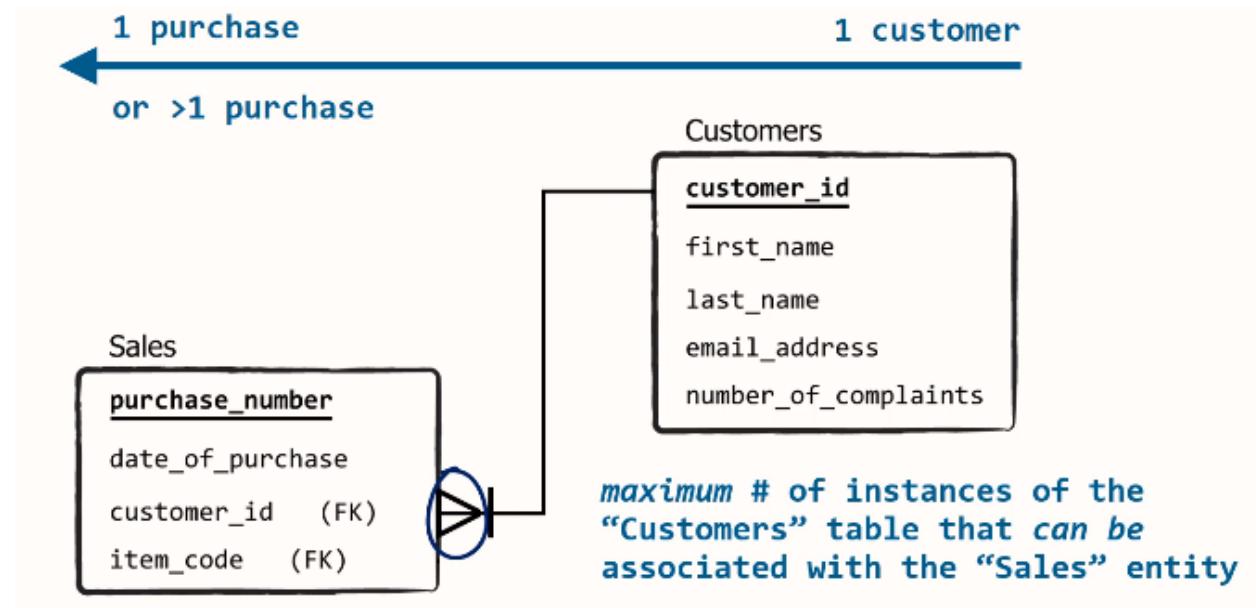
Customers				
customer_id	first_name	last_name	email_address	number_of_complaints
1	John	McKinley	john.mackinley@365careers.com	0
2	Elizabeth	McFarlane	e.mcfarlane@365careers.com	2
3	Kevin	Lawrence	kevin.lawrence@365careers.com	1
4	Catherine	Winnfield	c.winnfield@365careers.com	0

unique values

repeated values

Direction 1:

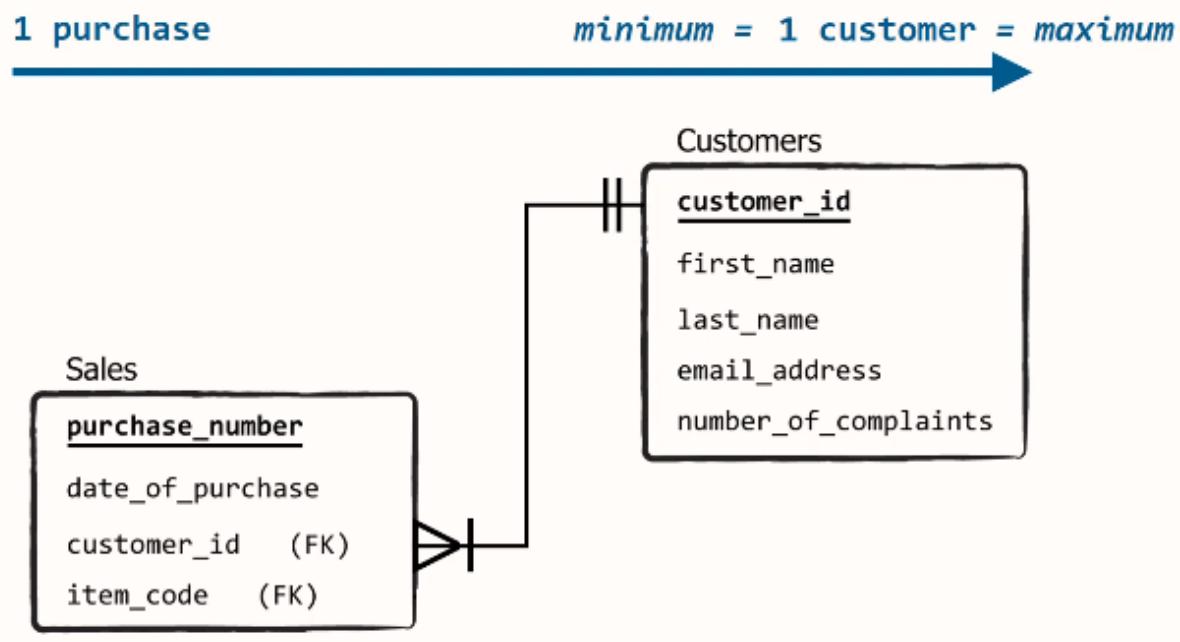
One to many



>1 : there is at least customer id in Sales

Direction 2:

Many to one



|| : max = min = 1

Other kinds of cardinality constraints: | > M N O

- (2) one -to-one
- (3) Many-to-many

Summary of relationships:

Relationships

- [Relational schemas](#)
 - represent the concept database administrators must implement
 - depict how a database is organized
 - = blueprints, or a plan for a database
 - *will help you immensely while writing your queries!*

14. Downloading softwares

Setting up connections (not required for this course)

Section 4: first steps in SQL

15. Creating a database in MySQL

Command template: CREATE DATABASE [IF NOT EXISTS] database_name;

[] means optional conditions

Database_name: the name is not case sensitive

Semicolon: must be included

CREATE DATABASE IF NOT EXISTS Sales;

#same as

CREATE SCHEMA IF NOT EXISTS Sales;

Use Sales;

16. Different data types

We must always specify the type of data that will be inserted in each column of the table.

String - the text format in SQL

	Surname of a person:	length	size
string	'James'	5 symbols	5 bytes

Although digits can be used, there is no computation

String is also called alphanumeric data type

17. String data subtypes

Char, varchar, enumerate

<u>string data type</u>		<u>Storage</u>	<u>Example</u>	<u>length</u> (symbols)	<u>size</u> (bytes)
<i>character</i>	CHAR	fixed	CHAR(5) ‘James’ 5 5 ‘Bob’ 3 5		
<i>variable character</i>	VARCHAR	variable	VARCHAR(5) ‘James’ 5 5 ‘Bob’ 3 3		

Although (5) means space limit, VARCHAR is more flexible.

<u>string data type</u>		<u>Maximum size</u> (bytes)	
<i>character</i>	CHAR	255	50% faster
<i>variable character</i>	VARCHAR	65,535	a lot more responsive to the data value inserted

CHAR is faster; VARCHAR is more responsive

Example: company code uses CHAR; password prefers VARCHAR

ENUM: limited choices

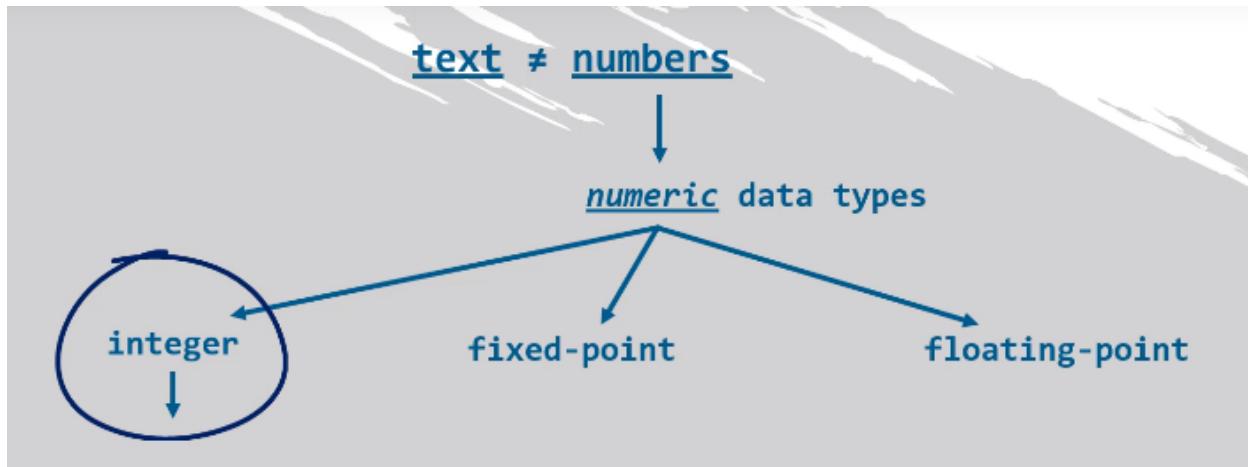
<u>string data type</u>	Example	
<i>character</i>	CHAR	CHAR(5)
<i>variable character</i>	VARCHAR	VARCHAR(5)
<i>ENUM ("enumerate")</i>	ENUM	ENUM('M','F')

ERROR

MySQL will show an error if you attempt to insert any value different from "M" or "F".

18. Integers

There are different kinds of numbers



Integers: numbers without decimal points.

Subtypes: size and min/max; these ints vary in different sizes (also processing speed)

<u>numeric data type</u>	<u>size (bytes)</u>	<u>minimum value (signed/unsigned)</u>	<u>maximum value (signed/unsigned)</u>
TINYINT	1	-128 0	127 255
SMALLINT	2	-32,768 0	32,767 65,535
MEDIUMINT	3	-8,388,608 0	8,388,607 16,777,215
INT	4	-2,147,483,648 0	2,147,483,647 4,294,967,295
BIGINT	8	-9,223,372,036,854,775,808 0	9,223,372,036,854,775,807 18,446,744,073,709,551,615

Signed - if the encompassed range includes both positive and negative; integer data are signed by default

Unsigned - if the values are only allowed to be positive; needs to be specified in your query

19. Fixed- and floating-point data types

[Precision]: number of digits

[Scale]: number of digits to the right of the decimal point

<u>number:</u>	<u>precision</u>	<u>scale</u>
10.523	5	3

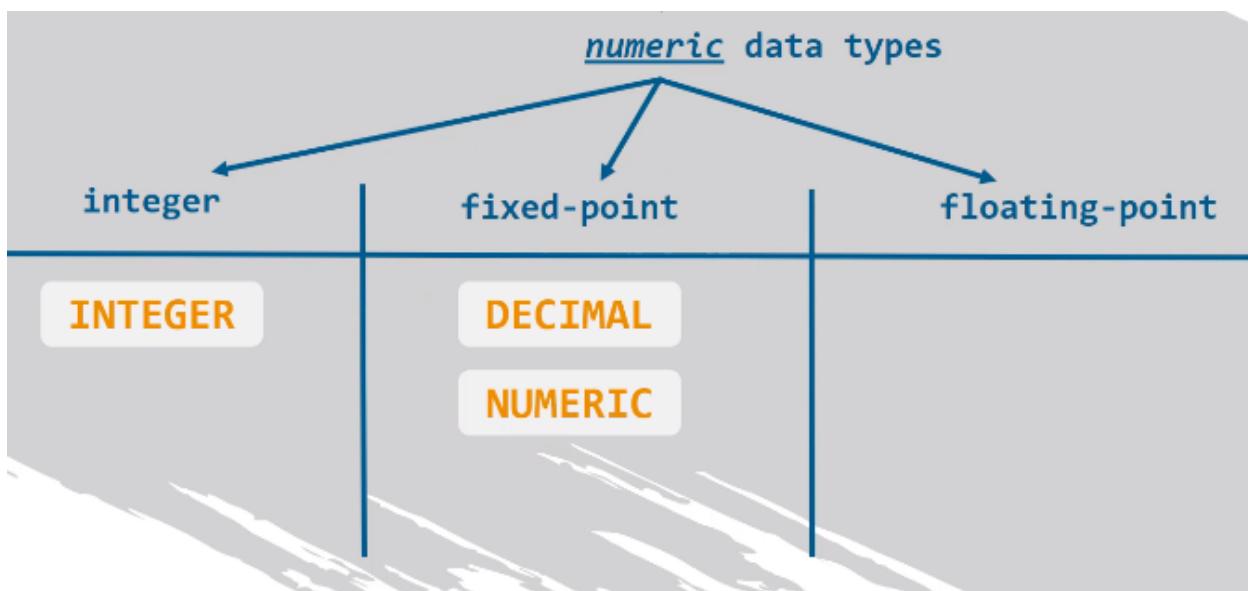
Another way to describe this number: DECIMAL(5,3)

Fixed-point: exact values

When only one digit is specified within the parentheses, it will be treated as the precision of the data type. **DECIMAL(7) = DECIMAL(7,0) for 123456**

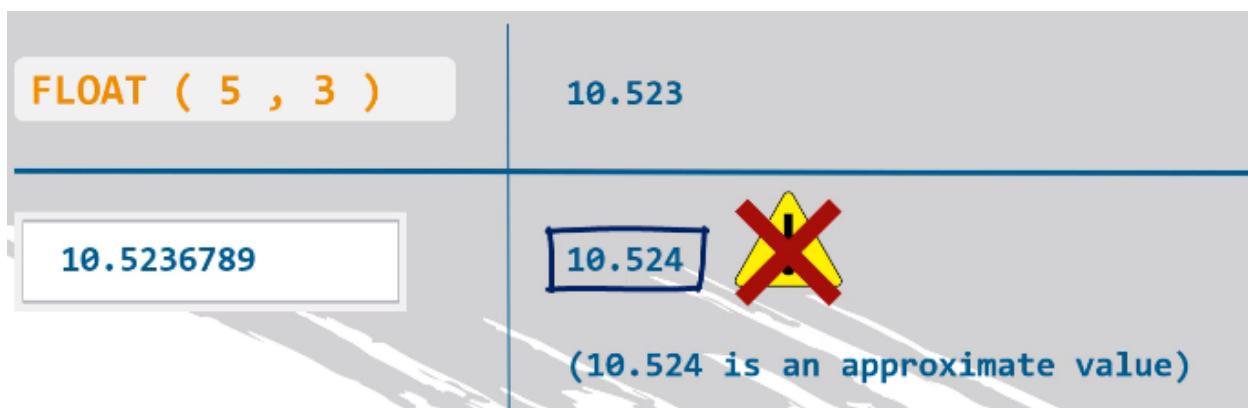


The pic above shows rounding will be given a warning for the fixed-point type

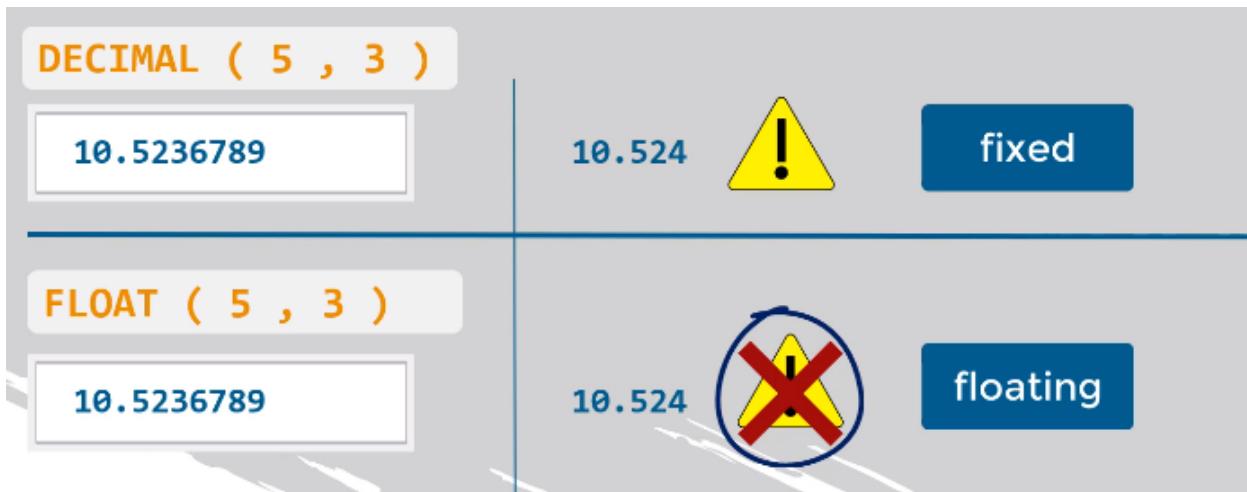


The NUMERIC function is the same as DECIMAL.

Floating-point data: used for approximate values only; aims to balance between range and precision (>= "floating")



As for the case of rounding (the actual number doesn't fit the structure defined), fixed- will return a warning, but float- doesn't.



FLOAT vs. DOUBLE (in range)

<u>Floating-point data type</u>	<u>size (bytes)</u>	<u>precision</u>	<u>maximum number of digits</u>
FLOAT	4	<i>single</i>	23
DOUBLE	8	<i>double</i>	53

20. Other data types

DATE: YYYY-MM-DD (range: 1000-1-1 ~ 9999-12-31)

DATETIME: YYYY-MM-DD:MM:SS[.fraction]

Represents the date shown on the calendar and the time shown on the clock

YYYY-MM-DD HH:MM:SS [.fraction]

0 - 23:59:59.999999

e.g. 25th of July 2018 9:30 a.m.: **'2018-07-25 9:30:00'**

TIMESTAMP:

- Used for a well-defined, exact point in time
- Representing a moment in time as a number allows you to easily obtain the difference between two TIMESTAMPS values
- Good at handling time zone differences

1st of January 1970 UTC - 19th of January 2038, 03:14:07 UTC

The 1970-1-1 is also the baseline time.

e.g. 25th of July 2018:

1,535,155,200

more than 48 years!

BLOB (Binary Large OBject)

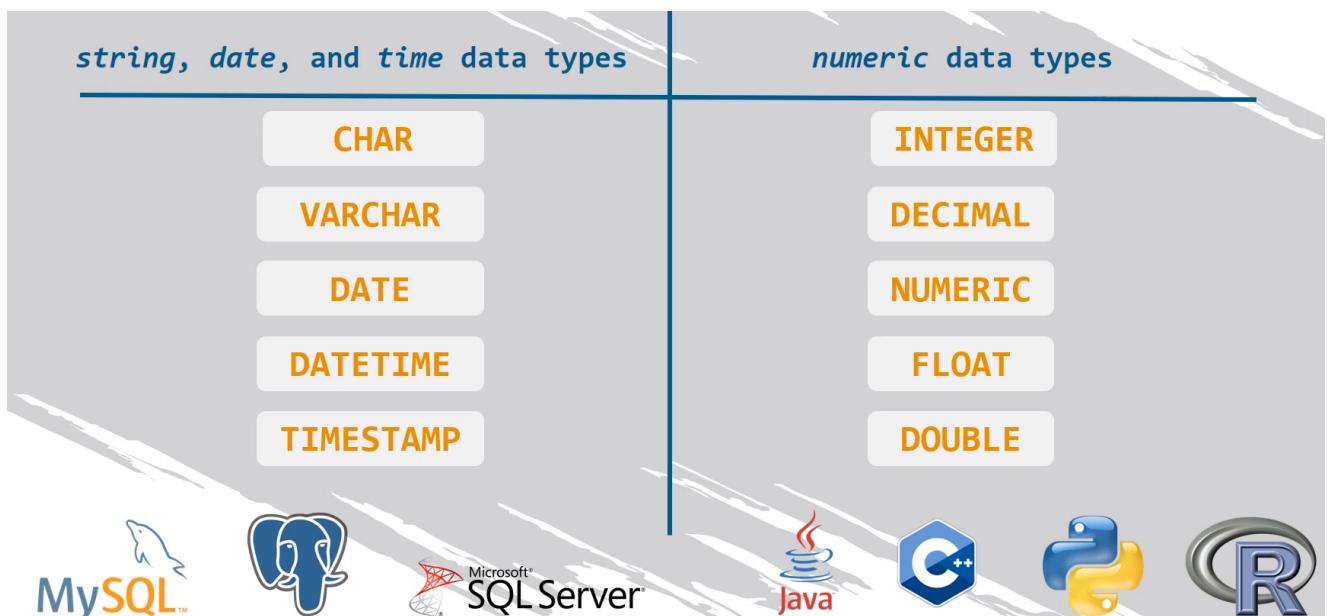
- Refers to a file of binary data - data with 1s and 0s
- Involves saving files in a record



Example:

Customers						
customer_id	first_name	last_name	email_address	number_of_complaints	photo	
1	John	McKinley	john.mackinley@365careers.com	0		
2	Elizabeth	McFarlane	e.mcfarlane@365careers.com	2		
3	Kevin	Lawrence	kevin.lawrence@365careers.com	1		
4	Catherine	Winnfield	c.winnfield@365careers.com	0		*.jpg

Summary:



21. Creating a table

```
CREATE DATABASE [IF NOT EXIST] sales;
-> CREATE TABLE table_name (); #() is column name
```

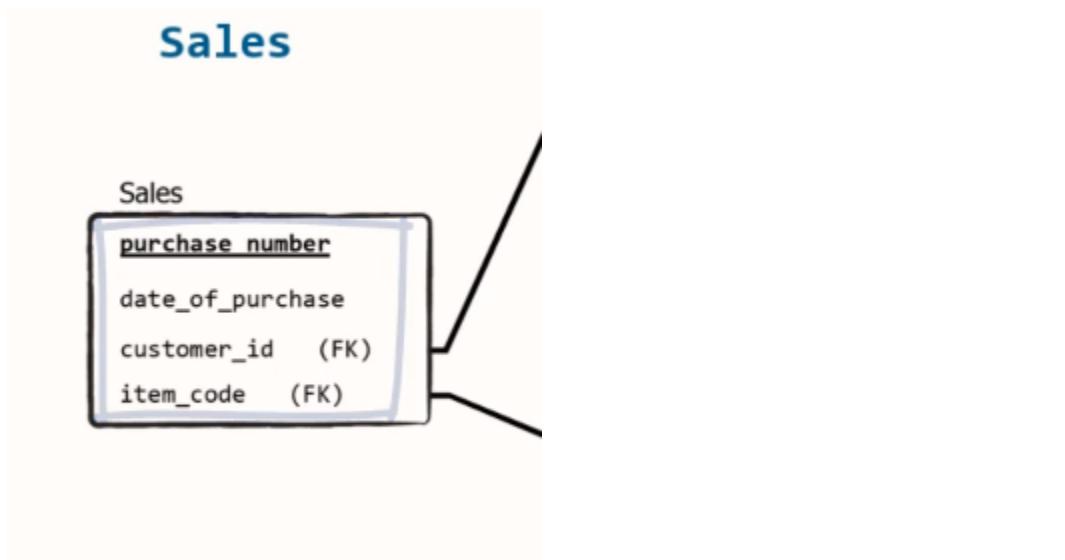
Compulsory requirement:P add at least one column

Full grammar:



```
SQL    CREATE TABLE table_name
      (
        column_1 data_type constraints,
        column_2 data_type constraints,
        ...
        column_n data_type constraints
      );
```

Example:



AUTO_INCREMENT

Frees you from having to insert all purchase numbers manually through the INSERT command at a later stage

- Assign 1 to the first record of the table and automatically increments by 1 for every subsequent row (similar to EXCEL)

Exercise_creating a table

```
CREATE DATABASE IF NOT EXISTS Sales;
USE Sales;
CREATE TABLE sales
(
  purchase_number INT NOT NULL PRIMARY KEY AUTO_INCREMENT,
```

```
data_of_purchse DATE NOT NULL,  
customer_id INT,  
item_code VARCHAR(10) NOT NULL  
);
```

22. Using databases and tables

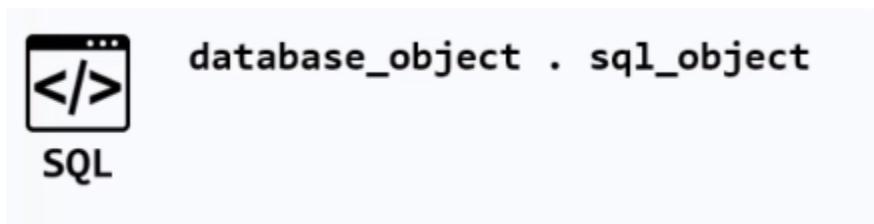
Whenever you would like to refer to an SQL subject in your queries, you must specify the database to which it is applied.

SQL objects: SQL tables, views, stored procedures, and functions.

- (1) Set a default database



- (2) Call a table from a certain database



“.” is called a “dot operator”

Signals the existence of a connection between the two object types

```
SELECT * FROM sales;
```

```
SELECT * FROM sales.sales;
```

23. Additional Notes on Using tables

Query is a command you write in SQL with the idea of either **retrieving information** from the database from the database on which you are working, or, alternatively, to **insert, update, or delete data** from it.

```
2 • CREATE TABLE sales
3   (
4     purchase_number INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
5     date_of_purchase DATE NOT NULL,
6     customer_id INT,
7     item_code VARCHAR(10) NOT NULL
8   ); |
```

"We ran the CREATE TABLE query"

A query contains the semi-colon, which is a statement terminator.

The DROP statement → **delete a table**

SQL

```
DROP TABLE table_name;
DROP TABLE sales;
```

Section 5: MySQL constraints

24. Primary key constraints

Constraints are specific rules or limits that we define our tables. The role of constraints is to outline the existing relationships between different tables in our database.

Example1: NOT NULL is a constraint

Example2: constraints on three keys

PRIMARY KEY Constraint

Sales

Sales
<u>purchase_number</u>
date_of_purchase
customer_id (FK)
item_code (FK)

Customers

<u>customer_id</u>
first_name
last_name
email_address
number_of_complaints

Items

primary key?
foreign key?
unique key?

Companies

you must define them in SQL through their respective **constraints**

Example3: constraints on purchase_number -AUTO_INCREMENT & PRIMARY KEY

```
1 • CREATE TABLE sales
2   (
3     purchase_number INT AUTO_INCREMENT PRIMARY KEY,
4     date_of_purchase DATE,
5     customer_id INT,
6     item_code VARCHAR(10)
7   );
8
```

Another way to write it:

```
CREATE TABLE sales
(
purchase_number INT AUTO_INCREMENT,
data_of_purchse DATE,
customer_id INT,
item_code VARCHAR(10) ,
PRIMARY KEY (purchase_number)
);
```

you can run the program by ctrl+shift+enter

Exercise: key constraints

Drop the “customers” table and re-create it using the following code:

```
CREATE TABLE customers

(
    customer_id INT,
    first_name varchar(255),
    last_name varchar(255),
    email_address varchar(255),
    number_of_complaints int,
    primary key (customer_id)
);
```

Then, create the “items” table

(columns - data types:

item_code – VARCHAR of 255,

item – VARCHAR of 255,

unit_price – NUMERIC of 10 and 2,

company_id – VARCHAR of 255),

and the “companies” table

(company_id – VARCHAR of 255,

company_name – VARCHAR of 255,

headquarters_phone_number – integer of 12).

A:

```
DROP TABLE customers;
```

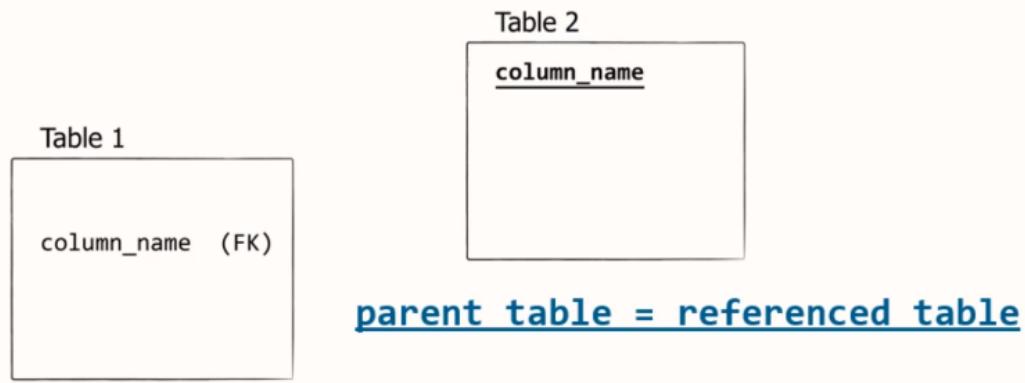
```
CREATE TABLE customers
(
    customer_id INT,
    first_name varchar(255),
    last_name varchar(255),
    email_address varchar(255),
    number_of_complaints int,
    primary key (customer_id)
);
```

```
CREATE TABLE items
(
    item_code VARCHAR(255),
    item VARCHAR(255),
    unit_price NUMERIC(10,2),
    company_id VARCHAR(255)
);
```

```
CREATE TABLE companies
(
    company_id VARCHAR(255),
    company_name VARCHAR(255),
    headquarters_phone_number INT(12)
);
```

25. Foreign key constraints

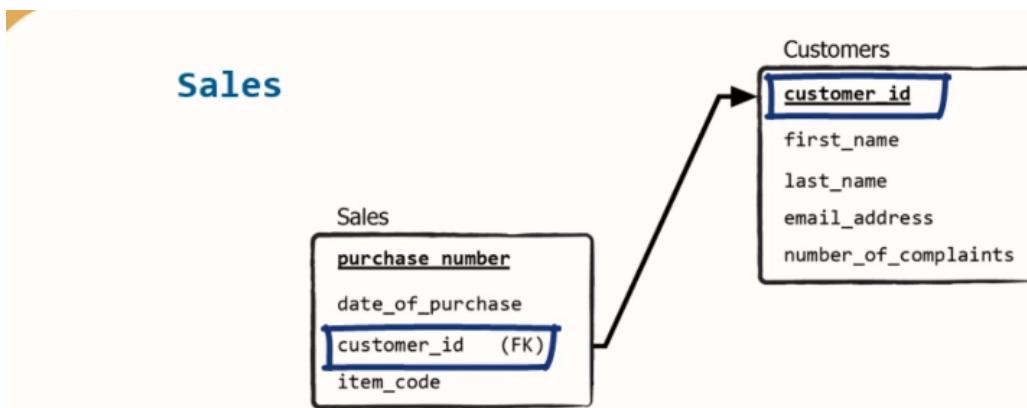
Foreign keys point to a column of another table, thus links two tables together



child table = referencing table

Usually the key in the referenced/parent table is the primary key.

Example1:



Remember, this is not an obligatory requirement - these two keys may have two completely different names. What's important is that the data types and the information match! It's just common practice to use, if not the same, then similar names for both keys.

A foreign key in SQL is defined through a foreign key constraint.

```
CREATE TABLE sales
(
    purchase_number INT AUTO_INCREMENT,
    date_of_purchase DATE,
    customer_id INT,
    item_code VARCHAR(10),
    PRIMARY KEY (purchase_number),
    FOREIGN KEY (customer_id) REFERENCES customers(customer_id) ON DELETE CASCADE
);
```

ON DELETE CASCADE

It is a specific value from the **parent** table's primary key has been deleted , all the records from the **child** table referring to this value will be removed as well (kinda like sync)

TIPS: (1) when you see “table sales already exists” you can “drop table sales” and rerun the other codes

(2) you can alter a foreign key in an existing table by using “alter table sales”

```
2 • drop table sales;
3
4 • CREATE TABLE sales
5   (
6       purchase_number INT AUTO_INCREMENT,
7       date_of_purchase DATE,
8       customer_id INT,
9       item_code VARCHAR(10),
10      PRIMARY KEY (purchase_number)
11      #foreign key (customer_id) references customers (customer_id) ON DELETE cascade
12  );
13
14 • alter table sales
15     add foreign key (customer_id) references customers (customer_id) on delete cascade;
```

Drop the foreign key using DDL code in a table:

```
alter table sales
drop foreign key sales_ibfk_1;
```

```
drop table sales;
```

```
CREATE TABLE sales
```

```

(
    purchase_number INT AUTO_INCREMENT,
    date_of_purchase DATE,
    customer_id INT,
    item_code VARCHAR(10),
PRIMARY KEY (purchase_number)
#foreign key (customer_id) references customers (customer_id) ON DELETE cascade
);

```

```

alter table sales
add foreign key (customer_id) references customers (customer_id) on delete cascade;

```

```

alter table sales
drop foreign key sales_ibfk_1;

```

(3) Using right click on the target table

The screenshot shows the MySQL Workbench interface for managing database objects. On the left, the Navigator pane shows the schema structure with 'sales' as the current schema. The main window displays the 'sales' table configuration. The 'Foreign Keys' tab is active, showing a single foreign key entry:

Foreign Key Name	Referenced Table
sales_ibfk_2	`sales`.`customers`

Under the 'Column' column, the 'customer_id' checkbox is checked, indicating it is the referenced column. The 'Referenced Column' column shows 'customer_id'. In the 'Foreign Key Options' section, 'On Delete' is set to 'CASCADE'. There is also a checkbox for 'Skip in SQL generation' which is unchecked.

Exercise: drop all sales tables

DROP TABLE sales;

```
DROP TABLE customers;
```

```
DROP TABLE items;
```

```
DROP TABLE companies;
```

26. Unique key constraints

Unique key is used whenever you'd like to specify that you don't want to see duplicate data in a given field. It ensures that all values in a column (or more than one) are different

With the constraint, if you attempt to insert existing, duplicate value in the unique column, SQL will display an error

Example: emails in customers

(A) Add unique key in create table

```
CREATE TABLE customers
```

```
(
```

```
customer_id INT,
```

```
first_name varchar(255),
```

```
last_name varchar(255),
```

```
email_address varchar(255),
```

```
number_of_complaints int,
```

```
primary key (customer_id),
```

Unique key (email_address)

```
);
```

(B) Use alter table

```
CREATE TABLE customers
```

```
(  
    customer_id INT,  
    first_name varchar(255),  
    last_name varchar(255),  
    email_address varchar(255),  
    number_of_complaints int,  
    primary key (customer_id)  
);
```

```
alter table customers  
add unique key (email_address);
```

[Indexes]

Unique keys in MySQL have the same function as indexes; but the reverse is not true

Index is an organizational unit that helps retrieve data more easily.

Drop a unique key:

```
ALTER TABLE table_name  
DROP INDEX unique_key_field;
```

```
alter table customers  
drop index (email_address);
```

Exercise: unique constraints

Drop the “customers” table, and then recreate it using the following code.

```
CREATE TABLE customers (
```

```
    customer_id INT AUTO_INCREMENT,
```

```
first_name VARCHAR(255),  
last_name VARCHAR(255),  
email_address VARCHAR(255),  
number_of_complaints INT,  
PRIMARY KEY (customer_id)  
);
```

Then run the following code that will add a “gender” column in the “customers” table, and will then insert a new record in it. Don’t worry if you don’t understand the meaning of the code perfectly – we will discuss these structures later on in the course in more detail. We will just use them now to insert a row in our “customers” table.

```
ALTER TABLE customers  
ADD COLUMN gender ENUM('M', 'F') AFTER last_name;
```

```
INSERT INTO customers (first_name, last_name, gender, email_address,  
number_of_complaints)  
VALUES ('John', 'Mackinley', 'M', 'john.mckinley@365careers.com', 0)  
;
```

27. Default constraints

Default constraint helps assign a particular default value to every row of a column

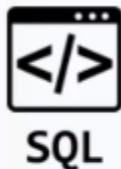
If a value is different, it can be stored in a field where the default constraint is applied, ONLY IF specifically indicated.

Example: #complaint

Customers					
customer_id	first_name	last_name	email_address	number_of_complaints	
1	John	McKinley	john.mackinley@365careers.com	0	
2	Elizabeth	McFarlane	e.mcfarlane@365careers.com	2	
3	Kevin	Lawrence	kevin.lawrence@365careers.com	1	
4	Catherine	Winnfield	c.winnfield@365careers.com	0	

DEFAULT Value: 0

Grammar: after the variable name (parenthesis is optional)



```
CREATE TABLE customers
(
    customer_id INT,
    first_name VARCHAR(255),
    last_name VARCHAR(255),
    email_address VARCHAR(255),
    number_of_complaints INT DEFAULT 0,
    PRIMARY KEY (customer_id)
);
```

(1) In a preexisting table, use “alter table NAME change column”

```
alter table customers
change column number_of_complaints number_of_complaints INT default 0;
```

(2) We can use new values to test the default value. Full code:

```
DROP TABLE customers;
```

```
CREATE TABLE customers (
```

```
customer_id INT AUTO_INCREMENT,
```

```

first_name VARCHAR(255),
last_name VARCHAR(255),
gender ENUM('M', 'F'),
email_address VARCHAR(255),
number_of_complaints INT,
PRIMARY KEY (customer_id)
);

alter table customers
change column number_of_complaints number_of_complaints INT default 0;

insert into customers (first_name, last_name, gender)
values ("Peter", "Figaro", "M")
;

select * from customers; #this shows the table itself

```

(3) Drop the default value

```

alter table customers
alter column number_of_complaints drop default;

```

(4) Data definition language

According to the lectures on constraints, there are three types:
CREATE, ALTER, and DROP.

Exercise: practice with default key

Recreate the “companies” table

```

(company_id – VARCHAR of 255,
company_name – VARCHAR of 255,
headquarters_phone_number – VARCHAR of 255),

```

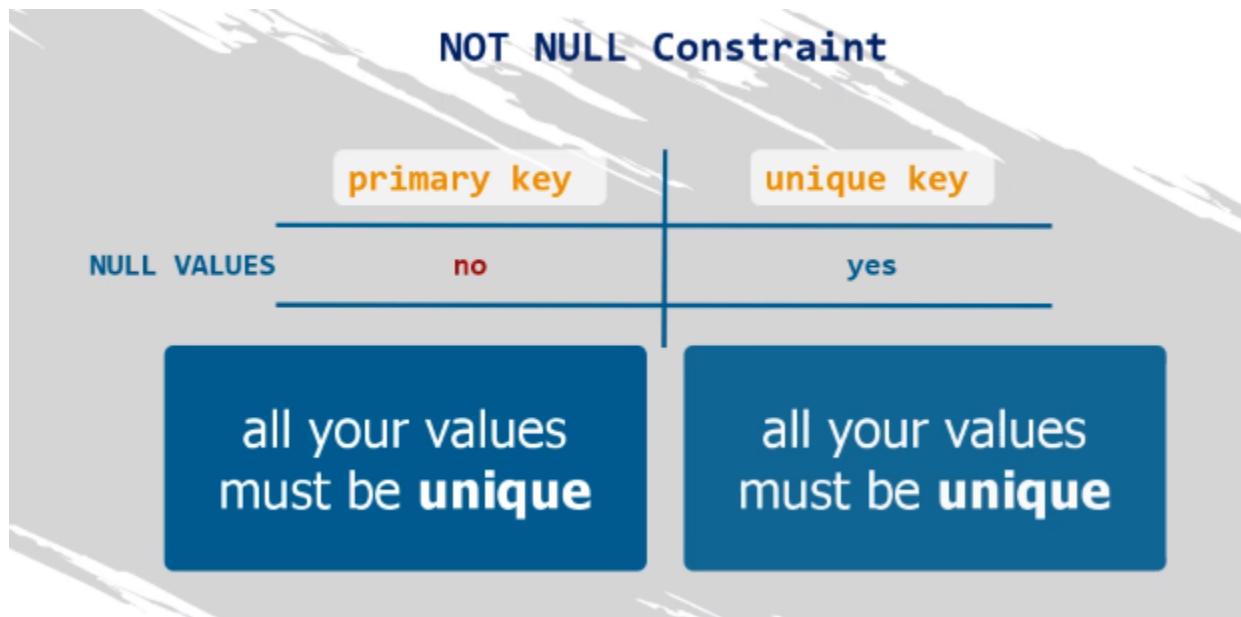
This time setting the “headquarters phone number” to be the unique key, and default value of the company's name to be “X”.

After you execute the code properly, drop the “companies” table.

```
create table companies(  
    company_id varchar(255),  
    company_name varchar(255) default "X",  
    headquaters_phone_number varchar(255),  
primary key (company_id),  
unique key (headquaters_phone_number)  
);  
drop table companies;
```

→ primary key must be defined here!

28. NOT NULL constraint



NOT NULL restriction means that when you insert values in the table, **you cannot leave the respective field empty**. Otherwise there will be an error.

Example: add NOT NULL to company_name

(1) Add NOT NULL

Companies		
company_id	headquarters_phone_number	company_name
1	+1 (202) 555-0196	Company A
2	+1 (202) 555-0152	Company B
3	+1 (229) 853-9913	Company C
4	+1 (618) 369-7392	Company D

NOT NULL

```
create table companies(
    company_id INT auto_increment,
    company_name varchar(255) NOT NULL,
    headquarters_phone_number varchar(255),
    primary key (company_id)
);
```

(2) Change NOT NULL to NULL (after creating the table)

```
alter table companies
modify company_name varchar(255) NULL;
```

(3) After (2) change NULL back to NOT NULL

```
alter table companies
change column company_name company_name varchar(255) NOT NULL;
```

(4) Test the NOT NULL restriction

```
insert into companies(headquarters_phone_number)
values ('+1 (202) 555 0196');
```

Error Code: 1364. Field 'company_name' doesn't have a default value

Correct:

```
insert into companies(headquarters_phone_number,company_name)
values ('+1 (202) 555 0196','Company A');
```

Exercise: Using ALTER TABLE, first add the NULL constraint to the *headquarters_phone_number* field in the “companies” table, and then drop that same constraint.

```
ALTER TABLE companies  
MODIFY headquarters_phone_number VARCHAR(255) NULL;
```

```
ALTER TABLE companies  
CHANGE COLUMN headquarters_phone_number headquarters_phone_number  
VARCHAR(255) NOT NULL;
```

29. More on NOT NULL

Don't confuse a NULL value with the value of 0 or with a NONE response
NULL value is more like a **missing/empty** value

0	NONE	NULL
assigned by the user		assigned by the computer

Section 7: SQL Best Practices

30. Coding techniques/styles

- You will always work as a team
- Good code is not the one computers understand; it is the one humans can understand
→ needs to be easy to read and understand

clean code

code that is *focused* and *understandable*, which means it must be readable, logical, and changeable

- Always choose shorter and meaningful names, conveying specific information

```
CREATE TABLE sales
(
    purchase_number INT,
    date_of_purchase DATE,
    customer_id VARCHAR(255),
    item_code VARCHAR(255),
    PRIMARY KEY (purcase_number)
);
```

- Usually you make the names lowercase and capitalize the SQL keywords
- Readability: horizontal and vertical codes; colors

suggestions

1

use ad-hoc software that re-organizes code and colours different words consistently

- time is a factor
- unification of coding style is a top-priority

it is unprofessional to merge code written in the same Language but in a different style

2

use the relevant analogical tool provided in Workbench

3

intervene manually and adjust your code as you like

#2: use **ctrl+B/the brush icon** in mySQL to beautify the codes

Exercise:

There are two lines of raw code:

```
use sales;  
create table if not exists test (numbers int(10), words varchar(10));
```

Answer: press **ctrl+B**

```
1 • use sales;           I  
2  
3 • CREATE TABLE IF NOT EXISTS test (  
4     numbers INT(10),  
5     words VARCHAR(10)  
6 );
```

Use TAB indentations to assign the parameters

- Make comments: large & one-line

comments

lines of text that Workbench will not run as code; they convey a message to someone who reads our code

/ ... */ (for Large comments)*
or -- (for one-line comments)

- Execution: lightening (run the selected or all) - **ctrl+shift+enter**; lightning and pointer (run the code under the pointer) - **ctrl+enter**

31. Loading a database

Load the employee.sql in MySQL

Section 8: Select statements

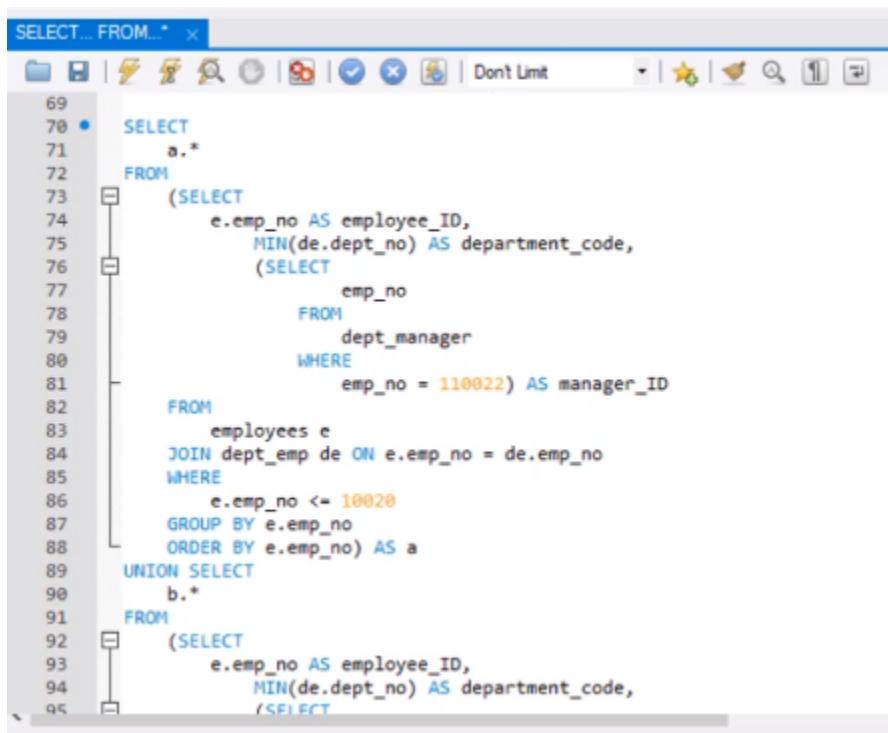
32. Select statements

The SELECT statement

- Allows you to extract a fraction of the entire dataset
- Used to retrieve data from database objects, like tables.

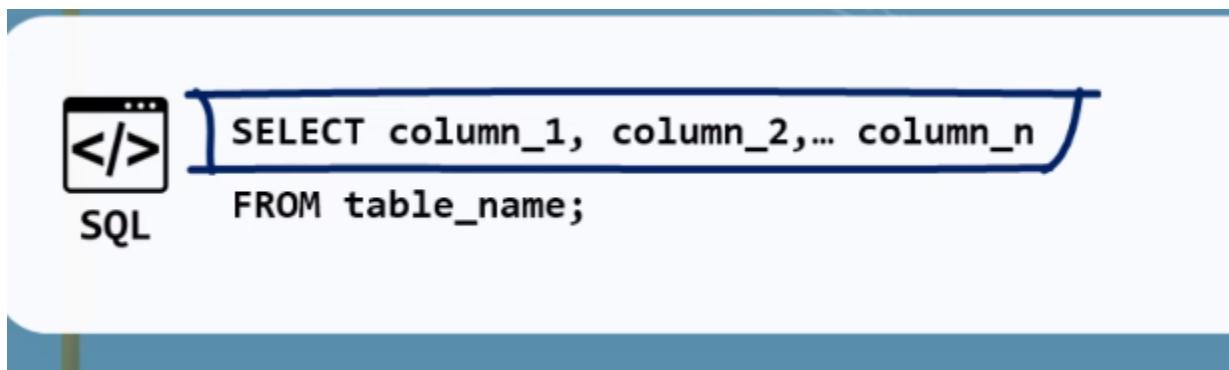
- Used to query data from a database

A complicated example of SELECT:



```
SELECT ... FROM... *  
69  
70 •  SELECT  
71     a.*  
72     FROM  
73         (SELECT  
74             e.emp_no AS employee_ID,  
75             MIN(de.dept_no) AS department_code,  
76             (SELECT  
77                 emp_no  
78                 FROM  
79                     dept_manager  
80                     WHERE  
81                         emp_no = 110022) AS manager_ID  
82             FROM  
83                 employees e  
84             JOIN dept_emp de ON e.emp_no = de.emp_no  
85             WHERE  
86                 e.emp_no <= 10020  
87             GROUP BY e.emp_no  
88             ORDER BY e.emp_no) AS a  
89         UNION SELECT  
90             b.*  
91             FROM  
92                 (SELECT  
93                     e.emp_no AS employee_ID,  
94                     MIN(de.dept_no) AS department_code,  
95                     /SELECT
```

SELECT ... FROM...



```
7 •  SELECT first_name, last_name  
8     FROM employees;
```

SELECT everything:

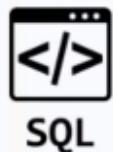
```
7 •  SELECT *
8   FROM employees;
```

Exercise: SELECT ... FROM ...

```
7 •  SELECT dept_no
8   FROM departments;
```

33. WHERE statements

The where clause will allow us to set a **condition** upon which we will specify what part of the data we want to retrieve from the database.



```
SELECT column_1, column_2,... column_n
FROM table_name
WHERE condition;
```

Example: extract anyone named Denis from employees

```

8 •  SELECT *
9   FROM employees
10  WHERE first_name = "Denis";

```

	emp_no	birth_date	first_name	last_name	gender	hire_date
▶	11688	1958-09-04	Denis	Couillard	F	1994-10-29
	15083	1958-11-24	Denis	Nicolson	M	1994-03-02
	15824	1957-07-28	Denis	Schwabacher	F	1988-02-14
	17116	1961-02-03	Denis	Mullainathan	F	1998-12-23
	17224	1955-04-06	Denis	Back	M	1985-07-29

For the condition checking, one equal sign is used.

Exercise: anyone named Elvis from employees

```

SELECT *
FROM employees
WHERE first_name = "Elvis";

```

34. AND

In SQL there are many other linking keywords and symbols like “=”, called operators, that you can use with the WHERE clause.

- AND
- OR
- IN - NOT IN
- LIKE -NOT LIKE
- BETWEEN ... AND ...
- EXISTS - NOT EXISTS
- IS NULL - NOT NULL
- Comparison operators
- Etc.

→ AND list multiple conditions



```
SELECT column_1, column_2,... column_n  
FROM table_name  
WHERE condition_1 AND condition_2;
```

Example: find Denis + male

```
8 •  SELECT *  
9    FROM employees  
10   WHERE first_name = "Denis" AND gender = "M";
```

Exercise: all females named Kellie

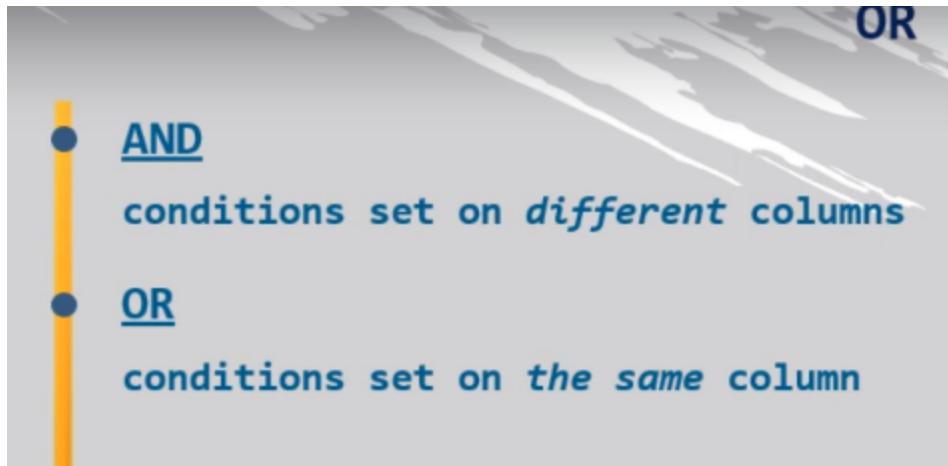
```
8 •  SELECT *  
9    FROM employees  
10   WHERE first_name = "Kellie" AND gender = "F";
```

35. OR

OR: either ... or

```
8 •  SELECT *  
9    FROM employees  
0     WHERE first_name = "Denis" OR first_name = "Elvis";
```

→ the result is sorted by employee ids



Exercise: Retrieve a list with all employees whose first name is either Kellie or Aruna.

```
8 •   SELECT *
9     FROM employees
10    WHERE first_name = "Kellie" OR first_name = "Aruna";
```

36. Operator precedence

The logical order with which you must comply when you use both operators in the same WHERE block

→ the operator AND is applied first and OR is applied second. (AND > OR)

Example: extract all denis regardless of their genders

Problem:

```
1 •   SELECT
2     *
3     FROM
4       employees
5     WHERE
6       last_name = 'Denis' AND gender = 'M' OR gender = 'F';
```

+ all female individuals in the data
table, regardless of their family name

→ use parenthesis like python

```
1 • SELECT *
2   FROM
3     employees
4 WHERE
5   last_name = 'Denis' AND (gender = 'M' OR gender = 'F');
```

condition 1 condition 2

Exercise: Retrieve a list with all female employees whose first name is either Kellie or Aruna.

```
8 • SELECT *
9   FROM employees
10  WHERE gender = "F" AND (first_name = "Aruna" OR first_name = "Kellie");
```

37. IN and NOT IN

Example: any of the three names/conditions

```
1 • SELECT *
2   *
3   FROM
4     employees
5 WHERE
6   first_name = 'Cathie'
7     OR first_name = 'Mark'
8     OR first_name = 'Nathan';
```

A FASTER way: IN (list) ; NOT IN is just its negation

```
8 • SELECT *
9   FROM employees
10  WHERE first_name IN ("Cathie","Mark","Nathan");
```

Exercise (1): Use the IN operator to select all individuals from the “employees” table, whose first name is either “Denis”, or “Elvis”.

```
8 •  SELECT *
9   FROM employees
10  WHERE first_name IN ("Denis","Elvis");
```

(2) Everyone other than John, mark, or Jacob

```
8 •  SELECT *
9   FROM employees
10  WHERE first_name NOT IN ("John","Mark","Jacob");
```

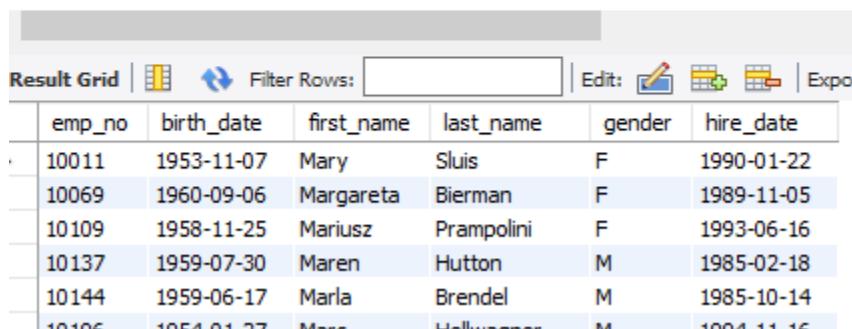
38. LIKE

Search for patterns. % means substitute of a sequence of characters.

%

Left to %: beginning

```
8 •  SELECT *
9   FROM employees
10  WHERE first_name LIKE ("Mar%");
```



The screenshot shows a database result grid with the following columns: emp_no, birth_date, first_name, last_name, gender, and hire_date. The data is as follows:

emp_no	birth_date	first_name	last_name	gender	hire_date
10011	1953-11-07	Mary	Sluis	F	1990-01-22
10069	1960-09-06	Margareta	Bierman	F	1989-11-05
10109	1958-11-25	Mariusz	Prampolini	F	1993-06-16
10137	1959-07-30	Maren	Hutton	M	1985-02-18
10144	1959-06-17	Marla	Brendel	M	1985-10-14
10196	1954-01-27	Marc	Hollingshead	M	1994-11-15

Right to %: ending

```

8 •  SELECT *
9   FROM employees
10  WHERE first_name LIKE ("%ar");

```

	emp_no	birth_date	first_name	last_name	gender	hire_date
▶	10029	1956-12-13	Otmar	Herbst	M	1985-11-20
	10074	1955-08-28	Mokhtar	Bernatsky	F	1990-08-13
	10175	1960-01-11	Aleksandar	Ananiadou	F	1988-01-11
	10266	1958-02-24	Sukumar	Rassart	M	1990-05-25
	10402	1953-07-23	Volkmar	Ebeling	M	1987-01-02

%ar%: ar is SOMEWHERE in the name strings

_: any single character

```

8 •  SELECT *
9   FROM employees
10  WHERE first_name LIKE ("Mar_");

```

	emp_no	birth_date	first_name	last_name	gender	hire_date
▶	10011	1953-11-07	Mary	Sluis	F	1990-01-22
	10196	1954-01-27	Marc	Hellwagner	M	1994-11-16
	10377	1954-08-19	Marl	Grospietsch	M	1990-05-07
	10415	1957-11-12	Mark	Coorg	M	1993-10-25
	10532	1959-08-31	Mary	Wossner	F	1986-05-18
	10921	1955-07-20	Mara	Rahi	F	1986-07-27

NOT LIKE:

The screenshot shows the MySQL Workbench interface. The SQL editor window has a title bar 'LIKE - NOT LIKE*' and contains the following query:

```
1 •  SELECT
2      *
3  FROM    employees
4  WHERE   first_name NOT LIKE ('%Mar%');
```

The Result Grid below displays the query results:

emp_no	birth_date	first_name	last_name	gender	hire_date
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10002	1964-06-02	Bezalel	Simmel	F	1985-11-21
10003	1959-12-03	Parto	Bamford	M	1986-08-28
10004	1954-05-01	Christian	Koblick	M	1986-12-01
10005	1955-01-21	Kvoichi	Maliniak	M	1989-09-12
10006	1953-04-20	Anneke	Preusia	F	1989-06-02
10007	1957-05-23	Tzvetan	Zielinski	F	1989-02-10
10008	1958-02-19	Saniva	Kalloufi	M	1994-09-15
10009	1952-04-19	Sumant	Peac	F	1985-02-18
10010	1963-06-01	Duanckaew	Piveteau	F	1989-08-24

The Result Grid toolbar includes buttons for 'Result Grid', 'Filter Rows', 'Export', 'Wrap Cell Content', and 'Fetch rows'.

MySQL is case insensitive!!!

Exercise: Working with the “employees” table, use the LIKE operator to select the data about all individuals, whose first name starts with “Mark”; specify that the name can be succeeded by any sequence of characters.

Retrieve a list with all employees who have been hired in the year 2000.

Retrieve a list with all employees whose employee number is written with 5 characters, and starts with “1000”.

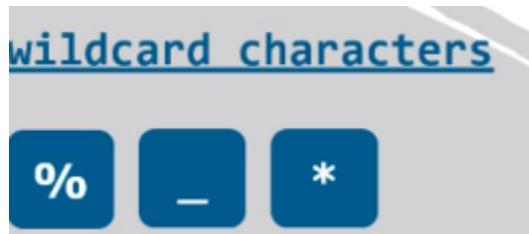
```

8 •  SELECT *
9   FROM employees
10  WHERE first_name LIKE ("Mark%");
11
12 •  SELECT *|  

13   FROM employees
14  WHERE hire_date LIKE ("2000%");
15
16 •  SELECT *
17   FROM employees
18  WHERE emp_no LIKE ("1000_");

```

39. Wildcard characters (regular expressions)



%: a sequence of characters

_: any single character

*: any character

Exercise: Extract all individuals from the ‘employees’ table whose first name contains “Jack”.

Once you have done that, extract another list containing the names of employees that do not contain “Jack”.

- `SELECT *`
`FROM employees`
`WHERE first_name LIKE ("%Jack%");`

- `SELECT *`
`FROM employees`
`WHERE first_name NOT LIKE ("%Jack%");`

40. BETWEEN ... AND ...

This expression helps us designate the interval to which a given value belongs

Example: between two dates:

```
SELECT
  *
FROM
  employees
WHERE
  hire_date BETWEEN '1990-01-01' AND '2000-01-01';
```

Both dates should be included

- ▶

```
SELECT *
FROM employees
WHERE hire_date BETWEEN "1990-01-01" AND "2000-01-01";
```

NOT BETWEEN ... AND ... has the opposite function.

```
SQL </> SELECT
  *
FROM
  employees
WHERE
  hire_date NOT BETWEEN '1990-01-01' AND '2000-01-01';
```

- the `hire_date` is *before* '1990-01-01'
or
- the `hire_date` is *after* '2000-01-01'

Exercise - use between and for strings

Select all the information from the “salaries” table regarding contracts from 66,000 to 70,000 dollars per year.

Retrieve a list with all individuals whose employee number is not between ‘10004’ and ‘10012’.

Select the names of all departments with numbers between ‘d003’ and ‘d006’.

```
8 •   SELECT *
9     FROM salaries
10    WHERE salary BETWEEN 66000 AND 70000;
11
12 •  SELECT *
13    FROM employees
14   WHERE emp_no NOT BETWEEN "10004" AND "10012";
15
16 •  SELECT *
17    FROM departments
18   WHERE dept_name BETWEEN "d003" AND "d006";
```

41. IS NOT NULL / IS NULL

IS NOT NULL
used to extract values that are not null

SQL

```
SELECT column_1, column_2,... column_n
FROM table_name
WHERE column_name IS NOT NULL;
```

Example: is anyone NULL?

```
16 •  SELECT *
17    FROM employees
18   WHERE first_name IS NOT NULL;
```

```
16 •  SELECT *
17    FROM employees
18   WHERE first_name IS NULL;
```

Result Grid					
emp_no	birth_date	first_name	last_name	gender	hire_date

→ no one was not entered

Exercise: Select the names of all departments whose department number value is not null.

```
17 •  SELECT *
18    FROM departments
19   WHERE dept_no IS NOT NULL;
```

Result Grid	
dept_no	dept_name
d009	Customer Service
d005	Development
d002	Finance
d003	Human Resources
d001	Marketing

42. Other comparison operators

So far:

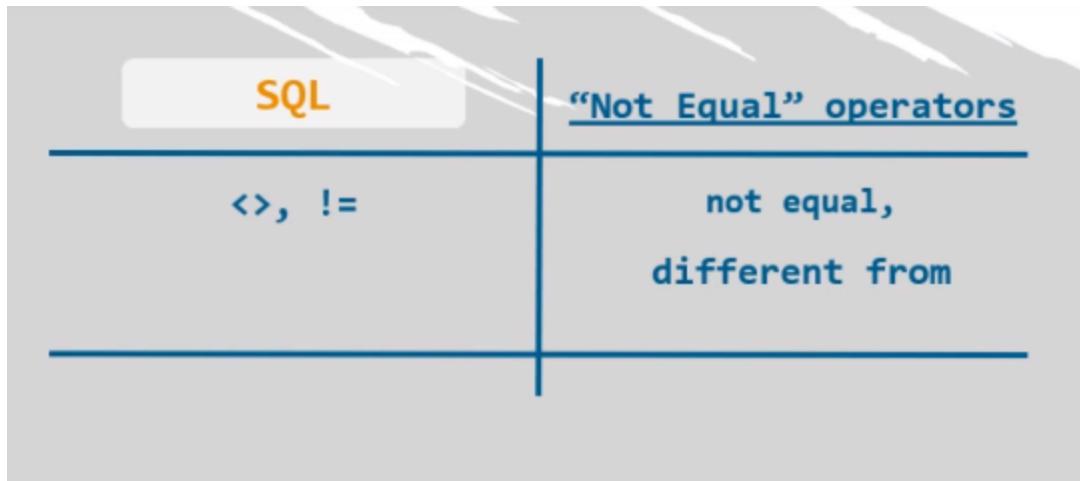
- BETWEEN... AND...
- LIKE
- IS NOT NULL
- NOT LIKE
- IS NULL

In this lecture:

=, >, >=, =<, <, <>, !=

SQL	
=	equal to
>	greater than
>=	greater than or equal to
<	less than
<=	less than or equal to

Not equal: either edgy brackets or !=



Ex1: no Mark

```
▶ SELECT
  *
FROM
  employees
WHERE
  first_name <> 'Mark';
```

Ex2: a time after 2000-1-1

- ```
SELECT
 *
FROM
 employees
WHERE
 hire_date > "2000-01-01";
```

Exercise:

Retrieve a list with data about all female employees who were hired in the year 2000 or after.

Hint: If you solve the task correctly, SQL should return 7 rows.

Extract a list with all employees' salaries higher than \$150,000 per annum.

```
SELECT
*
FROM
 employees
WHERE
 gender = "F" AND hire_date >= "2000-01-01";
```

```
SELECT
*
FROM
 salaries
WHERE
 salary >= 150000;
```

### 43. SELECT DISTINCT

The SELECT statement can retrieve from a designated column, given some criteria.

What if there are DUPLICATED values? The example shows repetition.

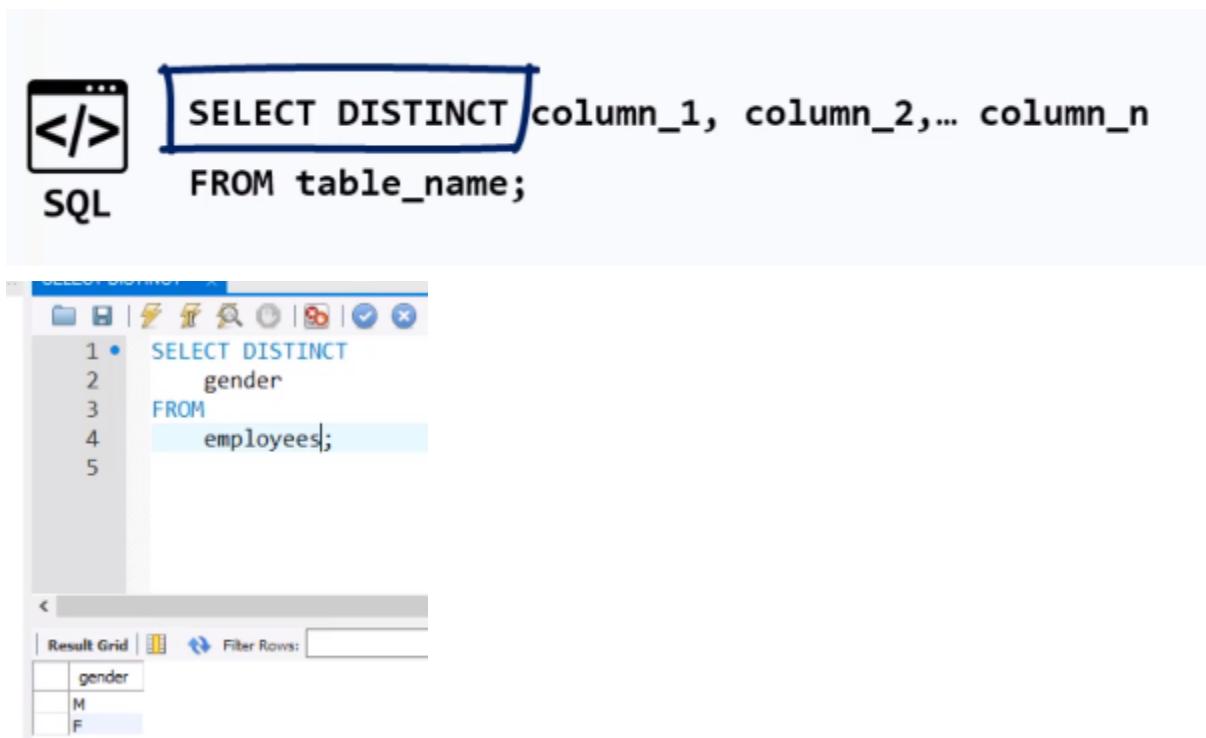
The screenshot shows the MySQL Workbench interface with a query editor titled "SELECT DISTINCT". The query is:

```
1 • SELECT
2 gender
3 FROM
4 employees;
```

The results grid displays the "gender" column with the following data:

| gender |
|--------|
| M      |
| M      |
| M      |
| F      |
| F      |
| F      |
| M      |
| F      |
| M      |
| M      |

## SELECT NON-REPETITIVE VALUES



The screenshot shows a SQL query window in SSMS. The query is:

```
1 • SELECT DISTINCT
2 gender
3 FROM
4 employees;
```

The result grid shows the following data:

| gender |
|--------|
| M      |
| F      |

Exercise: Obtain a list with all different “hire dates” from the “employees” table. Expand this list and click on “Limit to 1000 rows”. This way you will set the limit of output rows displayed back to the default of 1000.

In the next lecture, we will show you how to manipulate the limit rows count.

## SELECT DISTINCT

**hire\_date**

**FROM**

**employees;**

## 44. Introduction to Aggregate functions

Aggregate functions: they are applied on multiple rows of a single column of a table and return an output of a single value.

**These functions ignore NULL values unless told otherwise.**

COUNT()  
SUM()  
MIN()  
MAX()  
AVG()

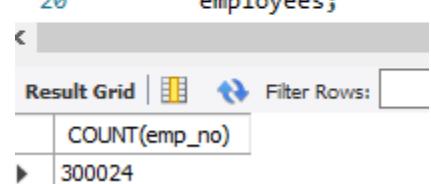
COUNT() is often used with DISTINCT

```
SELECT COUNT(column_name)
FROM table_name;
```

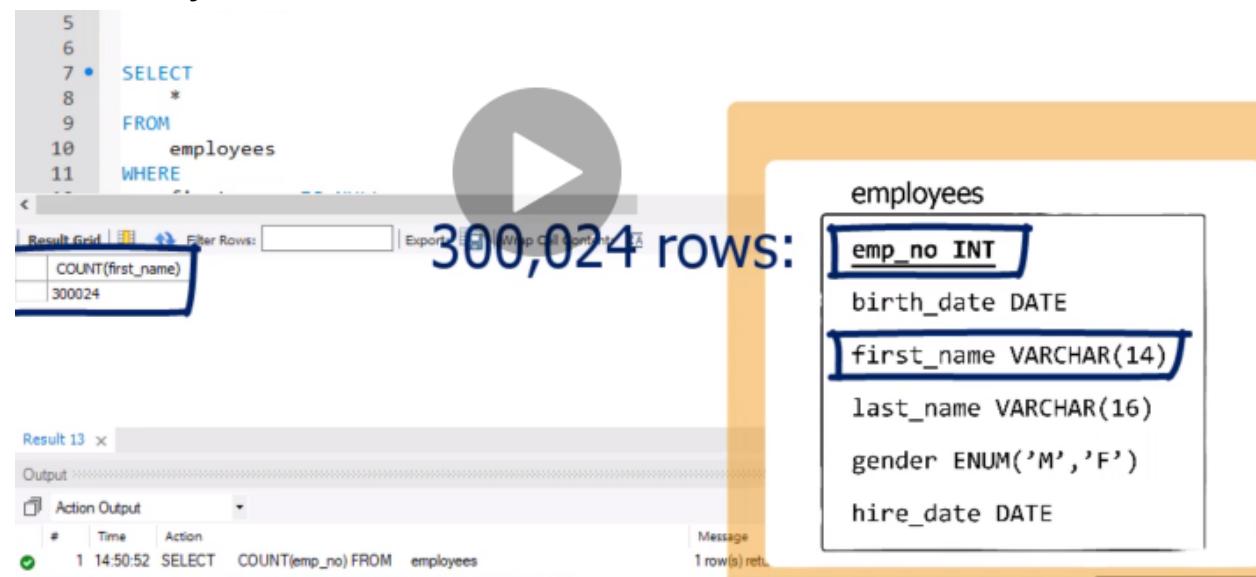
Note: the parenthesis must come right after the keyword, so there is no whitespace

Ex1: how many employees are registered in our database?

```
17 • SELECT
18 COUNT(emp_no)
19 FROM
20 employees;
```



A second way to show that there is no NULL value:



employees

|            |               |
|------------|---------------|
| emp_no     | INT           |
| birth_date | DATE          |
| first_name | VARCHAR(14)   |
| last_name  | VARCHAR(16)   |
| gender     | ENUM('M','F') |
| hire_date  | DATE          |

Result 13 x

Output

Action Output

# Time Action

1 14:50:52 SELECT COUNT(emp\_no) FROM employees

Message 1 row(s) ret.

employee# and first\_name# are the same

COUNT(DISTINCT) → must be in the parenthesis

```
SELECT
 COUNT(distinct first_name)
FROM
 employees;
```

**Exercise:**

**How many annual contracts with a value higher than or equal to \$100,000 have been registered in the salaries table?**

```
17 • SELECT
18 COUNT(salary)
19 FROM
20 salaries
21 WHERE salary >=100000;

| Result Grid | Filter Rows:
|
| COUNT(salary)|
| 32207 |
```

OR:

```
17 • SELECT
18
19 COUNT(*)
20
21 FROM
22
23 salaries
24
25 WHERE
26
27 salary >= 100000;
```

**How many managers do we have in the “employees” database? Use the star symbol (\*) in your code to solve this exercise.**

```
17 • SELECT
18
19 COUNT(*)
20
21 FROM
22
23 dept_manager
24 ;
```

The screenshot shows a MySQL Workbench interface. At the top, there's a toolbar with various icons. Below it is a query editor window containing the provided SQL code. To the right of the editor is a results grid titled "Result Grid". The results show a single row with the column name "COUNT(\*)" and the value "24". There are also navigation buttons like back, forward, and search.

| COUNT(*) |
|----------|
| 24       |

## 45. ORDER BY

Use ORDER BY to sort the results according to some column

**ASC:** ascending (default)

**DESC:** descending

The screenshot shows a MySQL Workbench interface. The query editor contains the following SQL code:

```
1 • SELECT
2 *
3 FROM
4 employees
5 ORDER BY first_name DESC;
```

To the right of the editor is a results grid titled "Result Grid". The results show a list of employee records sorted by "first\_name" in descending order. The columns are: emp\_no, birth\_date, first\_name, last\_name, gender, and hire\_date. The data includes entries for Zvonko Bultermann, Zvonko Lund, Zvonko Soataro, Zvonko Maverwieser, Zvonko Peron, Zvonko Lakshmanan, Zvonko Gente, and Zvonko Pirkh.

| emp_no | birth_date | first_name | last_name   | gender | hire_date  |
|--------|------------|------------|-------------|--------|------------|
| 497370 | 1962-03-30 | Zvonko     | Bultermann  | M      | 1986-02-26 |
| 498646 | 1957-10-25 | Zvonko     | Lund        | F      | 1989-03-09 |
| 493637 | 1958-10-27 | Zvonko     | Soataro     | M      | 1989-12-25 |
| 494621 | 1962-07-18 | Zvonko     | Maverwieser | F      | 1987-01-29 |
| 494935 | 1956-06-05 | Zvonko     | Peron       | F      | 1985-06-19 |
| 484995 | 1964-11-04 | Zvonko     | Lakshmanan  | M      | 1992-12-04 |
| 486170 | 1959-12-30 | Zvonko     | Gente       | M      | 1986-08-08 |
| 487925 | 1958-04-12 | Zvonko     | Pirkh       | M      | 1994-08-31 |

You can order the results by more than one field by using comma

```
17 • SELECT *
18 FROM employees
19 ORDER BY first_name, last_name ASC;
20
```

|   | emp_no | birth_date | first_name | last_name | gender | hire_date  |
|---|--------|------------|------------|-----------|--------|------------|
| ▶ | 69256  | 1962-04-14 | Aamer      | Anger     | M      | 1998-03-16 |
|   | 486584 | 1952-08-12 | Aamer      | Armand    | M      | 1990-09-15 |
|   | 237165 | 1962-02-23 | Aamer      | Azevedo   | F      | 1991-06-28 |
|   | 413688 | 1955-06-26 | Aamer      | Azuma     | M      | 1989-12-10 |
|   | 281363 | 1956-05-18 | Aamer      | Baak      | F      | 1994-03-10 |
|   | 242368 | 1959-07-26 | Aamer      | Baaleh    | F      | 1989-08-06 |

**Exercise:** Select all data from the “employees” table, ordering it by “hire date” in descending order.

```
17 • SELECT *
18 FROM employees
19 ORDER BY hire_date DESC;
20
```

|   | emp_no | birth_date | first_name | last_name | gender | hire_date  |
|---|--------|------------|------------|-----------|--------|------------|
| ▶ | 463807 | 1964-06-12 | Bikash     | Covnot    | M      | 2000-01-28 |
|   | 428377 | 1957-05-09 | Yucai      | Gerlach   | M      | 2000-01-23 |
|   | 499553 | 1954-05-06 | Hideyuki   | Delgrande | F      | 2000-01-22 |
|   | 222965 | 1959-08-07 | Volkmar    | Perko     | F      | 2000-01-13 |
|   | 47291  | 1960-09-09 | Ulf        | Flexer    | M      | 2000-01-12 |
|   | 422990 | 1953-04-09 | Jaana      | Verspoor  | F      | 2000-01-11 |

## 46. GROUP BY

### AN IMPORTANT TOPIC

When working with SQL, results can be grouped according to a specific field or fields. GROUP BY must be placed immediately after the WHERE conditions, if any, and just before the ORDER BY clause.

```
SELECT column_name(s)
FROM table_name
WHERE conditions

GROUP BY column_name(s)
ORDER BY column_name(s);
```

Example:

An unsorted list of names...

The screenshot shows the MySQL Workbench interface. In the SQL editor window, the following query is written:

```
1 • SELECT
2 first_name
3 FROM
4 employees;
```

The result grid below displays the 'first\_name' column from the 'employees' table. The data is as follows:

| first_name |
|------------|
| Georgi     |
| Bezalel    |
| Porto      |
| Christian  |
| Kvoichi    |
| Anneke     |
| Tzvetan    |
| Saniva     |

- Group by `first_name` → only distinct values will be selected

```
1 • SELECT
2 first_name
3 FROM
4 employees
5 GROUP BY first_name;
```

Result Grid | Filter Rows: [ ] | Export: [ ] | Wrap Cell Content:

| first_name |
|------------|
| Georgi     |
| Bezalel    |
| Parto      |
| Christian  |
| Kvoichik   |
| Anneke     |
| Tsvetan    |
| Saniva     |

employees 3 ×

Output ::::::::::::

Same as:

```
SELECT DISTINCT
 first_name
FROM
 employee;
```

Just sorted in a different way

In most cases, when you need an aggregate function, you must add a GROUP BY clause in the query too.

GROUP BY x

```
1 • SELECT
2 first_name, COUNT(first_name)
3 FROM
4 employees
5 GROUP BY first_name;
6
7
```

Result Grid | Filter Rows: [ ] | Export: | Wrap Cell Content: | Fetch rows: [ ]

| first_name | COUNT(first_name) |
|------------|-------------------|
| Aamer      | 228               |
| Aamod      | 216               |
| Abdelaziz  | 227               |
| Abdelhani  | 247               |
| Abdelkader | 222               |
| Abdelwahab | 241               |
| Abdulah    | 220               |
| Abdulla    | 226               |

Result 7 x

Output: .....

**Keep the names the same in both SELECT and GROUP BY.**

#### 47. Using Aliases (AS)

Problem with the second column name → need an intuitive name

```
17 • SELECT first_name, COUNT(first_na
18 FROM employees
19 GROUP BY first_name
20 ORDER BY first_name;
21
22
```

Result Grid | Filter Rows: [ ] | E

| first_name | COUNT(first_name) |
|------------|-------------------|
| Zvonko     | 258               |
| Zsolt      | 236               |
| Zorica     | 225               |
| Zongyan    | 254               |
| Ziyad      | 229               |

We need an alias name to rename a selection of the query. → AS

```
17 • SELECT first_name, COUNT(first_name) as names_count
18 FROM employees
19 GROUP BY first_name
20 ORDER BY first_name;
21
22
```

The screenshot shows a database query results grid. At the top, there are buttons for 'Result Grid', 'Filter Rows', 'Export', and 'Wrap Cell Content'. The grid has two columns: 'first\_name' and 'names\_count'. The data rows are:

| first_name | names_count |
|------------|-------------|
| Aamer      | 228         |
| Aamod      | 216         |
| Abdelaziz  | 227         |
| Abdelghani | 247         |
| Abdelkader | 222         |
| Abdelwaheb | 241         |
| Abdulrah   | 220         |

**Exercise:** This will be a slightly more sophisticated task.

Write a query that obtains two columns. The first column must contain annual salaries higher than 80,000 dollars. The second column, renamed to “emps\_with\_same\_salary”, must show the number of employees contracted to that salary. Lastly, sort the output by the first column.

```
17 • SELECT salary, COUNT(emp_no) as emps_with_same_salary
18 FROM salaries
19 WHERE salary > 80000
20 GROUP BY salary
21 ORDER BY salary;
```

The screenshot shows a database query results grid. At the top, there are buttons for 'Result Grid', 'Filter Rows', 'Export', and 'Wrap Cell Content'. The grid has two columns: 'salary' and 'emps\_with\_same\_salary'. The data rows are:

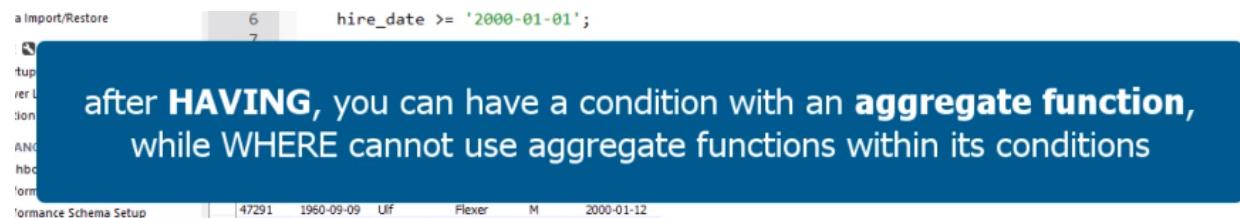
| salary | emps_with_same_salary |
|--------|-----------------------|
| 80001  | 12                    |
| 80002  | 10                    |
| 80003  | 10                    |
| 80004  | 13                    |
| 80005  | 15                    |
| 80006  | 8                     |
| 80007  | 8                     |

## 48. HAVING

Having refines the output from record that do not satisfy a certain condition; frequently implemented with GROUP BY.

```
SELECT column_name(s)
FROM table_name
WHERE conditions
GROUP BY column_name(s)
HAVING conditions
ORDER BY column_name(s);
```

It appears between GROUP BY and ORDER BY.



The screenshot shows a MySQL Workbench interface with a query editor. The query is:

```
6 hire_date >= '2000-01-01';
7
```

A callout box highlights the **HAVING** clause with the following text:

after **HAVING**, you can have a condition with an **aggregate function**,  
while WHERE cannot use aggregate functions within its conditions

Below the query editor, the status bar displays: Import/Restore, 47291, 1960-09-09, Uf, Flexer, M, 2000-01-12.

If putting WHERE with an aggregation function, there will be an error (invalid use of group function). This question asks “extract all first names that appears more than 250 times in the employee table.”

```

19
20
21
22
23 • SELECT
24 first_name, COUNT(first_name) AS names_count
25 FROM
26 employees
27 WHERE
28 COUNT(first_name) > 250
29 GROUP BY first_name
30 ORDER BY first_name;
31
32

```

Output:

| # | Time     | Action                                                                                  | Message                                          |
|---|----------|-----------------------------------------------------------------------------------------|--------------------------------------------------|
| 1 | 18:05:25 | SELECT * FROM employees WHERE hire_date >= '2000-01-01'                                 | 13 row(s) returned                               |
| 2 | 18:05:30 | SELECT * FROM employees HAVING hire_date >= '2000-01-01'                                | 13 row(s) returned                               |
| 3 | 18:06:19 | SELECT first_name, COUNT(first_name) AS names_count FROM employees ORDER BY first_name; | Error Code: 1111. Invalid use of group function. |

Need to insert having between group and order.

```

• SELECT
 first_name, COUNT(first_name) AS names_count
FROM
 employees
GROUP BY first_name
HAVING COUNT(first_name) > 250
ORDER BY first_name;

```

When you need HAVING: when the condition has an aggregation function

Startup / Shutdown  
Server Logs  
Options File  
**PERFORMANCE**  
Dashboard  
Performance Report  
Performance Schema

"Extract all first names that appear more than 250 times in the "employees" table."

COUNT() → an aggregate function

Information\_schema  
mysql  
performance\_schema

Action Output

### Exercise:

Select all employees whose average salary is higher than \$120,000 per annum.

Hint: You should obtain 101 records.

**Compare the output you obtained with the output of the following two queries:**

**SELECT**

**\*, AVG(salary)**

**FROM**

**salaries**

**WHERE**

**salary > 120000**

**GROUP BY emp\_no**

**ORDER BY emp\_no;**

**SELECT**

**\*, AVG(salary)**

**FROM**

**salaries**

**GROUP BY emp\_no**

**HAVING AVG(salary) > 120000;**

When using WHERE instead of HAVING, the output is larger because in the output we include individual contracts higher than \$120,000 per year. The output does not contain average salary values.

#### 49. WHERE VS. HAVING

WHERE allows us to set conditions that refer to **subsets of individual rows**. The conditions are applied before re-organizing the output into groups.



WHERE and HAVING are two kinds of sorting filters.

Example 1: extract a list of all names that are encountered less than 200 times. Let the data refer to people hired after the 1st of January 1999.

```

31 • SELECT first_name, COUNT(first_name) as names_count
32 FROM employees
33 WHERE hire_date > "1999-01-01"
34 GROUP BY first_name
35 HAVING count(first_name) < 200
36 ORDER BY first_name DESC;
37

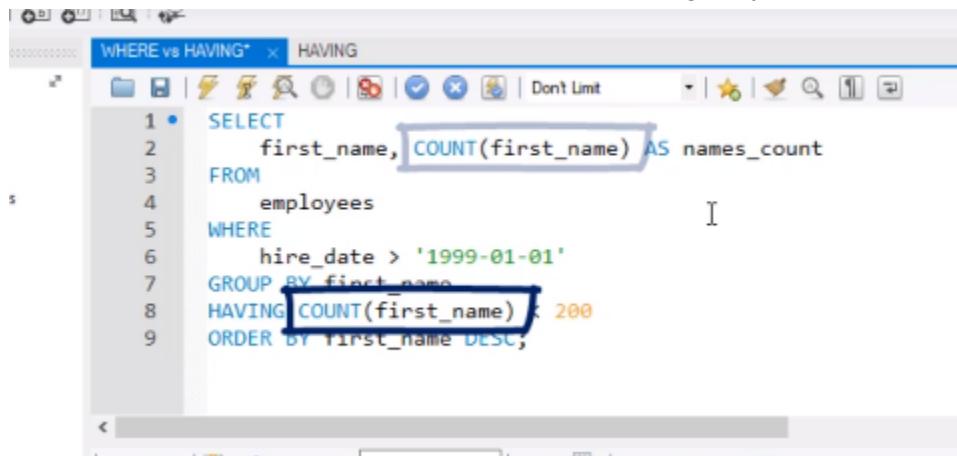
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

| first_name | names_count |
|------------|-------------|
| Zvonko     | 2           |
| Zsolt      | 2           |
| Zorica     | 1           |
| Zongyan    | 1           |
| Ziya       | 1           |
| Zissis     | 2           |

Observations:

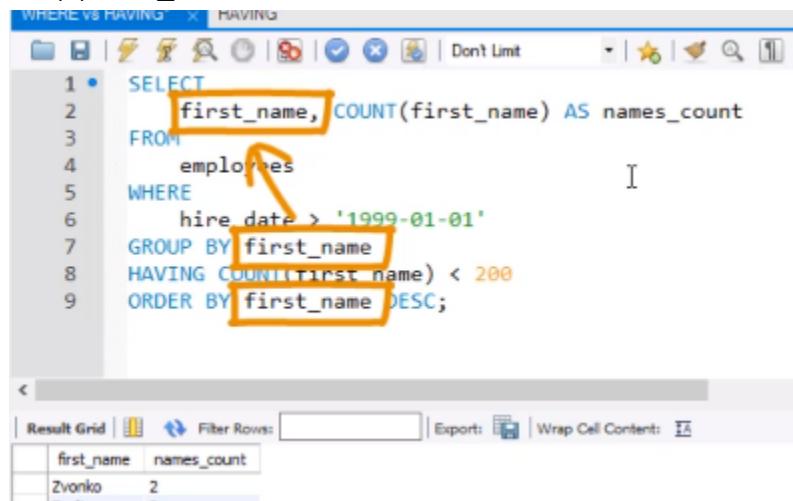
- (1) COUNT has been applied twice, which are logically related



The screenshot shows a SQL editor window titled "WHERE vs HAVING" with a tab labeled "HAVING". The query is as follows:

```
1 • SELECT
2 first_name, COUNT(first_name) AS names_count
3 FROM
4 employees
5 WHERE
6 hire_date > '1999-01-01'
7 GROUP BY first_name
8 HAVING COUNT(first_name) < 200
9 ORDER BY first_name DESC;
```

- (2) First\_name has been used three times



The screenshot shows a SQL editor window titled "WHERE vs HAVING" with a tab labeled "HAVING". The query is as follows:

```
1 • SELECT
2 first_name, COUNT(first_name) AS names_count
3 FROM
4 employees
5 WHERE
6 hire_date > '1999-01-01'
7 GROUP BY first_name
8 HAVING COUNT(first_name) < 200
9 ORDER BY first_name DESC;
```

A yellow arrow points from the first occurrence of "first\_name" in the SELECT clause to the first occurrence in the HAVING clause.

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| first_name | names_count |
|------------|-------------|
| Zvonko     | 2           |

- (3) The condition for WHERE cannot be put under HAVING because of two different types of aggregation

```

30
31 • SELECT first_name, COUNT(first_name) as names_count
32 from employees
33 #WHERE hire_date > "1999-01-01"
34 GROUP BY first_name
35 HAVING count(first_name) < 200 AND hire_date > "1999-01-01"
36 ORDER BY first_name DESC;
37
38
39

```

**Output**

| # | Time       | Action                                                                | Message                                                         |
|---|------------|-----------------------------------------------------------------------|-----------------------------------------------------------------|
| ✓ | 1 14:51:11 | SELECT first_name, COUNT(first_name) as names_count from employees... | 886 row(s) returned                                             |
| ✓ | 2 14:51:23 | SELECT first_name, COUNT(first_name) as names_count from employees... | 886 row(s) returned                                             |
| ✗ | 3 14:54:15 | SELECT first_name, COUNT(first_name) as names_count from employees... | Error Code: 1054. Unknown column 'hire_date' in 'having clause' |

### HAVING can take non-aggregated conditions below

The screenshot shows a MySQL Workbench interface with a query editor and a results pane.

**Query Editor:**

```

1 • SELECT
2 *
3 FROM
4 employees
5 HAVING
6 hire_date >= '2000-01-01';
7
8
9
10 •
11
12
13
14
15
16

```

A red box highlights the `HAVING hire_date >= '2000-01-01';` line. A large red X is drawn over the entire section from line 10 to line 16, covering the rest of the query.

**Results Pane:**

The results pane shows the output of the query:

```

non-aggregated

```

**Conclusion:**

Having conditions can be either aggregated or non-aggregated.

Conclusion:

Aggregate functions - GROUP BY and HAVING

General conditions - WHERE

```
SELECT column_name(s)
FROM table_name
WHERE conditions
GROUP BY column_name(s)
HAVING conditions
ORDER BY column_name(s);
```

**Exercise:** Select the employee numbers of all individuals who have signed more than 1 contract after the 1st of January 2000.

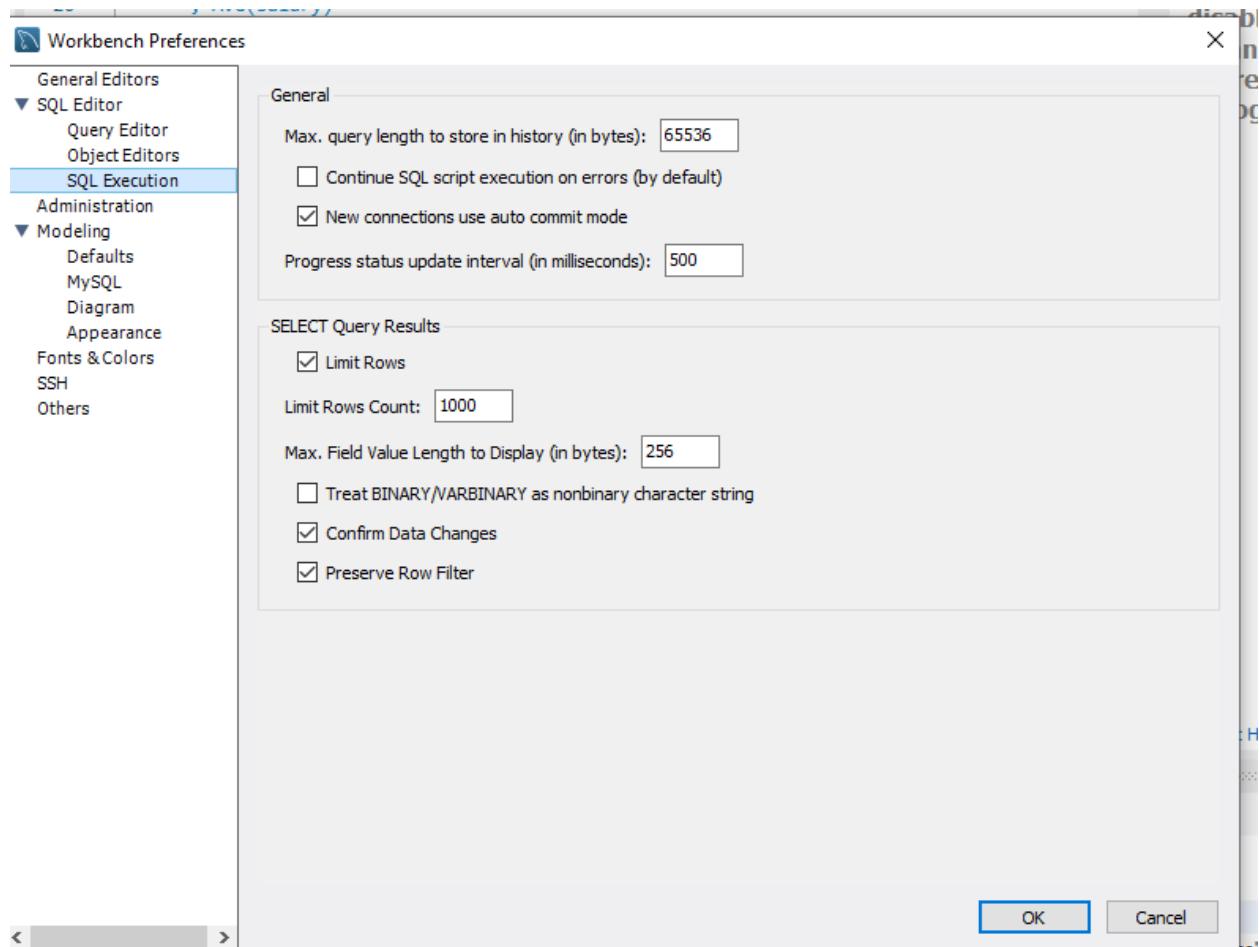
**Hint:** To solve this exercise, use the “dept\_emp” table.

```
31 • SELECT emp_no
32 from dept_emp
33 WHERE from_date > "2000-01-01"
34 GROUP BY emp_no
35 HAVING COUNT(from_date) > 1
36 ORDER BY emp_no;
37
38
```

One date is one contract by an employee!

50. Limit

You can change the default limit of outputs in edit>preferences



Ex1: please show me the employee numbers of the **10 highest** paid employees in the database

10 highest ORDER BY STANDARD DESC

```
31 • SELECT *
32 FROM salaries
33 ORDER BY salary DESC
34 LIMIT 10;
35
```

Result Grid | Filter Rows: Edit:

|   | emp_no | salary | from_date  | to_date    |
|---|--------|--------|------------|------------|
| ▶ | 43624  | 158220 | 2002-03-22 | 9999-01-01 |
| ▶ | 43624  | 157821 | 2001-03-22 | 2002-03-22 |
| ▶ | 47978  | 155709 | 2002-07-14 | 9999-01-01 |
| ▶ | 109334 | 155377 | 2000-02-12 | 2001-02-11 |
| ▶ | 109334 | 155190 | 2002-02-11 | 9999-01-01 |

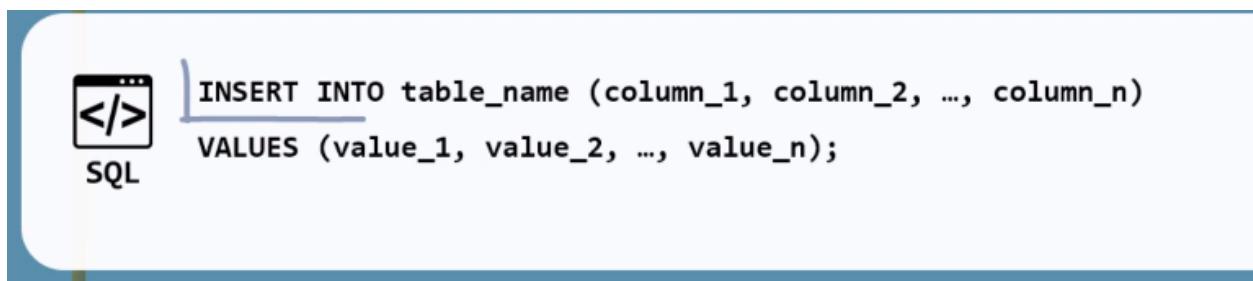
salaries 6 ×

Exercise: select the first 100 rows from the dept\_emp table

```
• SELECT *
 FROM dept_emp
 LIMIT 100;
```

## Section 9: INSERT statement

51. INSERT: Part 1



Ex:

(1) Add a record

**Integer can be written in quotes too (MySQL can automatically convert the strings into numbers)**

- `INSERT INTO employees`
  - `(`
    - `emp_no,`
    - `birth_date,`
    - `first_name,`
    - `last_name,`
    - `gender,`
    - `hire_date`
  - `) VALUES`
  - `(`
    - `9999901,`
    - `"1986-04-21",`
    - `"John",`
    - `"Smith",`
    - `"M",`
    - `"2011-01-01"`
  - `);`

(2) Check by using ORDER BY

```

31 • SELECT *
32 FROM employees
33 ORDER BY emp_no DESC
34 LIMIT 10
35
36
37 • INSERT INTO employees
38 (
39 emp_no,
40 birth_date,
41 first_name,
42 last_name,

```

Result Grid | Filter Rows:  | Edit: | Export/Import

|   | emp_no  | birth_date | first_name | last_name    | gender | hire_date  |
|---|---------|------------|------------|--------------|--------|------------|
| ▶ | 9999901 | 1986-04-21 | John       | Smith        | M      | 2011-01-01 |
|   | 999901  | 1986-04-21 | John       | Smith        | M      | 2011-01-01 |
|   | 499999  | 1958-05-01 | Sachin     | Tsukuda      | M      | 1997-11-30 |
|   | 499998  | 1956-09-05 | Patricia   | Breugel      | M      | 1993-10-13 |
|   | 499997  | 1961-08-03 | Berhard    | Lenart       | M      | 1986-04-21 |
|   | 499996  | 1953-03-07 | Zito       | Baaz         | M      | 1990-09-27 |
|   | 499995  | 1958-09-24 | Dekang     | Lichtner     | F      | 1993-01-12 |
|   | 499994  | 1952-02-26 | Navin      | Argence      | F      | 1990-04-24 |
|   | 499993  | 1963-06-04 | DeForest   | Mullainathan | M      | 1997-04-07 |
|   | 499992  | 1960-10-12 | Siamak     | Salverda     | F      | 1987-05-10 |
| * | NULL    | NULL       | NULL       | NULL         | NULL   | NULL       |

## 52. INSERT: Part 2

Within the columns, the order doesn't have to be the same as the tables

- specify as many data values as there are columns in the data table

- add them in the same order in which they appear in the table

#### Exercise:

Select ten records from the “titles” table to get a better idea about its content.

Then, in the same table, insert information about employee number 999903. State that he/she is a “Senior Engineer”, who has started working in this position on October 1st, 1997.

At the end, sort the records from the “titles” table in descending order to check if you have successfully inserted the new record.

**Hint: To solve this exercise, you'll need to insert data in only 3 columns!**

Don't forget, we assume that, apart from the code related to the exercises, you always execute all code provided in the lectures. This is particularly important for this exercise. If you have not run the code from the previous lecture, called ‘The INSERT Statement – Part II’, where you have to insert information about employee 999903, you might have trouble solving this exercise!

#### Code:

**INSERT INTO employees**

**VALUES**

**(**

**999903,**

**'1977-09-14',**

**'Johnathan',**

**'Creek',**

**'M',**

**'1999-01-01'**

**);**

No need to fill the to\_date column

```

31 • SELECT *
32 FROM titles
33 ORDER BY emp_no DESC
34 LIMIT 10
35
36
37 • INSERT INTO titles
38 (
39 emp_no,
40 from_date,
41 title
42) VALUES
43 (
44 999903,
45 "1997-10-01",
46 "Senior Engineer"
47);
48
49 • SELECT *
50 FROM titles
51 ORDER BY emp_no DESC;

```

The screenshot shows the MySQL Workbench interface with the 'Result Grid' tab selected. The grid displays the 'titles' table with the following data:

|   | emp_no | title           | from_date  | to_date    |
|---|--------|-----------------|------------|------------|
| ▶ | 999903 | Senior Engineer | 1997-10-01 | NULL       |
|   | 499999 | Engineer        | 1997-11-30 | 9999-01-01 |
|   | 499998 | Staff           | 1993-12-27 | 1998-12-27 |
|   | 499998 | Senior Staff    | 1998-12-27 | 9999-01-01 |
|   | 499997 | Senior Engineer | 1992-08-29 | 9999-01-01 |
|   | 499997 | Engineer        | 1987-08-30 | 1992-08-29 |
|   | 499996 | Senior Engineer | 2002-05-13 | 9999-01-01 |
|   | 499996 | Engineer        | 1996-05-13 | 2002-05-13 |
|   | 499995 | Engineer        | 1997-06-02 | 9999-01-01 |

### Exercise:

Insert information about the individual with employee number 999903 into the “dept\_emp” table. He/She is working for department number 5, and has started work on October 1st, 1997; her/his contract is for an indefinite period of time.

Hint: Use the date ‘9999-01-01’ to designate the contract is for an indefinite period.

```
31 • SELECT *
32 FROM dept_emp
33 ORDER BY emp_no DESC
34 LIMIT 10
35 ;
36
37 • INSERT INTO dept_emp
38 (
39 emp_no,
40 dept_no,
41 from_date,
42 to_date
43) VALUES
44 (
45 999903,
46 "d005",
47 "1997-10-01",
48 "9999-01-01"
49);
```

### 53. Inserting data INTO a new table

Insert data from one table into another

The diagram shows the SQL syntax for inserting data from one table into another. The code is:

```
INSERT INTO table_2(column_1, column_2, ..., column_n)
SELECT column_1, column_2, ..., column_n
FROM table_1
WHERE condition;
```

Annotations highlight specific parts of the query:

- A blue oval encloses the table name `table_2` in the `INSERT INTO` clause.
- A blue oval encloses the column names `column_1, column_2, ..., column_n` in the `SELECT` clause.
- An orange oval encloses the table name `table_1` in the `FROM` clause.

**Example: create a new departments table and insert data into it**

There is no condition for this data transfer/copying

First, create an empty table and check it

```
38 • CREATE TABLE departments_dup
39 ((
40 dept_no CHAR(4) NOT NULL,
41 dept_name VARCHAR(40) NOT NULL
42)
43 ;
44
45 • SELECT *
46 FROM departments_dup;
47
```

Then make an insertion

```
48 • INSERT INTO departments_dup
49 ((
50 dept_no,
51 dept_name
52)
53 SELECT *
54 FROM departments;
55
```

**Exercise:** Create a new department called “Business Analysis”. Register it under number ‘d010’.

**Hint:** To solve this exercise, use the “departments” table.

There is no table-to-table insertion. Also, you can omit the columns by using the default order

```
INSERT INTO departments
VALUES("d010","Business Analytics");
```

## Section 10: UPDATE statement

54. TCL’s COMMIT and ROLLBACK

Transaction control language

## COMMIT statement

Saves the transaction in the database

Changes cannot be undone

→ used to save the data in the database at the moment of execution

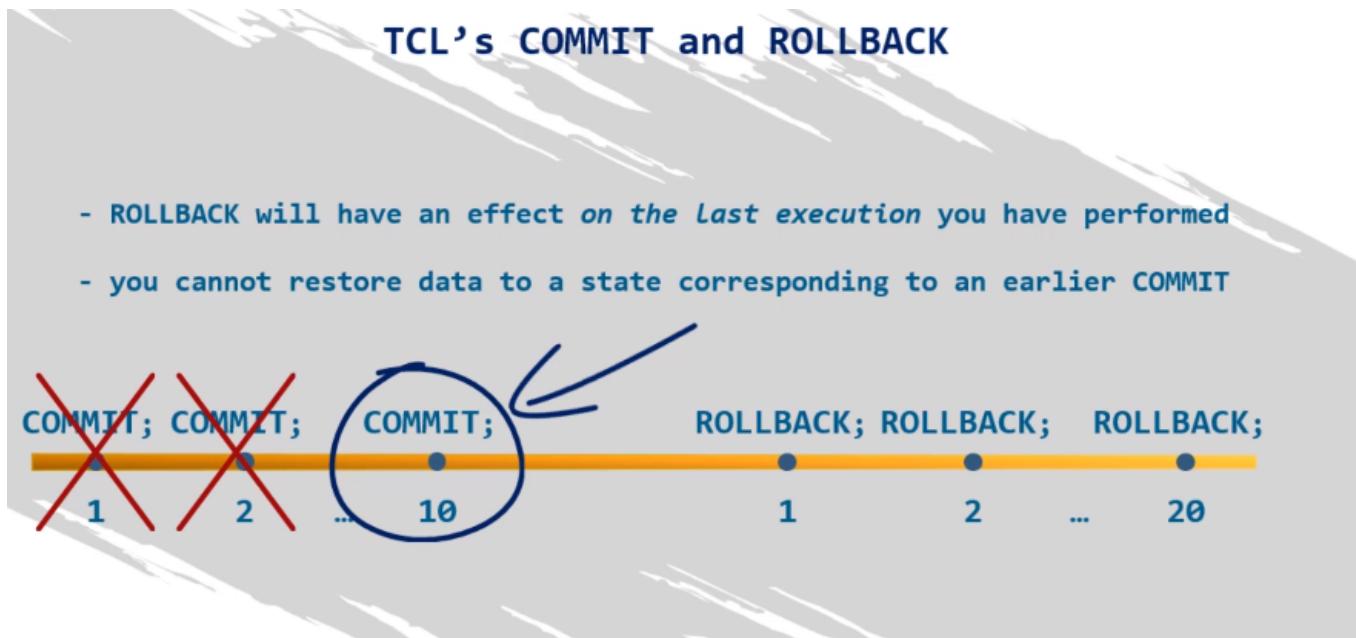
## ROLLBACK

Allows you to take a step back (undo)

The last changes made will not count

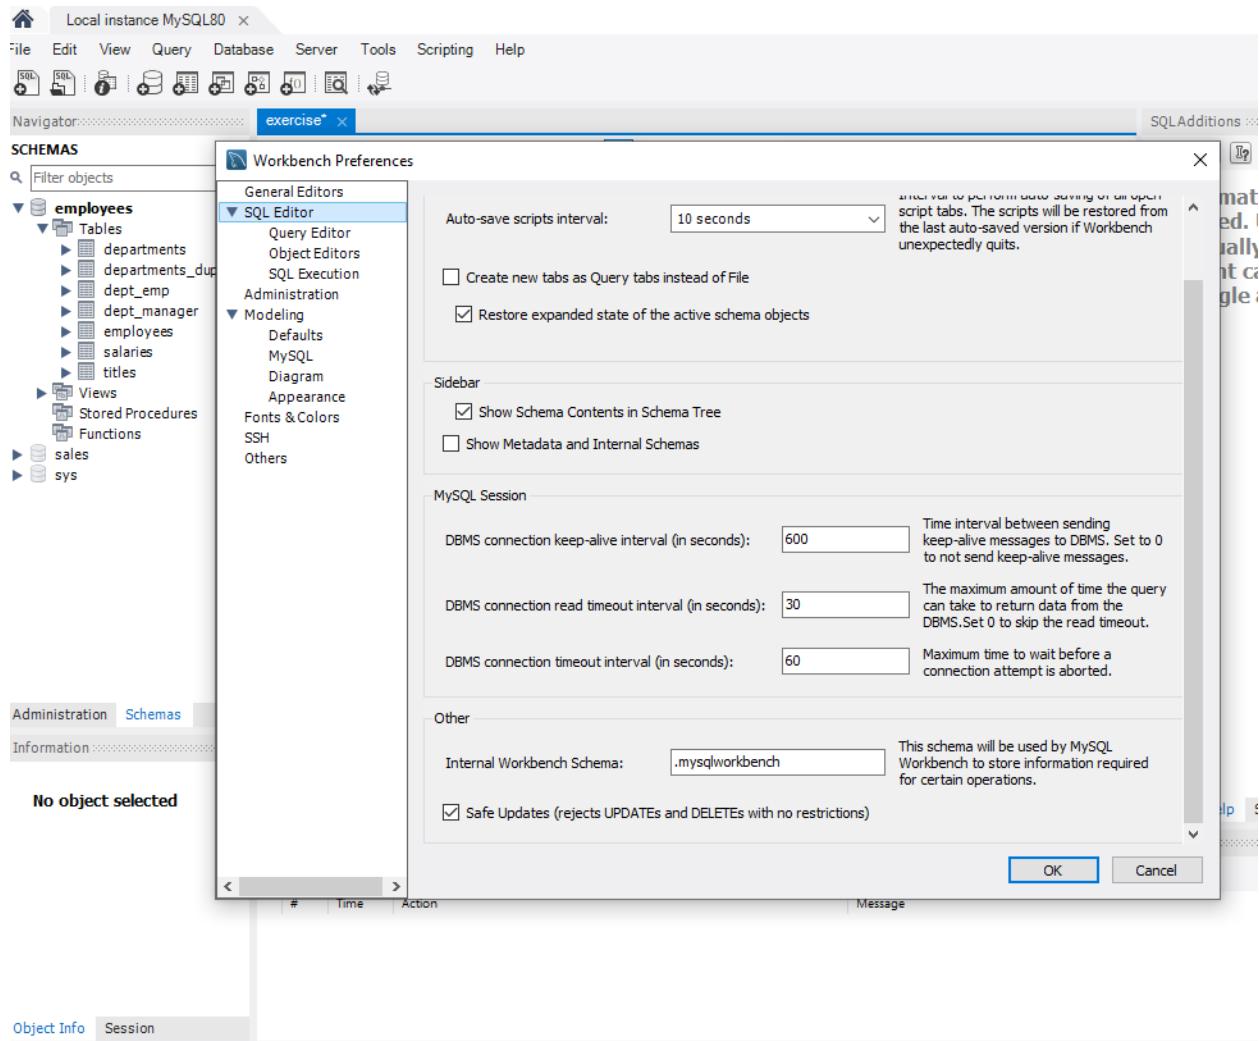
Recerts to the last non-committed state

→ it will refer to the state corresponding to the last time you execute the COMMIT



ROLLBACK (1-20) only points to the last commit, not earlier ones.

Edit > preferences > SQL editor > safe updates



Select yes means (1) the software can prevent you from losing significant amount of data, and (2) can eliminate the possibility of controlling the state of the data → you should turn it off for operations

## 55. UPDATE

- the UPDATE Statement

used to update the values of existing records in a table



```
UPDATE table_name
SET column_1 = value_1, column_2 = value_2 ...
WHERE conditions;
```



Look for an earlier record

```
37 • USE employees;
38 • SELECT *
39 FROM employees
40 WHERE emp_no = 9999901;
41
```

Result Grid | Filter Rows:  | Edit: | Export

|   | emp_no  | birth_date | first_name | last_name | gender | hire_date  |
|---|---------|------------|------------|-----------|--------|------------|
| . | 9999901 | 1986-04-21 | John       | Smith     | M      | 2011-01-01 |
| • | NULL    | NULL       | NULL       | NULL      | NULL   | NULL       |

```

37 • USE employees;
38 • SELECT *
39 FROM employees
40 WHERE emp_no = 9999901;
41
42 • UPDATE employees
43 SET
44 first_name = "Stella",
45 last_name = "Parkinson",
46 birth_date = "1990-12-31",
47 gender = "F"
48 WHERE
49 emp_no = 9999901;

```

The screenshot shows a MySQL Workbench interface with a result grid. The grid has columns for emp\_no, birth\_date, first\_name, last\_name, gender, and hire\_date. There are two rows: one for the primary key (9999901) and one for a secondary row marked with an asterisk (\*). The primary key row contains valid data: birth\_date 1990-12-31, first\_name Stella, last\_name Parkinson, gender F, and hire\_date 2011-01-01. The secondary row contains NULL values for all columns.

|   | emp_no  | birth_date | first_name | last_name | gender | hire_date  |
|---|---------|------------|------------|-----------|--------|------------|
| ▶ | 9999901 | 1990-12-31 | Stella     | Parkinson | F      | 2011-01-01 |
| * | NULL    | NULL       | NULL       | NULL      | NULL   | NULL       |

- If “WHERE emp\_no = 9999922111” (wrong #), there is no change in the table.
- If the WHERE is not written, all rows of the table will be updated

The UPDATE Statement - Part I      The UPDATE Statement - Part II

```
1 • SELECT
2 *
3 FROM
4 departments_dup
5 ORDER BY dept_no;
6
7
8 • COMMIT;
```

Result Grid | Filter Rows: [ ]

| dept_no | dept_name          |
|---------|--------------------|
| d001    | Marketing          |
| d002    | Finance            |
| d003    | Human Resources    |
| d004    | Production         |
| d005    | Development        |
| d006    | Quality Management |
| d007    | Sales              |
| d008    | Research           |
| d009    | Customer Service   |

Commit

The line of COMMIT is the same as toggling off the commit button

Example: without WHERE clause

```
1 • SELECT
2 *
3 FROM
4 departments_dup
5 ORDER BY dept_no;
6
7
8 • COMMIT;
9
10
11 • UPDATE departments_dup
12 SET
13 dept_no = 'd011',
14 dept_name = 'Quality Control';
```

The UPDATE will change all the table into the same data.

Result Grid | Filter Rows:

| dept_no | dept_name       |
|---------|-----------------|
| d011    | Quality Control |

departments\_dup 7 ✓

[USING ROLLBACK TO CANCEL THE PREVIOUS STEP]

```

4 departments_dup
5 ORDER BY dept_no;
6
7
8 • COMMIT;
9
10
11 • UPDATE departments_dup
12 SET
13 dept_no = 'd011',
14 dept_name = 'Quality Control';
15
16 • ROLLBACK;
17
18 • COMMIT;
19
20
21
22
23

```

**Exercise:** Change the “Business Analysis” department name to “Data Analysis”.

**Hint:** To solve this exercise, use the “departments” table.

```

--
37 • UPDATE departments
38 SET dept_name = "Data Analysis"
39 WHERE dept_no = "d010"
40 ;
41

```

## Section 11: DELETE statement

56. The DELETE statement

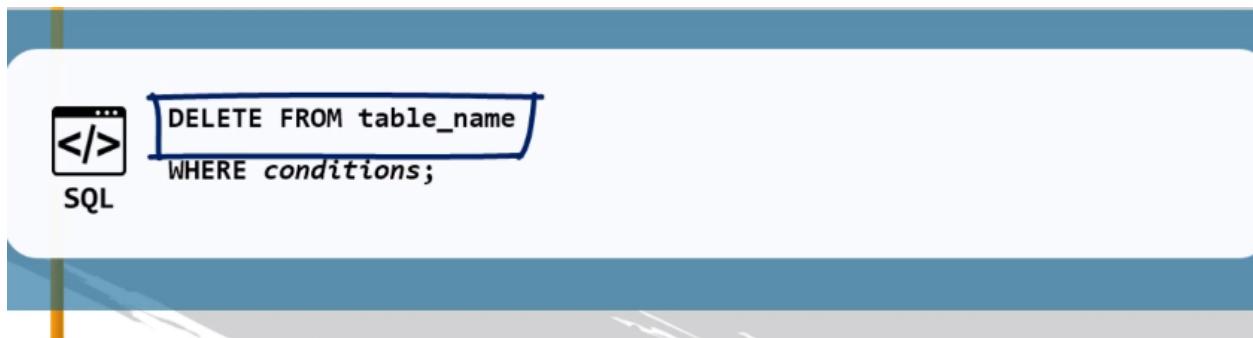
Part 1:

Example: Jonathan

```
37 • USE employees;
38 • COMMIT;
39
40 • SELECT *
41 FROM titles
42 WHERE emp_no = 999903;
43
44
```

| emp_no | title           | from_date  | to_date |
|--------|-----------------|------------|---------|
| 999903 | Senior Engineer | 1997-10-01 | NULL    |

## DELETE CLAUSE



```
DELETE FROM employees
WHERE emp_no = 999903;
```

After deletion

A screenshot of MySQL Workbench showing the 'titles' table. The table has columns: emp\_no, birth\_date, first\_name, last\_name, gender, and hire\_date. All rows are empty (NULL).

|   | emp_no | birth_date | first_name | last_name | gender | hire_date |
|---|--------|------------|------------|-----------|--------|-----------|
| 1 | NULL   | NULL       | NULL       | NULL      | NULL   | NULL      |

Also, the record on the titles table is deleted because it says “ON DELETE CASCADE”

#### DDL for employees.titles

```

1 CREATE TABLE `titles` (
2 `emp_no` int NOT NULL,
3 `title` varchar(50) NOT NULL,
4 `from_date` date NOT NULL,
5 `to_date` date DEFAULT NULL,
6 PRIMARY KEY (`emp_no`,`title`,`from_date`),
7 CONSTRAINT `titles_ibfk_1` FOREIGN KEY (`emp_no`) REFERENCES `employees` (`emp_no`) ON DELETE CASCADE
8) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci

```

#### Foreign constraint - ON DELETE CASCADE

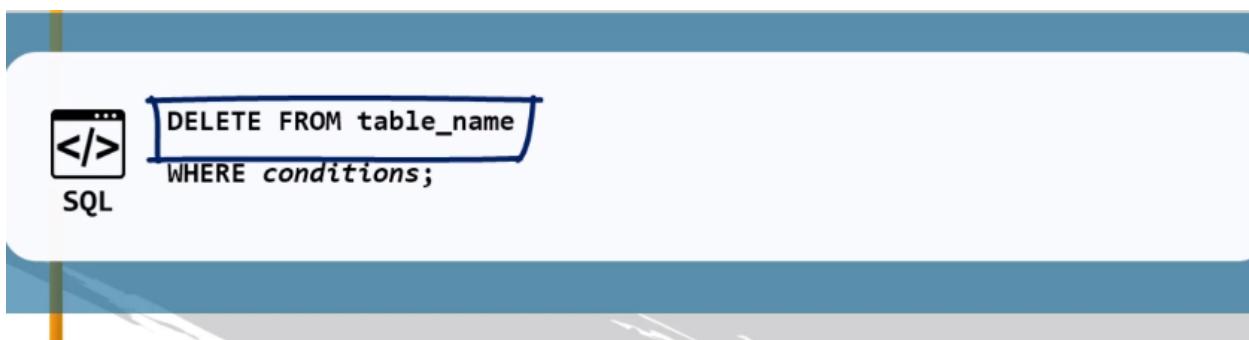
If a specific value from the parent table’s primary key has been deleted, all the records from the child table referring to this value will be removed as well

[If one item is deleted]

ROLLBACK;

#### Part 2:

WHERE is very important in the DELETE statement



Example: without WHERE everything on this table is deleted

The screenshot shows the MySQL Workbench interface. In the top tab bar, there are two tabs: "The DELETE Statement - Part II" and "The DELETE Statement - Part I". The main area displays a SQL editor with the following code:

```
1 • SELECT *
2 *
3 FROM
4 departments_dup
5 ORDER BY dept_no;|
6
7 • DELETE FROM departments_dup;
```

Below the editor is a "Result Grid" window. It has a header row with columns "dept\_no" and "dept\_name". The body of the grid is currently empty, indicated by a large blue bracket on the left side.

**Exercise: Remove the department number 10 record from the “departments” table.**

```
38 • SELECT *
39 FROM departments;
40
41 • DELETE FROM departments
42 WHERE dept_no = "d010";
```

## 57. DROP vs. TRUNCATE vs. DELETE

### (1) DROP: too powerful

- You will lose every detail in the record, including index, constraints, and the whole tables
- You won't be able to roll back to its initial state, or to the last COMMIT statement
- DROP it only when you are sure you aren't going to use the table in question any more

### (2) TRUNCATE

DELETE without WHERE; when truncating, auto-increment values will be reset

## TRUNCATE

when truncating, auto-increment values will be reset

## AUTO\_INCREMENT

| column_1 |
|----------|
| 1        |
| 2        |
| 3        |
| 4        |
| ...      |
| 10       |

TRUNCATE

| column_1 |
|----------|
| X 1      |
| 12       |
|          |
|          |
|          |

## (3) DELETE

Removes records row by row



`DELETE FROM table_name  
WHERE conditions;`

## (4) TRUNCATE vs. DELETE without WHERE

The SQL optimizer will implement different programmatic approaches when we are using TRUNCATE or DELETE

- (i) TRUNCATE (whole table) delivers the output much quicker than DELETE (row by row)
- (ii) auto-increment values are not reset with DELETE

- auto-increment values are *not* reset with `DELETE`

## AUTO\_INCREMENT

| column_1 |
|----------|
| 1        |
| 2        |
| 3        |
| 4        |
| ...      |
| 10       |



| column_1 |
|----------|
| 11       |
| 12       |
|          |
|          |
|          |
|          |

### Exercise:

Question 1:

You want to remove a table from your database, together with its structure and all related objects, like indexes and constraints. Which of the following commands would allow you to do that?

DELETE without a WHERE clause

DELETE with a WHERE clause

DROP

TRUNCATE

Question 2:

You want to remove the entire data from a table and keep its structure. Furthermore, you want auto-increment values to be reset. Which of the following commands should you use?

DELETE without a WHERE clause

DELETE with a WHERE clause

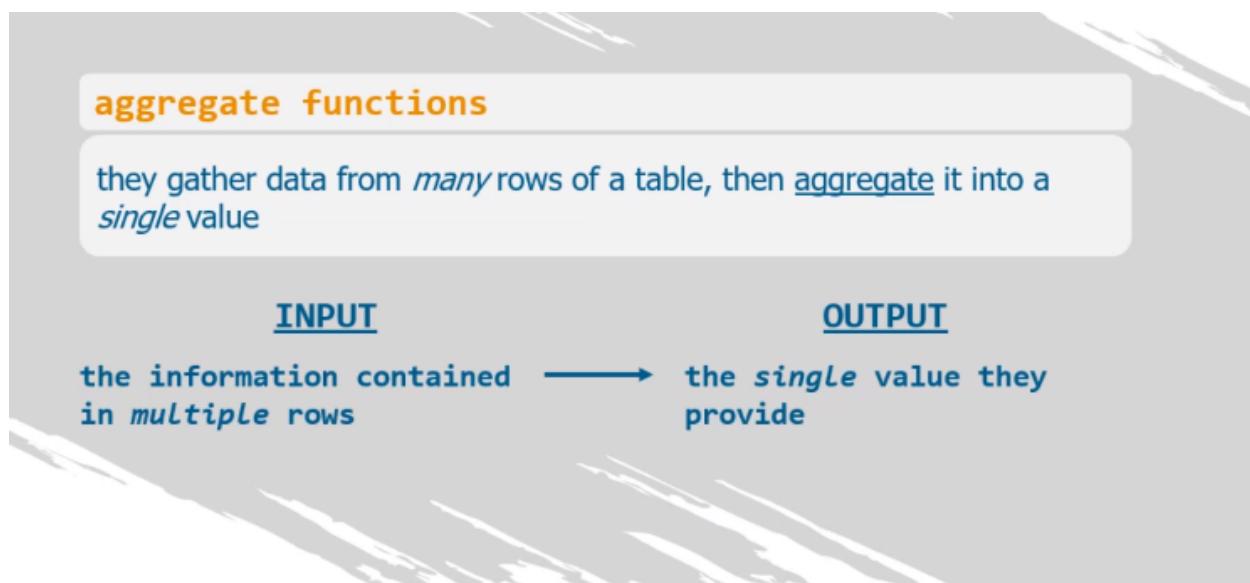
DROP

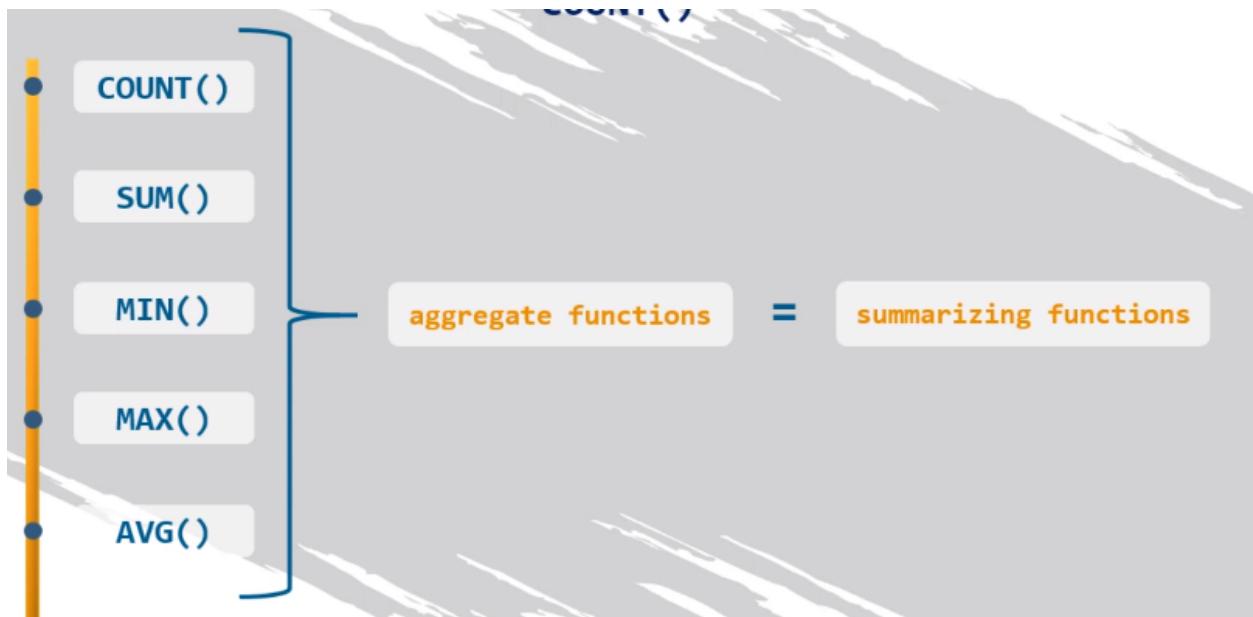
TRUNCATE

## Section 12: More aggregate functions

### 58. COUNT()

Review:





Why do these functions exist?

- They are a response to the information requirements of a company's different organizational levels
- Top management executives are typically interested in summarized figures and rarely in detailed data

COUNT()

- Applicable to both numeric and non-numeric data

Example: how many employee start dates are in the database?

```

38 • SELECT *
39 FROM salaries
40 ORDER BY salary DESC
41 LIMIT 10;
42
43 • SELECT
44 COUNT(salary)
45 FROM salaries;
```

If COUNT(salary) is changed into COUNT(from\_date), the result is the same. This is because the link to the columns are the same and there is no missing data.

## COUNT(DISTINCT)

- Helps us find the number of times unique values are encountered in a given column

```
39 • SELECT
40 COUNT(DISTINCT from_date)
41 FROM salaries;
42 |
43
```

Result Grid | Filter Rows: Export:

|                           |
|---------------------------|
| COUNT(DISTINCT from_date) |
| 6392                      |

There is no from\_date as repetition.

Usually COUNT ignores any NULL values, but COUNT(\*) returns the number of all rows of the table.

Also, there is no space before the parenthesis.



**Exercise: How many departments are there in the “employees” database? Use the ‘dept\_emp’ table to answer the question.**

```
39 • SELECT
40 COUNT(DISTINCT emp_no)
41 FROM dept_emp;
42
43
```

Result Grid | Filter Rows:

|                        |
|------------------------|
| COUNT(DISTINCT emp_no) |
| 300025                 |

## 59. SUM()

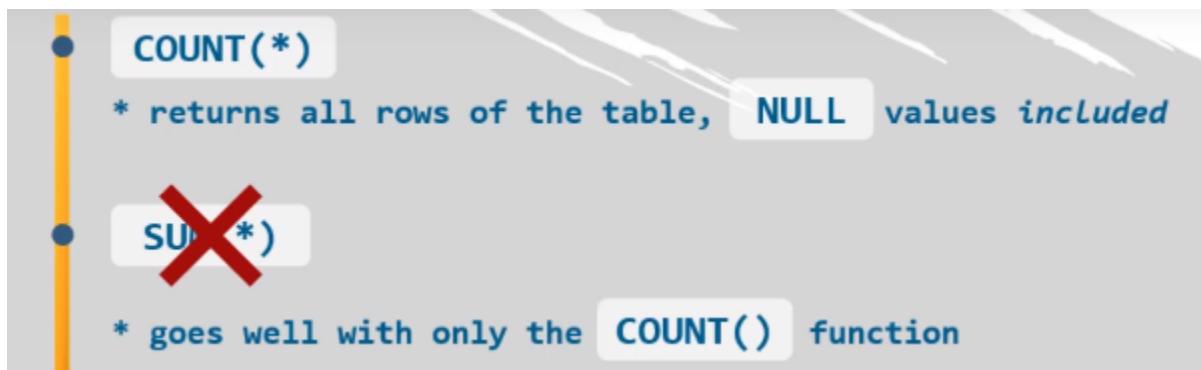
Example:

```
39 • SELECT SUM(salary)
40 FROM salaries;
41
```

Result Grid | Filter Rows: [ ]

| SUM(salary) |
|-------------|
| 61678125784 |

- SUM() does not allow \*



Exercise: What is the total amount of money spent on salaries for all contracts starting after the 1st of January 1997?

```
39 • SELECT SUM(salary)
40 FROM salaries
41 WHERE from_date > "1997-01-01";
42
```

Result Grid | Filter Rows: [ ] | Export: [ ]

| SUM(salary) |
|-------------|
| 31909143195 |

## 60. MAX() and MIN()

Example: which is the highest salary we offer?

```
39 • SELECT MAX(salary)
40 FROM salaries
41 ;
42
```

| Result Grid |             |
|-------------|-------------|
|             | MAX(salary) |
| ▶           | 158220      |

```
39 • SELECT MIN(salary)
40 FROM salaries
41 ;
42
```

| Result Grid |             |
|-------------|-------------|
|             | MIN(salary) |
| ▶           | 38735       |

**Exercise:**

1. Which is the lowest employee number in the database?
2. Which is the highest employee number in the database?

```
SELECT MIN(emp_no)
FROM salaries
;
```

```
SELECT MAX(emp_no)
FROM salaries
;
```

61. AVG()

Extract all the average values of all non-NULL values in a table

Exercise: which is the average annual salary the company's employee received?

```
39 • SELECT AVG(salary)
40 FROM salaries
41 ;
42 |
```

< Result Grid Filter Rows: [ ]

|   | AVG(salary) |
|---|-------------|
| ▶ | 63761.2043  |

**Exercise: What is the average annual salary paid to employees who started after the 1st of January 1997?**

```
38
39 • SELECT AVG(salary)
40 FROM salaries
41 WHERE from_date > "1997-01-01"
42 ;
```

< Result Grid Filter Rows: [ ]

|   | AVG(salary) |
|---|-------------|
| ▶ | 67717.7450  |

## 62. ROUND()

- Usually applied to a single value an aggregate function returns

ROUND(#, decimal\_places)

```

39 • SELECT ROUND(AVG(salary),2)
40 FROM salaries
41 ;
42

```

| Result Grid          |              |
|----------------------|--------------|
|                      | Filter Rows: |
| ROUND(AVG(salary),2) |              |
| 63761.20             |              |

**Exercise: Round the average amount of money spent on salaries for all contracts that started after the 1st of January 1997 to a precision of cents.**

```

39 • SELECT ROUND(AVG(salary),2)
40 FROM salaries
41 WHERE from_date > "1997-01-01"
42 ;
43

```

| Result Grid          |              |
|----------------------|--------------|
|                      | Filter Rows: |
| ROUND(AVG(salary),2) |              |
| 67717.75             |              |

### 63. IFNULL() and COALESCE

Using some strings to replace NULL values

| Result Grid |                    |              |
|-------------|--------------------|--------------|
| dept_no     | dept_name          | dept_manager |
| d001        | Marketing          | NULL         |
| d002        | Finance            | NULL         |
| d003        | Human Resources    | NULL         |
| d004        | Production         | NULL         |
| d005        | Development        | NULL         |
| d006        | Quality Management | NULL         |
| d007        | Sales              | NULL         |
| d008        | Research           | NULL         |
| d009        | Customer Service   | NULL         |
| d010        | NULL               | NULL         |
| d011        | NULL               | NULL         |

"department name not provided"

### IFNULL(exp1,exp2)

It returns the first of the two indicated values if the data values found in the table is not NULL, and returns the second value if there is a NULL value

- Prints the returned value in the column of the output

```

11
12 • SELECT
13 dept_no,
14 IFNULL(dept_name,
15 'Department name not provided')
16 FROM
17 departments_dup;]
18

```

| Result Grid |                                                      |
|-------------|------------------------------------------------------|
| dept_no     | IFNULL(dept_name,<br>'Department name not provided') |
| d009        | Customer Service                                     |
| d005        | Development                                          |
| d002        | Finance                                              |
| d003        | Human Resources                                      |
| d001        | Marketing                                            |
| d004        | Production                                           |
| d006        | Quality Management                                   |
| d008        | Research                                             |
| d007        | Sales                                                |
| d010        | Department name not provided                         |
| d011        | Department name not provided                         |

One more step is to use alias to replace the long query

### COALESCE (ex1,exp2,...)

- Functions like IFNULL with more than one parameters
- COALESCE will always return a single value of the ones we have within parentheses, and this value **will be the first non-NULL value of this list**, reading the values from left to right.

Two args

```

11
12 • SELECT
13 dept_no,
14 COALESCE(dept_name,
15 'Department name not provided') as dept_name
16 FROM
17 departments_dup; I
18

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| dept_no | dept_name                    |
|---------|------------------------------|
| d009    | Customer Service             |
| d005    | Development                  |
| d002    | Finance                      |
| d003    | Human Resources              |
| d001    | Marketing                    |
| d004    | Production                   |
| d006    | Quality Management           |
| d008    | Research                     |
| d007    | Sales                        |
| d010    | Department name not provided |
| d011    | Department name not provided |

Result 27 ×

Three args:

```

IFNULL() and COALESCE() * x

```

31
32 • SELECT
33 dept\_no,
34 dept\_name,
35 COALESCE(dept\_manager, dept\_name, 'N/A') AS dept\_manager
36 FROM
37 departments\_dup
38 ORDER BY dept\_no ASC; I

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| dept_no | dept_name          | dept_manager       |
|---------|--------------------|--------------------|
| d001    | Marketing          | Marketing          |
| d002    | Finance            | Finance            |
| d003    | Human Resources    | Human Resources    |
| d004    | Production         | Production         |
| d005    | Development        | Development        |
| d006    | Quality Management | Quality Management |
| d007    | Sales              | Sales              |
| d008    | Research           | Research           |
| d009    | Customer Service   | Customer Service   |
| d010    | NULL               | N/A                |
| d011    | NULL               | N/A                |

The AS dept\_manager part means that the NON-NULL information is displayed in the column of dept\_manager

IFNULL() and COALESCE() do not make any changes to the dataset. They merely create an output where certain data values appear in place of NULL values.

**Another way to use COALESCE(): create a supplementary column; it helps visualization of a fake column**

```
33 • SELECT
34 dept_no, dept_name,
35 COALESCE("Department name not provided") AS fake_col
36 FROM departments_dup;
37
38 #SELECT * FROM departments_dup;
39
```

The screenshot shows a database query result grid. At the top, there are buttons for 'Result Grid', 'Filter Rows:', 'Export:', and 'Wrap Cell Content:'. The result grid has three columns: 'dept\_no', 'dept\_name', and 'fake\_col'. The data rows are:

|   | dept_no | dept_name          | fake_col                     |
|---|---------|--------------------|------------------------------|
| ▶ | d010    | Business Analytics | Department name not provided |
|   | d009    | Customer Service   | Department name not provided |
|   | d005    | Development        | Department name not provided |
|   | d002    | Finance            | Department name not provided |
|   | d003    | Human Resources    | Department name not provided |
|   | d001    | Marketing          | Department name not provided |

**Exercise1:** Select the department number and name from the 'departments\_dup' table and add a third column where you name the department number ('dept\_no') as 'dept\_info'. If 'dept\_no' does not have a value, use 'dept\_name'.

```
33 • SELECT
34 dept_no, dept_name,
35 IFNULL(dept_no,dept_name) AS dept_inf
36 FROM departments_dup
37 ORDER BY dept_no ASC;
```

The screenshot shows a database query result grid. At the top, there are buttons for 'Result Grid', 'Filter Rows:', 'Export:', and 'Wrap Cell Content:'. The result grid has three columns: 'dept\_no', 'dept\_name', and 'dept\_info'. The data rows are:

|   | dept_no | dept_name          | dept_info |
|---|---------|--------------------|-----------|
| ▶ | d001    | Marketing          | d001      |
|   | d002    | Finance            | d002      |
|   | d003    | Human Resources    | d003      |
|   | d004    | Production         | d004      |
|   | d005    | Development        | d005      |
|   | d006    | Quality Management | d006      |

**Exercise 2:** Modify the code obtained from the previous exercise in the following way. Apply the IFNULL() function to the values from the first and second column, so that 'N/A' is displayed whenever a department number has no value, and 'Department name not provided' is shown if there is no value for 'dept\_name'.

**SELECT**

```
IFNULL(dept_no, 'N/A') as dept_no,
IFNULL(dept_name,
 'Department name not provided') AS dept_name,
COALESCE(dept_no, dept_name) AS dept_info
```

**FROM**

**departments\_dup**

**ORDER BY dept\_no ASC;**

### **Exercise3:**

Question 1:

The output of which of the following queries will be the string 'Second'?

SELECT COALESCE(NULL, NULL, 'Third') AS coalesce\_test;

SELECT COALESCE('First', null, 'Third') AS coalesce\_test;

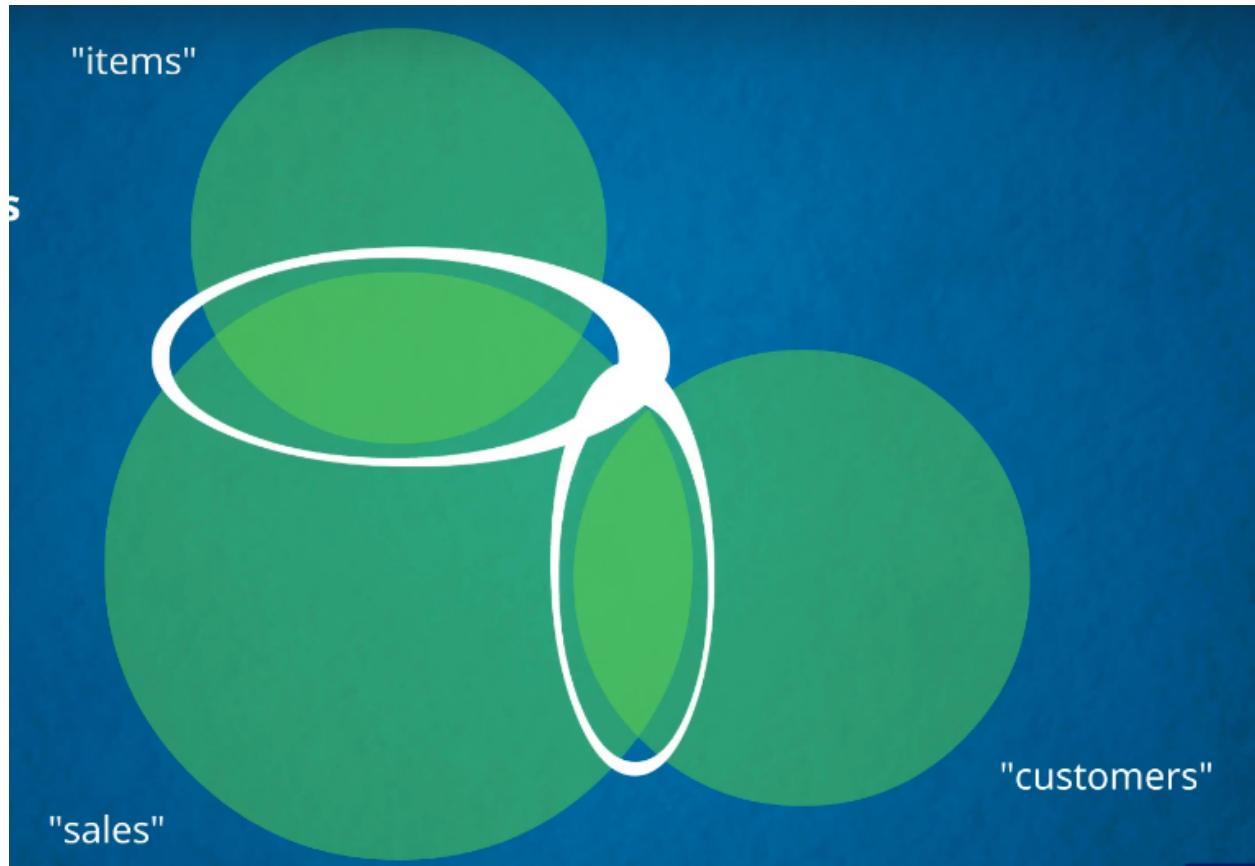
SELECT COALESCE('First', 'Second', 'Third') AS coalesce\_test;

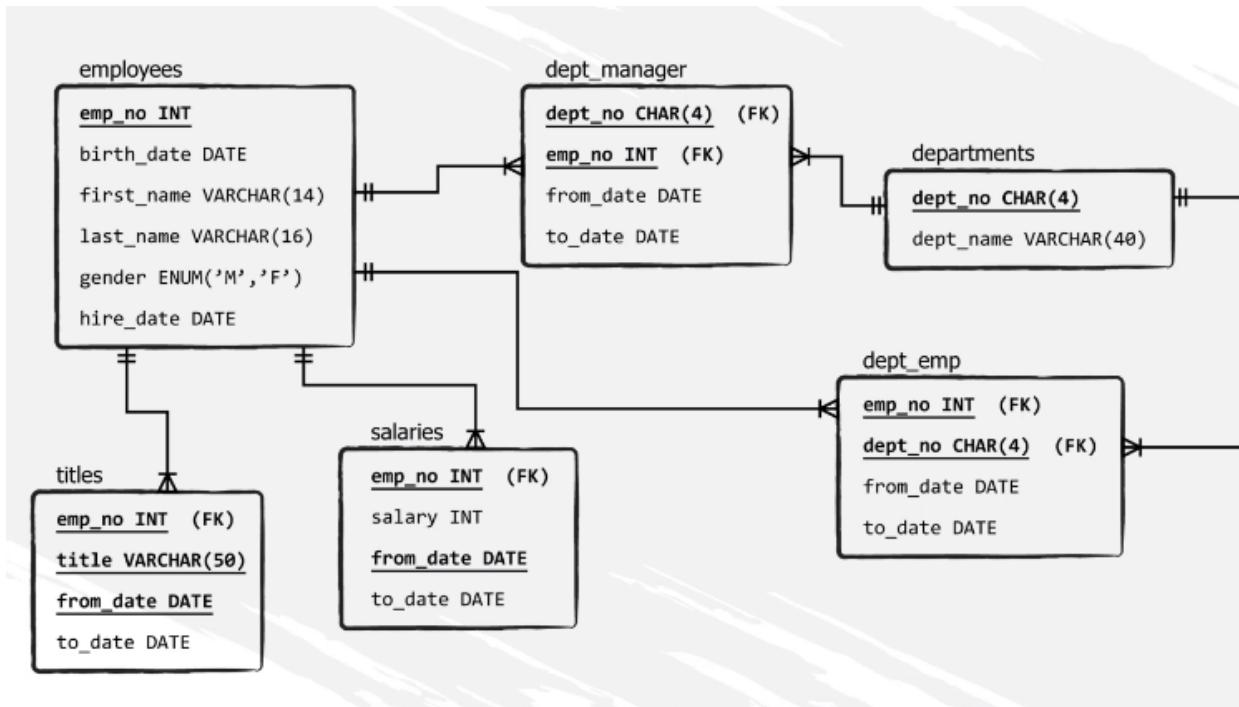
SELECT COALESCE(NULL, 'Second', null) AS coalesce\_test;

## Section 13: SQL joins

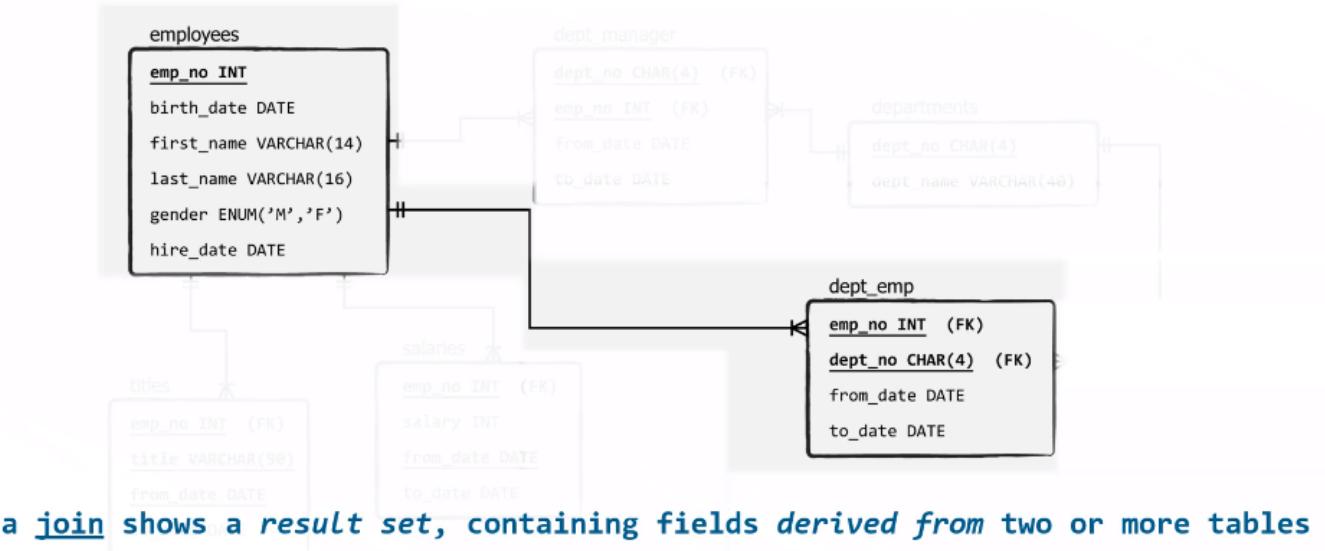
### 64. Introduction to Joins

Relational database: sharing columns in tables can speed up retrieving data



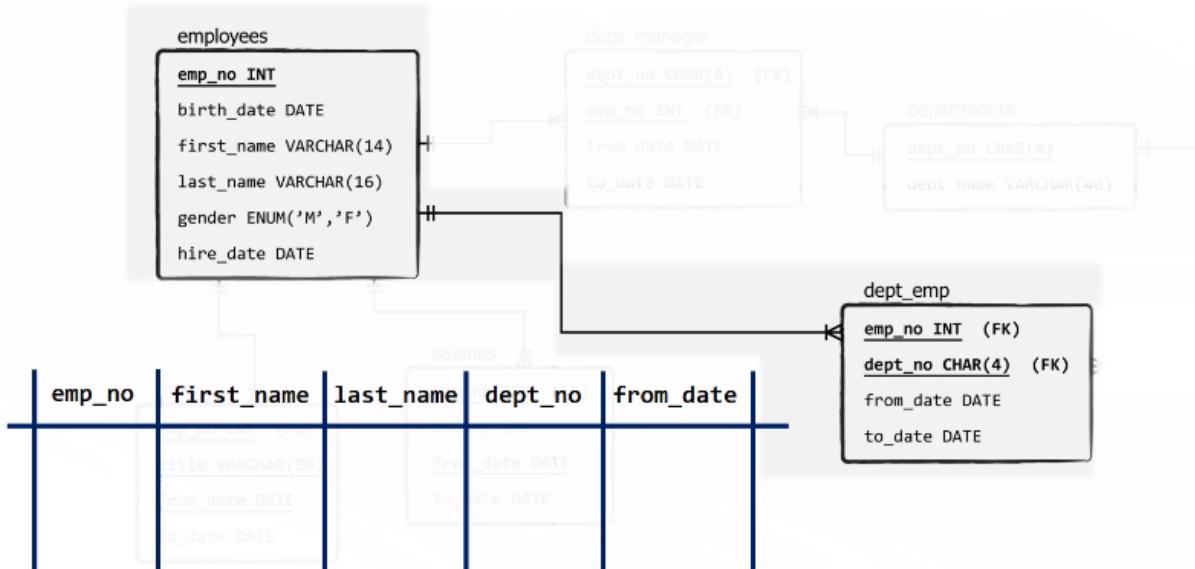


Relational schemas are the tools that will help you find a strategy for linking tables

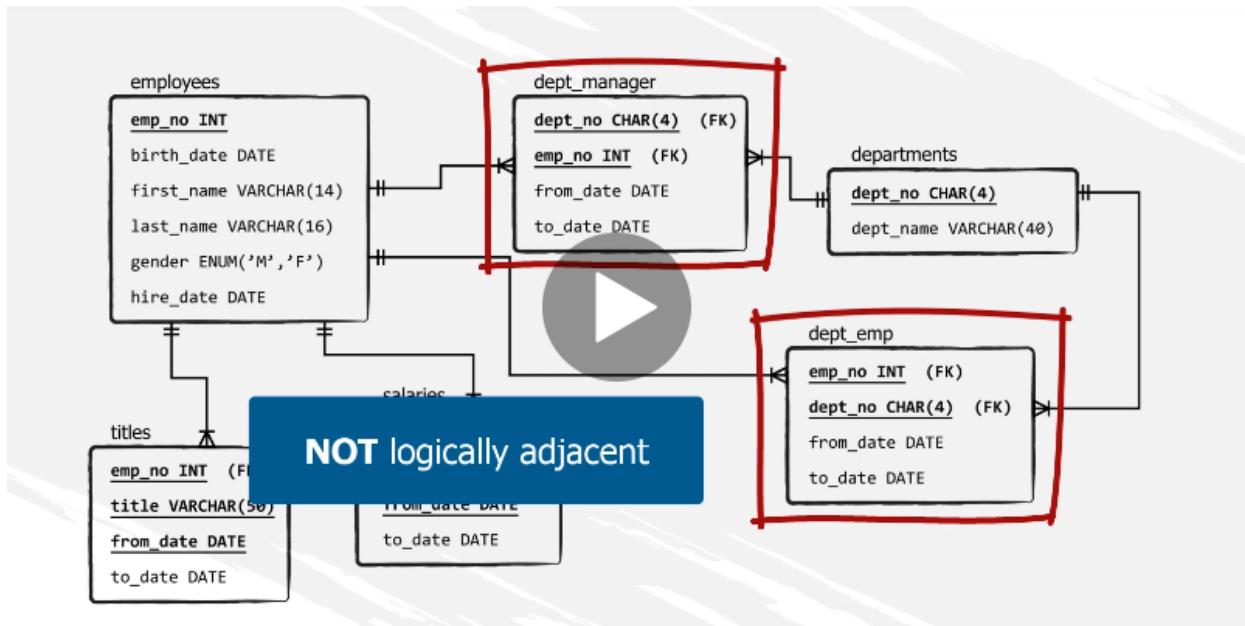


### Joins:

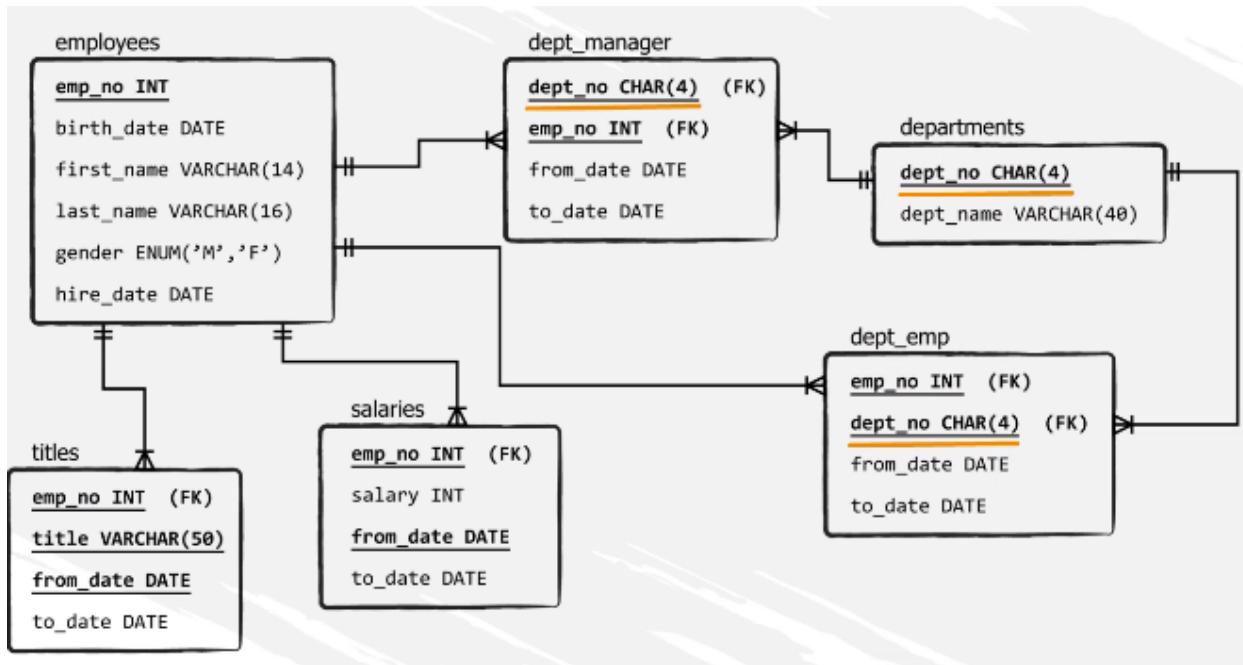
- We must find the shared column from the table, e.g., employee\_no
- We will be free to add columns to the outputs



- The column you use to relate tables must represent the same object, such as `id`
- But the tables you are considering need not be logically adjacent



By logically adjacent, it means that two tables are connected by a line. However, these two tables can be joined by `emp_no`.



**Exercise1:** If you currently have the ‘departments\_dup’ table set up, use **DROP COLUMN** to remove the ‘dept\_manager’ column from the ‘departments\_dup’ table.

Then, use **CHANGE COLUMN** to change the ‘dept\_no’ and ‘dept\_name’ columns to **NULL**.

(If you don’t currently have the ‘departments\_dup’ table set up, create it. Let it contain two columns: dept\_no and dept\_name. Let the data type of dept\_no be CHAR of 4, and the data type of dept\_name be VARCHAR of 40. Both columns are allowed to have null values. Finally, insert the information contained in ‘departments’ into ‘departments\_dup’.)

Then, insert a record whose department name is “Public Relations”.

Delete the record(s) related to department number two.

Insert two new records in the “departments\_dup” table. Let their values in the “dept\_no” column be “d010” and “d011”.

```
if you currently have 'departments_dup' set up:
```

```
ALTER TABLE departments_dup
```

```
DROP COLUMN dept_manager;
```

```
ALTER TABLE departments_dup
CHANGE COLUMN dept_no dept_no CHAR(4) NULL;
```

```
ALTER TABLE departments_dup
CHANGE COLUMN dept_name dept_name VARCHAR(40) NULL;
```

**Exercise2:**

**Create and fill in the ‘dept\_manager\_dup’ table, using the following code:**

```
DROP TABLE IF EXISTS dept_manager_dup;

CREATE TABLE dept_manager_dup (
 emp_no int(11) NOT NULL,
 dept_no char(4) NULL,
 from_date date NOT NULL,
 to_date date NULL
);
```

```
INSERT INTO dept_manager_dup
select * from dept_manager;
```

```
INSERT INTO dept_manager_dup (emp_no, from_date)
VALUES (999904, '2017-01-01'),
 (999905, '2017-01-01'),
```

```
(999906, '2017-01-01'),
```

```
(999907, '2017-01-01');
```

```
DELETE FROM dept_manager_dup
```

```
WHERE
```

```
dept_no = 'd001';
```

```
INSERT INTO departments_dup (dept_name)
```

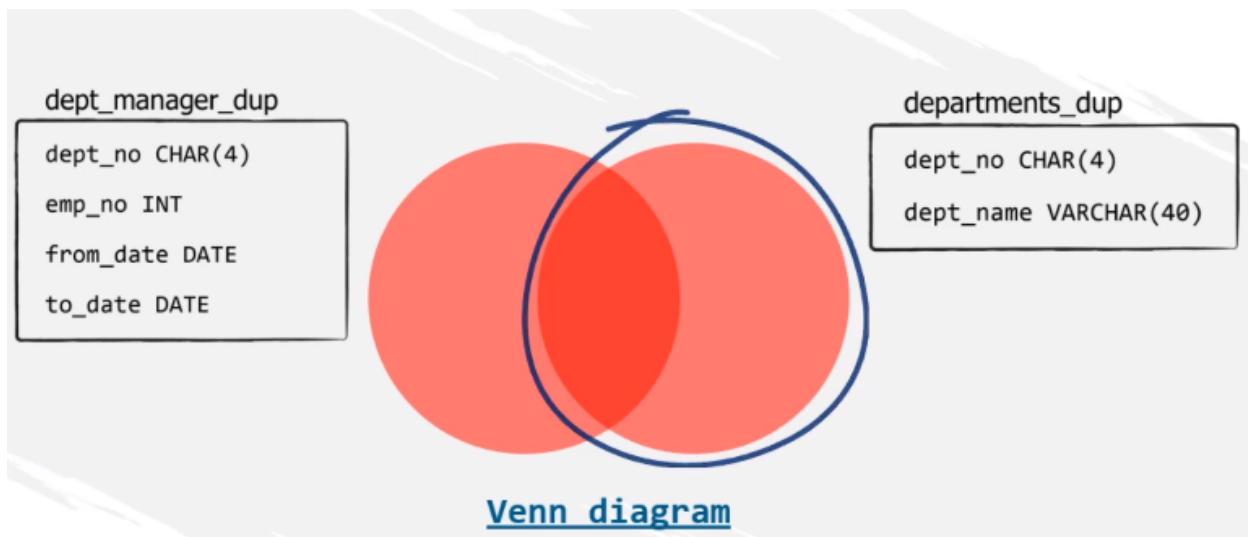
```
VALUES ('Public Relations');
```

```
DELETE FROM departments_dup
```

```
WHERE
```

```
dept_no = 'd002';
```

## 65. Inner JOIN



The related column is dept\_no (red). The red part is called matching values or matching columns (or linking column).

Data observations:

After checking the dept\_manager\_dup, we can see (1) there are some NULL cells, and (2) there is no d001.

```
32 • SELECT *
33 FROM dept_manager_dup
34 ORDER BY dept_no;
```

The screenshot shows a database query results grid titled 'dept\_manager\_dup 2'. The grid has columns: emp\_no, dept\_no, from\_date, and to\_date. The data includes several rows with dept\_no values like 'NULL', 'd002', and 'd003', indicating missing or incorrect department numbers. The grid interface includes a 'Result Grid' tab, a 'Filter Rows:' input field, and an 'Exp' button.

|   | emp_no | dept_no | from_date  | to_date    |
|---|--------|---------|------------|------------|
| ▶ | 999904 | NULL    | 2017-01-01 | NULL       |
|   | 999905 | NULL    | 2017-01-01 | NULL       |
|   | 999906 | NULL    | 2017-01-01 | NULL       |
|   | 999907 | NULL    | 2017-01-01 | NULL       |
|   | 110085 | d002    | 1985-01-01 | 1989-12-17 |
|   | 110114 | d002    | 1989-12-17 | 9999-01-01 |
|   | 110183 | d003    | 1985-01-01 | 1992-03-21 |

After checking the departments\_dup, we can see (1) no dept\_no for public\_relations, and (2) d010 and 011 has NULL

```
7
8
9 # departments_dup
10 • SELECT
11 *
12 FROM
13 departments_dup
14 ORDER BY dept_no;
```

The screenshot shows a database query results grid titled 'departments\_dup'. The grid has columns: dept\_no and dept\_name. It lists various department names with their corresponding dept\_no. There are two entries with NULL dept\_no values: one for 'Public Relations' and another for 'Customer Service'. The grid interface includes a 'Result Grid' tab, a 'Filter Rows:' input field, and an 'Exp' button.

|  | dept_no | dept_name          |
|--|---------|--------------------|
|  | NULL    | Public Relations   |
|  | d001    | Marketing          |
|  | d003    | Human Resources    |
|  | d004    | Production         |
|  | d005    | Development        |
|  | d006    | Quality Management |
|  | d007    | Sales              |
|  | d008    | Research           |
|  | d009    | Customer Service   |
|  | d010    | NULL               |
|  | d011    | NULL               |

Inner JOIN syntax:

**SQL**

```

</>
SELECT
 table_1.column_name(s), table_2.column_name(s)
FROM
 table_1
JOIN
 table_2 ON table_1.column_name = table_2.column_name;

```

Another version using alias for table\_n (no “table\_1 AS t1”)

**SQL**

```

</>
SELECT
 t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
 table_1 t1
JOIN
 table_2 t2 ON t1.column_name = t2.column_name;

```

INNER JOIN - Part II

## INNER JOIN

**M**

```

dept_manager_dup
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE

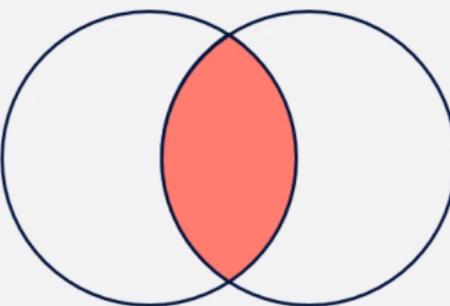
```

**D**

```

departments_dup
dept_no CHAR(4)
dept_name VARCHAR(40)

```



Code:

```
6
7 • SELECT
8 m.dept_no, m.emp_no, d.dept_name
9 FROM
10 dept_manager_dup m
11 INNER JOIN
12 departments_dup d ON m.dept_no = d.dept_no
13 ORDER BY m.dept_no;
```

The screenshot shows a database query results grid. At the top, there are buttons for 'Result Grid' (selected), 'Filter Rows', 'Export' (with icons for CSV, XML, and PDF), and 'Wrap Cell'. The grid itself has three columns: 'dept\_no', 'emp\_no', and 'dept\_name'. The data is as follows:

| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| d003    | 110183 | Human Resources    |
| d003    | 110228 | Human Resources    |
| d004    | 110303 | Production         |
| d004    | 110344 | Production         |
| d004    | 110386 | Production         |
| d004    | 110420 | Production         |
| d005    | 110511 | Development        |
| d005    | 110567 | Development        |
| d006    | 110725 | Quality Management |
| d006    | 110765 | Quality Management |
| d006    | 110800 | Quality Management |
| d006    | 110854 | Quality Management |
| d007    | 111035 | Sales              |
| d007    | 111133 | Sales              |
| d008    | 111400 | Research           |

Result Grid | Filter Rows:  Export: Wrap Cell Content:

| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| d003    | 10228  | Human Resources    |
| d003    | 10183  | Human Resources    |
| d004    | 10344  | Production         |
| d004    | 10420  | Production         |
| d004    | 10303  | Production         |
| d004    | 10386  | Production         |
| d005    | 10567  | Development        |
| d005    | 10511  | Development        |
| d006    | 10800  | Quality Management |
| d006    | 10765  | Quality Management |
| d006    | 10854  | Quality Management |
| d006    | 10725  | Quality Management |
| d007    | 11035  | Sales              |
| d007    | 11133  | Sales              |
| d008    | 11140  | Research           |
| d008    | 111534 | Research           |
| d009    | 111784 | Customer Service   |
| d009    | 111939 | Customer Service   |
| d009    | 111692 | Customer Service   |
| d009    | 111877 | Customer Service   |

Result 4 / X

Output:  Action Output

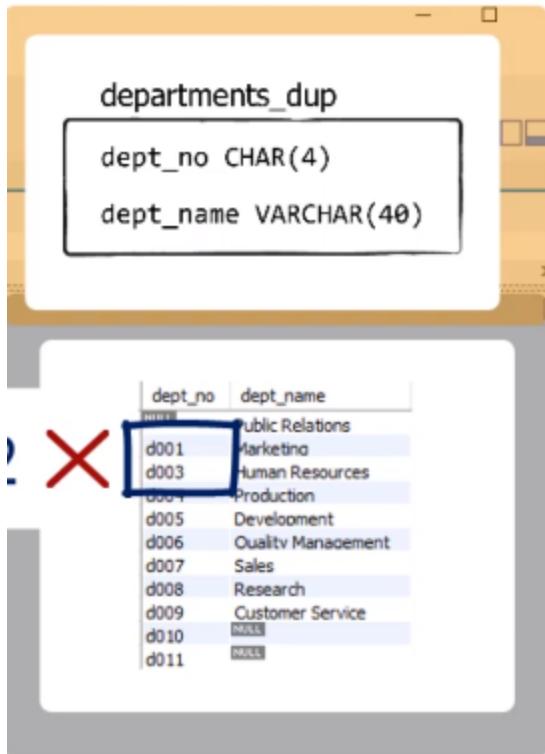
**INNER JOINS extract only records in which the values in the related columns match; if values not in the intersection, they are not displayed.**

**dept\_manager\_dup**

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```

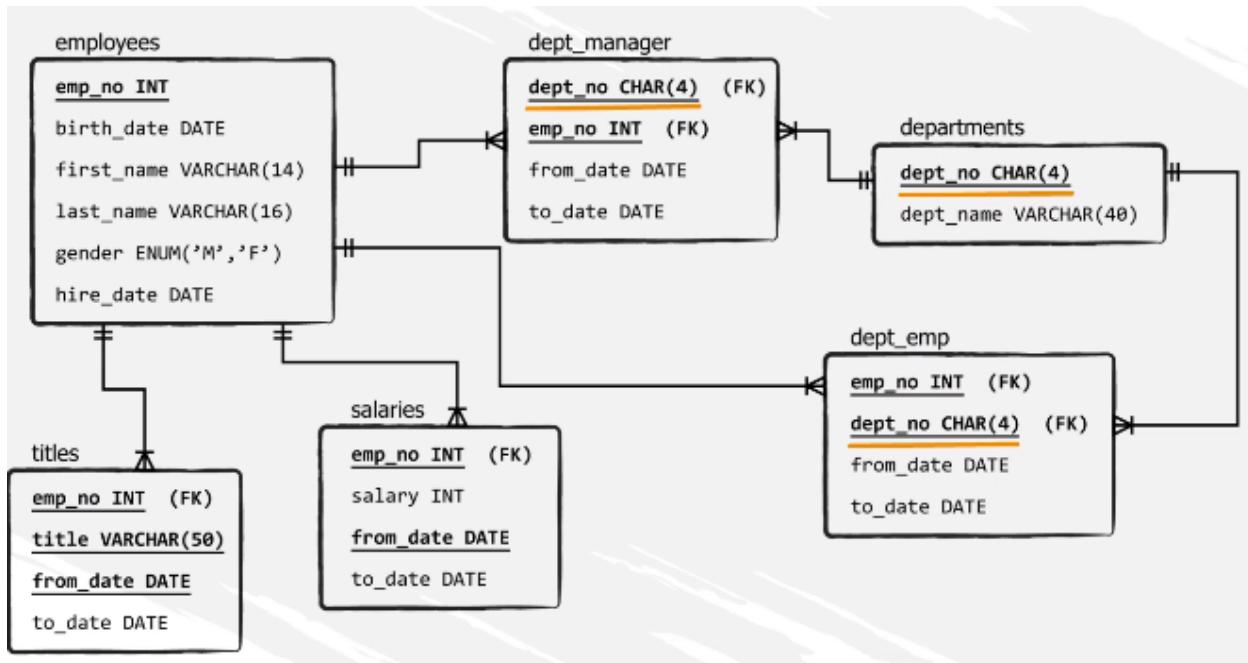
| emp_no | dept_no | from_date  | to_date    |
|--------|---------|------------|------------|
| 999904 | NULL    | 2017-01-01 | NULL       |
| 999905 | NULL    | 2017-01-01 | NULL       |
| 999906 | NULL    | 2017-01-01 | NULL       |
| 999907 | NULL    | 2017-01-01 | NULL       |
| 110085 | d002    | 1985-01-01 | 1989-12-17 |
| 110114 | d002    | 1985-01-01 | 1991-01-01 |
| 110183 | d003    | 1985-01-01 | 1985-03-01 |
| 110228 | d003    | 1992-03-21 | 9999-01-01 |
| 110303 | d004    | 1985-01-01 | 1988-09-09 |
| 110344 | d004    | 1988-09-09 | 1992-08-02 |
| 110386 | d004    | 1985-01-01 | 1991-08-30 |
| 110420 | d004    | 1985-01-01 | 9999-01-01 |
| 110511 | d005    | 1985-01-01 | 1992-04-25 |
| 110567 | d005    | 1992-04-25 | 9999-01-01 |
| 110725 | d006    | 1985-01-01 | 1994-05-01 |
| 110765 | d006    | 1985-05-06 | 1991-09-17 |
| 110800 | d006    | 1991-09-12 | 1994-06-28 |
| 110854 | d006    | 1994-06-28 | 9999-01-01 |
| 111035 | d007    | 1985-01-01 | 1991-03-07 |
| 111133 | d007    | 1991-03-07 | 9999-01-01 |
| 111400 | d008    | 1985-01-01 | 1991-04-08 |
| 111534 | d008    | 1991-04-08 | 9999-01-01 |
| 111692 | d009    | 1985-01-01 | 1988-10-17 |
| 111784 | d009    | 1988-10-17 | 1992-09-08 |
| 111877 | d009    | 1992-09-08 | 1996-01-03 |
| 111939 | d009    | 1996-01-03 | 9999-01-01 |

do01  
do10  
do11



If there is matching value, then there is no result

**Exercise:** Extract a list containing information about all managers' employee number, first and last name, department number, and hire date.



# the key point is to find the shared tables with matching values

SELECT

e.emp\_no, e.first\_name, e.last\_name, dm.dept\_no, e.hire\_date

FROM

dept\_manager\_dup dm

INNER JOIN

employees e ON e.emp\_no = dm.emp\_no

;

## 66. Some notes about JOINS

- Add any combination of columns in the joined tables to display more info

```
7 • SELECT
8 m.dept_no, m.emp_no, m.from_date, m.to_date, d.dept_name
9 FROM
10 dept_manager_dup m
11 INNER JOIN
12 departments_dup d ON m.dept_no = d.dept_no
13 ORDER BY m.dept_no;
14
15 |
```

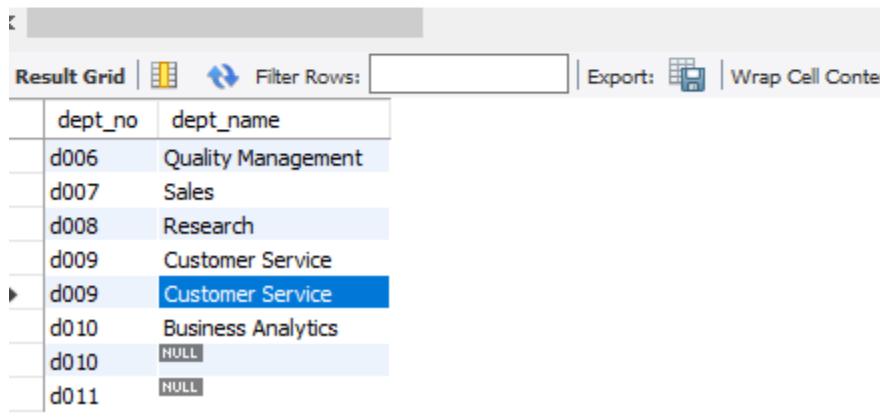
- Although you can enter FROM info about the table sharing, it is obligatory to follow the ordering like “SELECT... FROM... INNER JOIN... ON...”
- **INNER JOIN = JOIN here**
- **There is no difference between the equation after ON**
- **In ORDER BY, must specify the table of the column value like m or n**

## 67. Duplicate records

Duplicate records is also called duplicate rows. For a pair of duplicate records, the values in each column coincide.

- Duplication is always allowed in a database; they are sometimes encountered, especially in new, raw, or uncontrolled data

```
7 • INSERT INTO dept_manager_dup
8 VALUES ("110228","d003","1992-03-21","9999-01-01");
9
10 • INSERT INTO departments_dup
11 VALUES ("d009","Customer Service");
12
13 #check these two tables
14 • SELECT *
15 FROM dept_manager_dup
16 ORDER BY dept_no ASC;
17
18 • SELECT *
19 FROM departments_dup
20 ORDER BY dept_no ASC;
21
```



The screenshot shows a database result grid with the following columns: dept\_no and dept\_name. The data is as follows:

|   | dept_no | dept_name          |
|---|---------|--------------------|
|   | d006    | Quality Management |
|   | d007    | Sales              |
|   | d008    | Research           |
|   | d009    | Customer Service   |
| ▶ | d009    | Customer Service   |
|   | d010    | Business Analytics |
|   | d010    | NULL               |
|   | d011    | NULL               |

After applying these duplications, the JOIN results changed too (with duplicates):

```

22 • SELECT
23
24 m.dept_no, m.emp_no,d.dept_no
25
26 FROM
27
28 dept_manager_dup m
29
30 JOIN
31
32 departments_dup d ON m.dept_no = d.dept_no

```

Result Grid | Filter Rows:

|   | dept_no | emp_no | dept_no |
|---|---------|--------|---------|
|   | d003    | 110183 | d003    |
|   | d003    | 110228 | d003    |
|   | d004    | 110303 | d004    |
|   | d004    | 110344 | d004    |
|   | d004    | 110386 | d004    |
| ▶ | d004    | 110420 | d004    |
|   | d005    | 110511 | d005    |
|   | d005    | 110567 | d005    |
|   | d006    | 110725 | d006    |
|   | d006    | 110765 | d006    |
|   | d006    | 110800 | d006    |
|   | d006    | 110854 | d006    |
|   | d007    | 111035 | d007    |
|   | d007    | 111133 | d007    |
|   | d008    | 111400 | d008    |
|   | d008    | 111534 | d008    |
|   | d009    | 111692 | d009    |
|   | d009    | 111692 | d009    |
|   | d009    | 111784 | d009    |
|   | d009    | 111784 | d009    |
|   | d009    | 111877 | d009    |
|   | d009    | 111877 | d009    |
|   | d009    | 111939 | d009    |
|   | d009    | 111939 | d009    |
|   | d003    | 110228 | d003    |

|   |      |        |                  |
|---|------|--------|------------------|
| ▶ | d009 | 111692 | Customer Service |
|   | d009 | 111692 | Customer Service |

One way of not showing duplicates is to use GROUP BY (25 → 20)

```
18 • SELECT *
19 FROM departments_dup
20 ORDER BY dept_no ASC;
21
22 • SELECT
23
24 m.dept_no, m.emp_no,d.dept_name
25
26 FROM
27
28 dept_manager_dup m
29
30 JOIN
31
32 departments_dup d ON m.dept_no = d.dept_no
33 GROUP BY m.emp_no
34 ORDER BY dept_no
35 ;
36
```

Result Grid | Filter Rows: [ ] | Export: [ ] | Wrap Cell Content: [ ]

|   | dept_no | emp_no | dept_name          |
|---|---------|--------|--------------------|
| ▶ | d003    | 110183 | Human Resources    |
|   | d003    | 110228 | Human Resources    |
|   | d004    | 110303 | Production         |
|   | d004    | 110344 | Production         |
|   | d004    | 110386 | Production         |
|   | d004    | 110420 | Production         |
|   | d005    | 110511 | Development        |
|   | d005    | 110567 | Development        |
|   | d006    | 110725 | Quality Management |

Result 9 ×

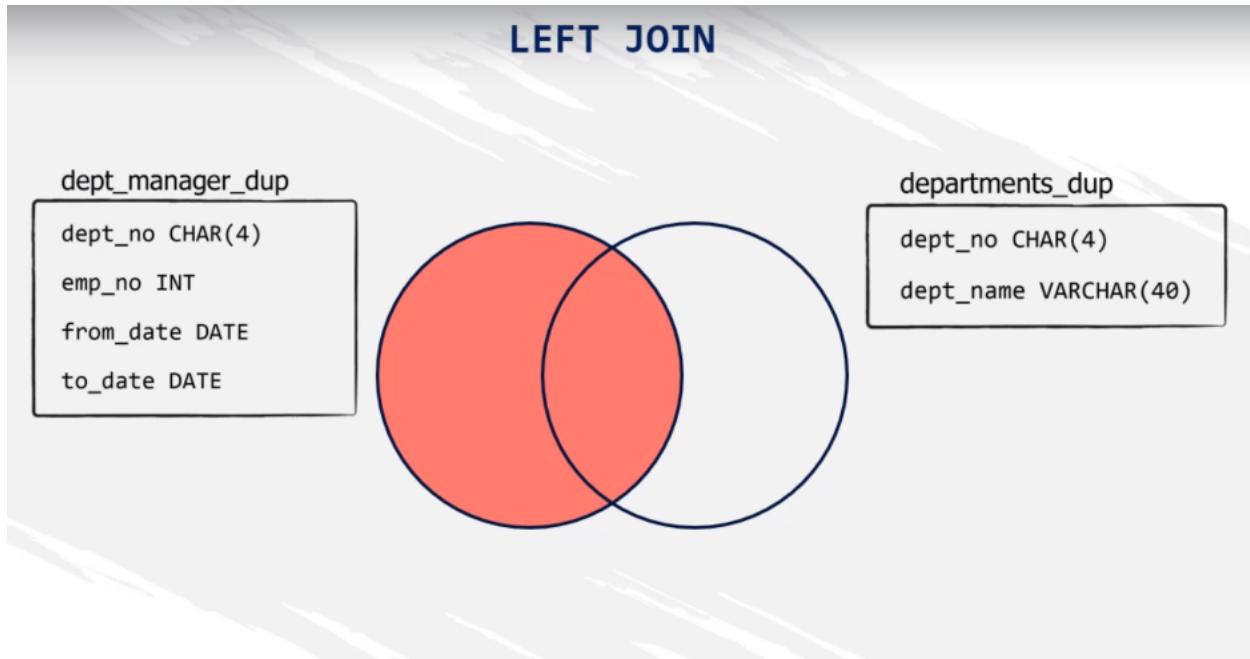
Output ::::::::::::

Action Output

| #  | Time     | Action                                                     | Message            |
|----|----------|------------------------------------------------------------|--------------------|
| 11 | 14:53:44 | SELECT m.dept_no,m.emp_no,d.dept_name FROM dept_manager... | 25 row(s) returned |
| 12 | 15:24:12 | SELECT m.dept_no,m.emp_no,d.dept_name FROM dept_manager... | 20 row(s) returned |

## 68. LEFT JOIN

LEFT JOIN is indicated in the set graph below:



- 1) all matching values of the two tables +
- 2) all values from the left table that match no values from the right table

Example: make a case of left join using the example from #67

Step1: delete the duplicates and add the same ones

```

7 # delete the duplicates
8 • DELETE FROM dept_manager_dup
9 WHERE emp_no = "110228";
10
11 • DELETE FROM departments_dup
12 WHERE dept_no = "d009";
13
14 # add back initial values
15 • INSERT INTO dept_manager_dup
16 VALUES ("110228","d003","1992-03-21","9999-01-01");
17
18 • INSERT INTO departments_dup
19 VALUES ("d009","Customer Service");
20

```

Step2: use LEFT JOIN instead of JOIN

```

22 # left join
23 • SELECT
24 m.dept_no, m.emp_no,d.dept_name
25 FROM
26 dept_manager_dup m
27 LEFT JOIN
28 departments_dup d ON m.dept_no = d.dept_no
29 GROUP BY m.emp_no
30 ORDER BY dept_no
31 ;

```

Comparison: the blue part refers to the 2) in the van diagram

10. VALUES ('110228', 'd003', '1992-03-21', '9999-01-01')

**26 rows (left join)**

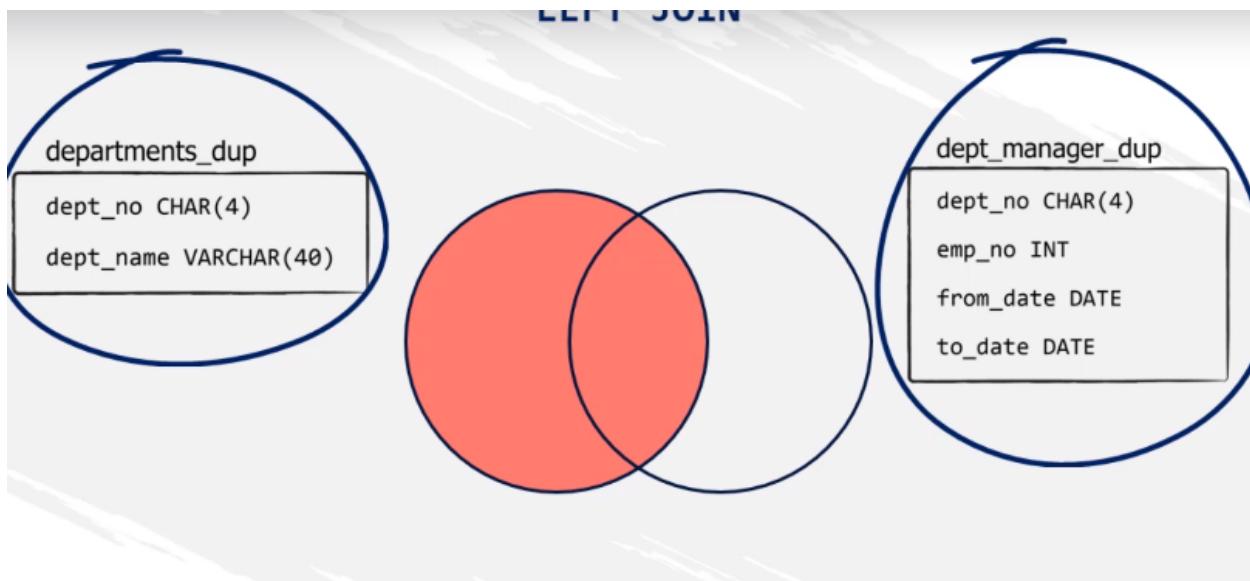
| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| HULL    | 999905 | HULL               |
| HULL    | 999907 | HULL               |
| HULL    | 999904 | HULL               |
| HULL    | 999906 | HULL               |
| d002    | 110085 | HULL               |
| d002    | 110114 | HULL               |
| d003    | 110183 | Human Resources    |
| d003    | 110228 | Human Resources    |
| d004    | 110420 | Production         |
| d004    | 110303 | Production         |
| d004    | 110386 | Production         |
| d004    | 110344 | Production         |
| d005    | 110511 | Development        |
| d005    | 110567 | Development        |
| d006    | 110765 | Quality Management |
| d006    | 110854 | Quality Management |
| d006    | 110725 | Quality Management |
| d006    | 110800 | Quality Management |
| d007    | 111133 | Sales              |
| d007    | 111035 | Sales              |
| d008    | 111534 | Research           |
| d008    | 111400 | Research           |
| d009    | 111692 | Customer Service   |
| d009    | 111877 | Customer Service   |
| d009    | 111784 | Customer Service   |
| d009    | 111939 | Customer Service   |

**20 rows (inner join)**

| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| d003    | 110228 | Human Resources    |
| d003    | 110183 | Human Resources    |
| d004    | 110344 | Production         |
| d004    | 110420 | Production         |
| d004    | 110303 | Production         |
| d004    | 110386 | Production         |
| d005    | 110567 | Development        |
| d005    | 110511 | Development        |
| d006    | 110800 | Quality Management |
| d006    | 110765 | Quality Management |
| d006    | 110854 | Quality Management |
| d006    | 110725 | Quality Management |
| d007    | 111035 | Sales              |
| d007    | 111133 | Sales              |
| d008    | 111400 | Research           |
| d008    | 111534 | Research           |
| d009    | 111784 | Customer Service   |
| d009    | 111939 | Customer Service   |
| d009    | 111692 | Customer Service   |
| d009    | 111877 | Customer Service   |

## 69. LEFT JOIN - MORE

The ordering of LEFT JOIN matters



After switching `departments_dup` and `dept_manager_dup`, the result# changed

```

21
22 # left join
23 • SELECT
24 m.dept_no, m.emp_no,d.dept_name
25 FROM
26 departments_dup d
27 LEFT JOIN
28 dept_manager_dup m ON m.dept_no = d.dept_no
29 GROUP BY m.emp_no
30 ORDER BY dept_no
31 ;

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| NULL    | NULL   | Business Analytics |
| d003    | 110183 | Human Resources    |
| d003    | 110228 | Human Resources    |
| d004    | 110303 | Production         |
| d004    | 110344 | Production         |
| d004    | 110386 | Production         |
| d004    | 110420 | Production         |
| d005    | 110511 | Development        |
| d005    | 110567 | Development        |

Result 2 ×

Output

Action Output

| #  | Time     | Action                                                        | Message            |
|----|----------|---------------------------------------------------------------|--------------------|
| 7  | 16:29:07 | DELETE FROM departments_dup WHERE dept_no = "d009"            | 1 row(s) affected  |
| 8  | 16:29:07 | INSERT INTO dept_manager_dup VALUES ("110228","d003","...")   | 1 row(s) affected  |
| 9  | 16:29:07 | INSERT INTO departments_dup VALUES ("d009","Customer Ser...") | 1 row(s) affected  |
| 10 | 16:29:07 | SELECT m.dept_no,m.emp_no,d.dept_name FROM departme...        | 21 row(s) returned |



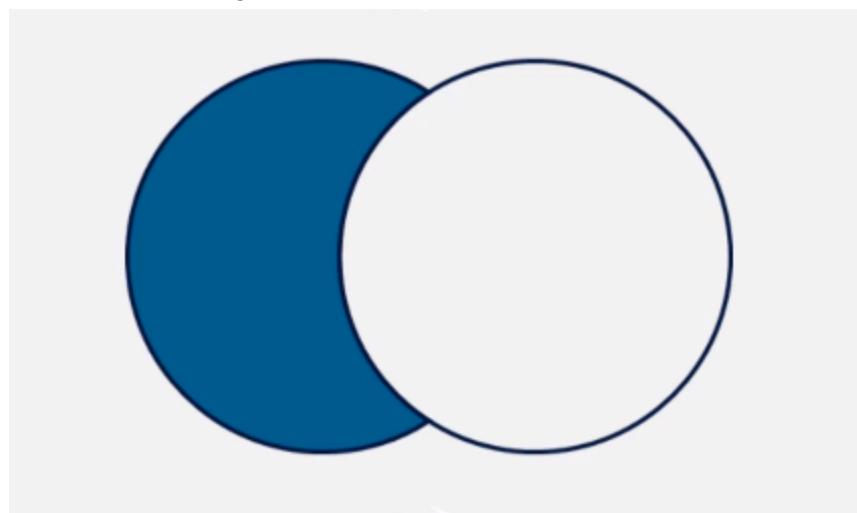
```
dept_name
.dept_no = d.dept_no
departments_dup d
```

| dept_no | dept_name          |
|---------|--------------------|
| NULL    | Public Relations   |
| d001    | Marketing          |
| d003    | Human Resources    |
| d004    | Production         |
| d005    | Development        |
| d006    | Quality Management |
| d007    | Sales              |
| d008    | Research           |
| d009    | Customer Service   |
| d010    | NULL               |
| d011    | NULL               |

same FROM departments\_dup d ... 24 row(s) returned

3

- **LEFT OUTER JOIN = LEFT JOIN**
- Left joins can deliver a list with all records from the left table that do not match any rows from the right table



For example, our goal is to ONLY show the blue parts in the code above. Since we know NULL is the common column value, it is important to share.

```
1 • SELECT
2 m.dept_no, m.emp_no, d.dept_name
3 FROM
4 dept_manager_dup m
5 LEFT JOIN
6 departments_dup d ON m.dept_no = d.dept_no
7
8 ORDER BY m.dept_no;
```

Result Grid | Filter Rows: [ ] | Export: [ ] | Wrap Cell Contents: [ ]

| dept_no | emp_no | dept_name       |
|---------|--------|-----------------|
| NULL    | 999905 | NULL            |
| NULL    | 999907 | NULL            |
| NULL    | 999904 | NULL            |
| NULL    | 999906 | NULL            |
| d002    | 110114 | NULL            |
| d002    | 110085 | NULL            |
| d003    | 110228 | Human Resources |
| d003    | 110183 | Human Resources |
| d004    | 110386 | Production      |
| d004    | 110344 | Production      |
| d004    | 110420 | Production      |

Result 14 ×

Output

| # | Time     | Action                                                           | Message            |
|---|----------|------------------------------------------------------------------|--------------------|
| 4 | 15:40:24 | SELECT d.dept_no,m.emp_no,d.dept_name FROM departments_dup d ... | 24 row(s) returned |
| 5 | 15:41:39 | SELECT m.dept_no,m.emp_no,d.dept_name FROM dept_manager_dup m... | 26 row(s) returned |

```

22 # left join
23 • SELECT
24 m.dept_no, m.emp_no,d.dept_name
25 FROM
26 dept_manager_dup m
27 LEFT JOIN
28 departments_dup d ON m.dept_no = d.dept_no
29 WHERE dept_name IS NULL
30 #GROUP BY m.emp_no
31 ORDER BY m.dept_no
32 ;
33

```

The screenshot shows a database query results grid. At the top, there are buttons for 'Result Grid' (selected), 'Filter Rows:', 'Export:', and 'Wrap Cell'. The grid has three columns: 'dept\_no', 'emp\_no', and 'dept\_name'. The data is as follows:

|   | dept_no | emp_no | dept_name |
|---|---------|--------|-----------|
| ▶ | NULL    | 999904 | NULL      |
|   | NULL    | 999905 | NULL      |
|   | NULL    | 999906 | NULL      |
|   | NULL    | 999907 | NULL      |
|   | d002    | 110085 | NULL      |
|   | d002    | 110114 | NULL      |

**Exercise:** Join the 'employees' and the 'dept\_manager' tables to return a subset of all the employees whose last name is Markovitch. See if the output contains a manager with that name.

**Hint:** Create an output containing information corresponding to the following fields: 'emp\_no', 'first\_name', 'last\_name', 'dept\_no', 'from\_date'. Order by 'dept\_no' descending, and then by 'emp\_no'.

```

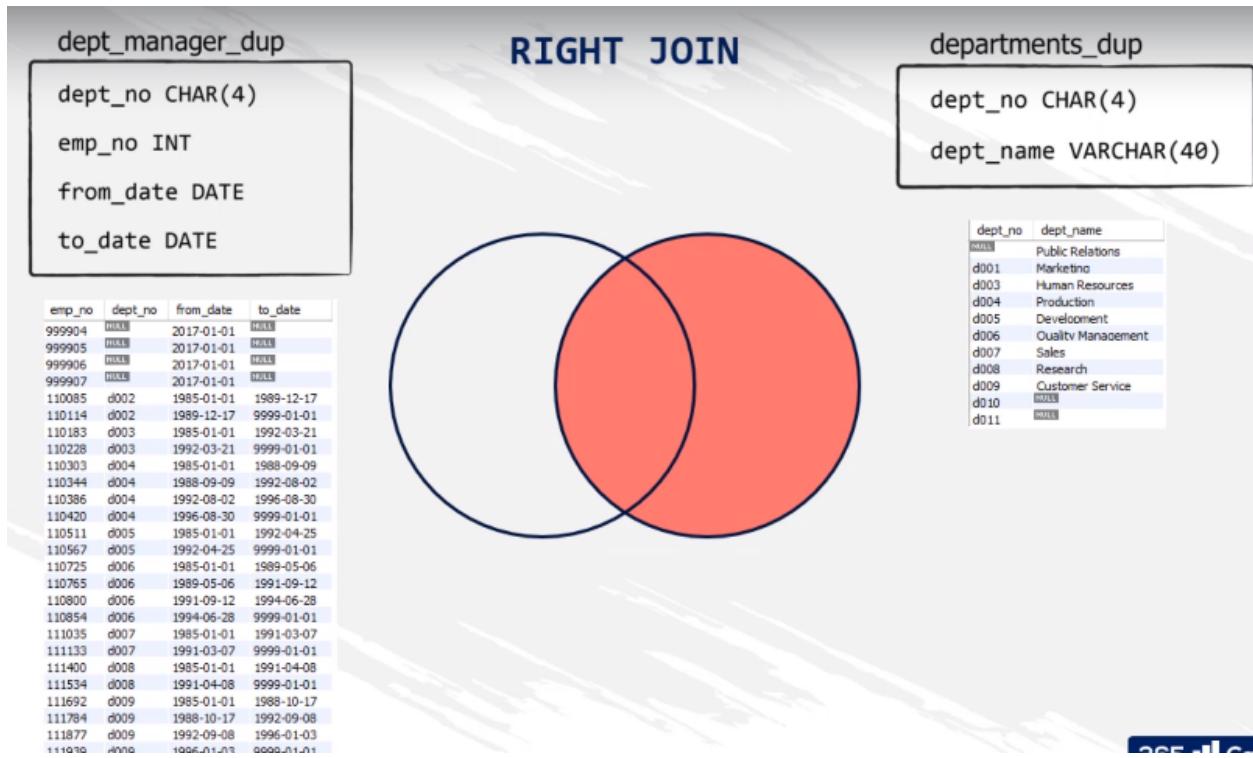
left join
SELECT
 e.emp_no,e.first_name,e.last_name,dm.dept_no,dm.from_date
FROM
 employees e
 LEFT JOIN
 dept_manager_dup dm ON e.emp_no = dm.emp_no
 WHERE e.last_name = "Markovitch"
ORDER BY dm.dept_no DESC, e.emp_no
;

```

**Two things:** (1) two parameters for ORDER BY, (2) ASC vs. DESC

## 70. RIGHT JOIN (RIGHT OUT JOIN)

Logically similar to LEFT JOIN



```

21 • SELECT
22 m.dept_no, m.emp_no, d.dept_name
23 FROM
24 dept_manager_dup m
25 RIGHT JOIN
26 departments_dup d ON m.dept_no = d.dept_no
27 ORDER BY dept_no;
28

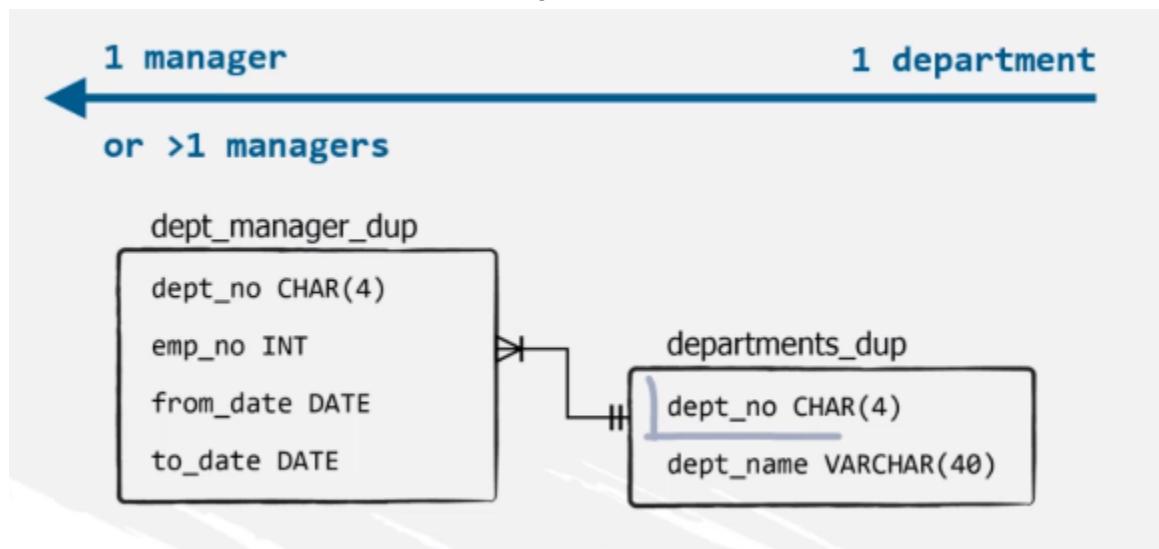
```

Result Grid | Filter Rows: [ ] | Export: [ ] | Wrap Cell Cont

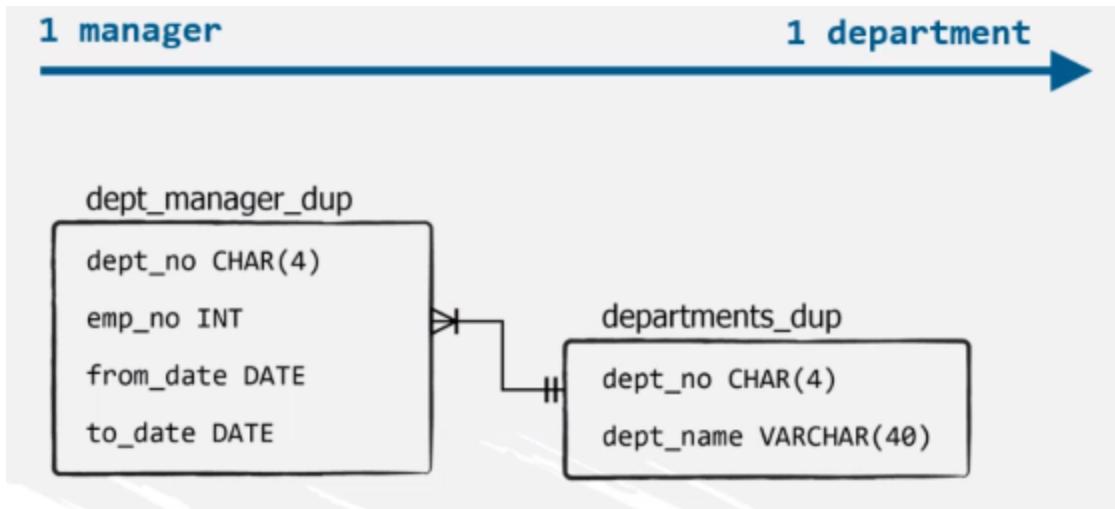
| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| NULL    | NULL   | Business Analytics |
| NULL    | NULL   | Marketing          |
| NULL    | NULL   | Public Relations   |
| NULL    | NULL   | NULL               |
| NULL    | NULL   | NULL               |
| d003    | 110183 | Human Resources    |
| d003    | 110228 | Human Resources    |
| d004    | 110303 | Production         |
| d004    | 110344 | Production         |

- A LEFT JOIN B = B RIGHT JOIN A; RIGHT JOIN is rarely applied
- LEFT/RIGHT JOINS are examples of one-to-many relationships

One department can have multiple managers



One manager only has one department



## 71. The NEW and the OLD JOIN syntax

New contents coming soon: (1) subqueries, self-joins, (2) stored routines, and (3) advanced SQL tools

*The same example of using JOIN*

```

9 • SELECT
10 m.dept_no, m.emp_no, d.dept_name
11 FROM
12 dept_manager_dup m
13 INNER JOIN
14 departments_dup d ON m.dept_no = d.dept_no
15 ORDER BY m.dept_no;

```

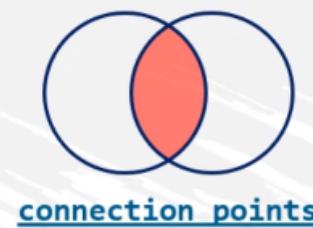
## The New and the Old Join Syntax

| emp_no | dept_no | from_date  | to_date    |
|--------|---------|------------|------------|
| 999904 | NULL    | 2017-01-01 | NULL       |
| 999905 | NULL    | 2017-01-01 | NULL       |
| 999906 | NULL    | 2017-01-01 | NULL       |
| 999907 | NULL    | 2017-01-01 | NULL       |
| 110085 | d002    | 1985-01-01 | 1989-12-17 |
| 110114 | d002    | 1989-12-17 | 9999-01-01 |
| 110183 | d003    | 1985-01-01 | 1992-03-21 |
| 110228 | d003    | 1992-03-21 | 9999-01-01 |
| 110303 | d004    | 1985-01-01 | 1988-09-09 |
| 110344 | d004    | 1988-09-09 | 1992-08-02 |
| 110386 | d004    | 1992-08-02 | 1996-08-30 |
| 110420 | d004    | 1996-08-30 | 9999-01-01 |
| 110511 | d005    | 1985-01-01 | 1992-04-25 |
| 110567 | d005    | 1992-04-25 | 9999-01-01 |
| 110725 | d006    | 1985-01-01 | 1989-05-06 |
| 110765 | d006    | 1989-05-06 | 1991-09-12 |
| 110800 | d006    | 1991-09-12 | 1994-06-28 |
| 110854 | d006    | 1994-06-28 | 9999-01-01 |
| 111035 | d007    | 1985-01-01 | 1991-03-07 |
| 111133 | d007    | 1991-03-07 | 9999-01-01 |
| 111400 | d008    | 1985-01-01 | 1991-04-08 |
| 111534 | d008    | 1991-04-08 | 9999-01-01 |
| 111692 | d009    | 1985-01-01 | 1988-10-17 |
| 111784 | d009    | 1988-10-17 | 1992-09-08 |
| 111877 | d009    | 1992-09-08 | 1996-01-03 |
| 111939 | d009    | 1996-01-03 | 9999-01-01 |

dept\_manager\_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```

| dept_no | emp_no | dept_name          |
|---------|--------|--------------------|
| d003    | 110228 | Human Resources    |
| d003    | 110183 | Human Resources    |
| d004    | 110344 | Production         |
| d004    | 110420 | Production         |
| d004    | 110303 | Production         |
| d005    | 110386 | Production         |
| d005    | 110567 | Development        |
| d005    | 110511 | Development        |
| d006    | 110800 | Quality Management |
| d006    | 110765 | Quality Management |
| d006    | 110854 | Quality Management |
| d006    | 110725 | Quality Management |
| d007    | 111035 | Sales              |
| d007    | 111133 | Sales              |
| d008    | 111400 | Research           |
| d008    | 111534 | Research           |
| d009    | 111784 | Customer Service   |
| d009    | 111939 | Customer Service   |
| d009    | 111692 | Customer Service   |
| d009    | 111877 | Customer Service   |



| dept_no | dept_name          |
|---------|--------------------|
| NULL    | Public Relations   |
| d001    | Marketing          |
| d003    | Human Resources    |
| d004    | Production         |
| d005    | Development        |
| d006    | Quality Management |
| d007    | Sales              |
| d008    | Research           |
| d009    | Customer Service   |
| d010    | NULL               |
| d011    | NULL               |

departments\_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```

This same result can be achieved by WHERE

```
SELECT
 t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
 table_1 t1,
 table_2 t2
WHERE
 t1.column_name = t2.column_name;
```

```
SELECT
 m.dept_no, m.emp_no, d.dept_name
FROM
 dept_manager_dup m,
 departments_dup d
WHERE
 m.dept_no = d.dept_no
```

```
ORDER BY m.dept_no;
```

#### JOIN or WHERE?

- The retrieved result is the same
- Using WHERE is more time consuming; WHERE is old and slow
- JOIN allows you to modify the table-to-table relationships easily

#### Exercise:

Extract a list containing information about all managers' employee number, first and last name, department number, and hire date. Use the old type of join syntax to obtain the result.

```
JOIN - new
SELECT
 e.emp_no,e.first_name,e.last_name,dm.dept_no,e.hire_date
FROM
 employees e
 JOIN
 dept_manager dm ON e.emp_no = dm.emp_no
;
```

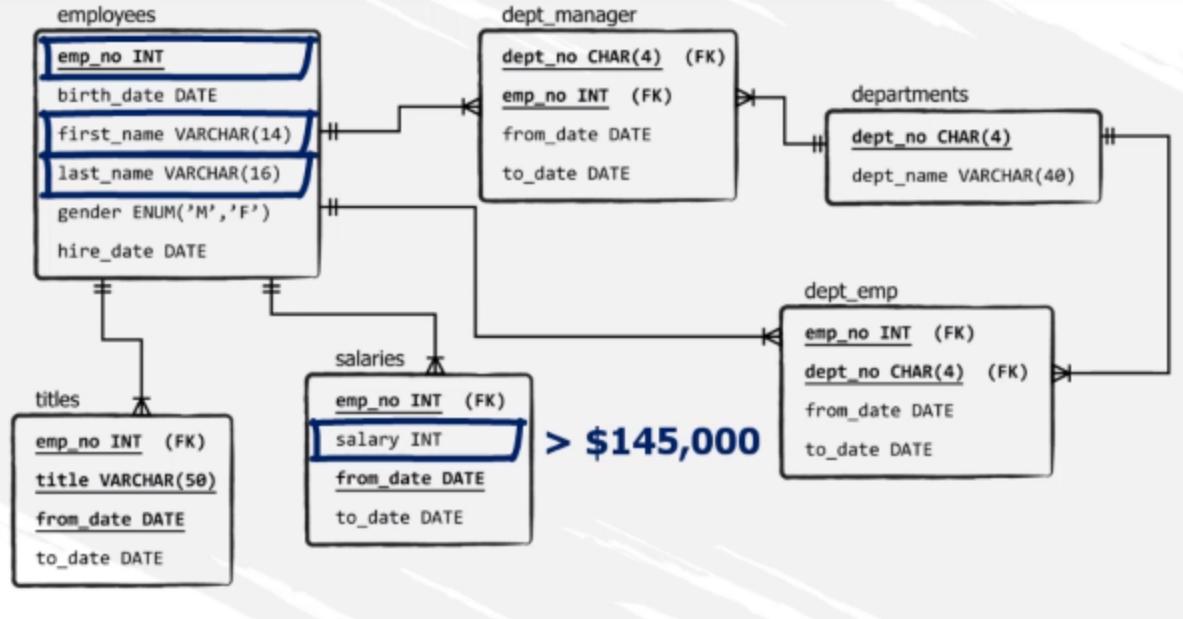
```
WHERE - old
SELECT
 e.emp_no,e.first_name,e.last_name,dm.dept_no,e.hire_date
FROM
 employees e,
 dept_manager dm
WHERE
 e.emp_no = dm.emp_no
;
```

#### 72. JOIN and WHERE used together

When WHERE is used for JOIN, JOIN is used for connecting the “employees” and “Salaries” tables, and **WHERE is used to define the condition or conditions that will determine which will be connecting points(intersection) between the two tables**

Example:

## JOIN and WHERE Used Together



The question is to find out the connecting point data with a salary higher than 145k.

```

SELECT
 e.emp_no, e.first_name, e.last_name, s.salary
FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no
WHERE
 s.salary > 145000;

```

**Exercise:**

**Important – Prevent Error Code: 1055!**

Depending on your operating system and version of MySQL, you will be working with different SQL settings.

To make sure you can take some of the remaining lectures of the course without unnecessary interruption, we strongly advise you to execute the following query now.

```

set @@global.sql_mode := replace(@@global.sql_mode,
'ONLY_FULL_GROUP_BY', '') ;

undo this change:
set @@global.sql_mode := concat('ONLY_FULL_GROUP_BY',
@@global.sql_mode) ;

```

Select the first and last name, the hire date, and the job title of all employees whose first name is “Margareta” and have the last name “Markovitch”.

```

SELECT
 e.first_name, e.last_name, e.hire_date, t.title
FROM employees e
 JOIN
titles t ON e.emp_no=t.emp_no
WHERE e.first_name = "margareta" AND e.last_name = "Markovitch"
ORDER BY e.emp_no
;

```

## 73. CROSS JOIN

- (1) A cross join will take the values from a certain table and connect them with all the values from the tables we want to join with.

INTER JOIN - typically connects only the matching values

CROSS JOIN

- connects all the values, not just those that match
- The cartesian product of the values of two or more sets
- Particularly useful when the tables in a database are not well connected

Example: fully join values from two tables

| dept_manager |         |            |            | departments |                    |
|--------------|---------|------------|------------|-------------|--------------------|
| emp_no       | dept_no | from_date  | to_date    | dept_no     | dept_name          |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d001        | Marketing          |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d002        | Finance            |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d003        | Human Resources    |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d004        | Production         |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d005        | Development        |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d006        | Quality Management |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d007        | Sales              |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d008        | Research           |
| 110022       | d001    | 1985-01-01 | 1991-10-01 | d009        | Customer Service   |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d001        | Marketing          |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d002        | Finance            |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d003        | Human Resources    |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d004        | Production         |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d005        | Development        |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d006        | Quality Management |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d007        | Sales              |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d008        | Research           |
| 110039       | d001    | 1991-10-01 | 9999-01-01 | d009        | Customer Service   |
| 110085       | d002    | 1985-01-01 | 1989-12-17 | d001        | Marketing          |
| 110085       | d002    | 1985-01-01 | 1989-12-17 | d002        | Finance            |
| 110085       | d002    | 1985-01-01 | 1989-12-17 | d003        | Human Resources    |
| 110085       | d002    | 1985-01-01 | 1989-12-17 | d004        | Production         |

```
SELECT
 dm.* , d.*
FROM
 dept_manager dm
 CROSS JOIN
 departments d
ORDER BY dm.emp_no, d.dept_no;
```

**dept\_manager    departments**

| emp_no  | dept_no |
|---------|---------|
| 110022  | d001    |
| 110039  | d002    |
| 110085  | d003    |
| 110114  | d004    |
| 110183  | d005    |
| 110228  | d006    |
| 110303  | d007    |
| 110344  | d008    |
| 110386  | d009    |
| 110420  |         |
| 110511  |         |
| 110567  |         |
| 110725  |         |
| 110765  |         |
| 110800  |         |
| 110854  |         |
| 111035  |         |
| 111133  |         |
| 111400  |         |
| 111534  |         |
| 111692  |         |
| Message |         |
| 111784  |         |
| 111877  |         |
| 111939  |         |

365

The result is a column x column multiplication

RESULT GRID | FILTER ROWS | EXPAND | WRAP CELL CONTENTS |

|   | emp_no | dept_no | from_date  | to_date    | dept_no | dept_name          |
|---|--------|---------|------------|------------|---------|--------------------|
| ▶ | 110022 | d001    | 1985-01-01 | 1991-10-01 | d001    | Marketing          |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d002    | Finance            |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d003    | Human Resources    |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d004    | Production         |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d005    | Development        |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d006    | Quality Management |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d007    | Sales              |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d008    | Research           |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d009    | Customer Service   |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d010    | Business Analytics |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d001    | Marketing          |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d002    | Finance            |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d003    | Human Resources    |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d004    | Production         |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d005    | Development        |

(2) Another way to write it without JOIN (this version is just slower)

```
SELECT
 dm.* , d.*
FROM
 dept_manager dm,
 departments d
```

```
ORDER BY dm.emp_no, d.dept_no;
```

Since there is no assigned condition, using JOIN, CROSS JOIN, and INNER JOIN are the same thing here

**JOIN without ON = not considered best practice**

**CROSS JOIN = best practice**

- (3) Also, we can specify that the dept\_no should be different in the combined tables by using <> in WHERE

**JOIN + ON clause = CROSS JOIN +WHERE clause**

```
SELECT
 dm.* , d.*
FROM
 dept_manager dm,
 departments d
WHERE d.dept_no <> dm.dept_no
ORDER BY dm.emp_no, d.dept_no;
```

|   | emp_no | dept_no | from_date  | to_date    | dept_no | dept_name          |
|---|--------|---------|------------|------------|---------|--------------------|
| ▶ | 110022 | d001    | 1985-01-01 | 1991-10-01 | d002    | Finance            |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d003    | Human Resources    |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d004    | Production         |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d005    | Development        |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d006    | Quality Management |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d007    | Sales              |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d008    | Research           |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d009    | Customer Service   |
|   | 110022 | d001    | 1985-01-01 | 1991-10-01 | d010    | Business Analytics |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d002    | Finance            |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d003    | Human Resources    |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d004    | Production         |
|   | 110039 | d001    | 1991-10-01 | 9999-01-01 | d005    | Development        |

- (4) Using CROSS JOIN whenever it is possible
- (5) CROSS JOIN can be used for more than two tables; it's possible that the result can be too big

Example: three tables

```

SELECT
 e.*, d.*
FROM
 departments d
 CROSS JOIN
 dept_manager dm
 JOIN
 employees e ON dm.emp_no = e.emp_no
WHERE d.dept_no <> dm.dept_no
ORDER BY dm.emp_no, d.dept_no;

```

|  | emp_no | birth_date | first_name | last_name  | gender | hire_date  | dept_no | dept_name          |
|--|--------|------------|------------|------------|--------|------------|---------|--------------------|
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d005    | Development        |
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d006    | Quality Management |
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d007    | Sales              |
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d008    | Research           |
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d009    | Customer Service   |
|  | 110022 | 1956-09-12 | Margareta  | Markovitch | M      | 1985-01-01 | d010    | Business Analytics |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d002    | Finance            |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d003    | Human Resources    |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d004    | Production         |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d005    | Development        |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d006    | Quality Management |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d007    | Sales              |
|  | 110039 | 1963-06-21 | Vishwani   | Minakawa   | M      | 1986-04-12 | d008    | Research           |

### Exercise 1:

Use a CROSS JOIN to return a list with all possible combinations between managers from the dept\_manager table and department number 9.

```

SELECT
 dm.*, d.*
FROM
 dept_manager dm
CROSS JOIN
 departments d

```

```
WHERE dm.dept_no = "d009"
ORDER BY d.dept_no;
```

**Exercise 2:**

Return a list with the first 10 employees with all the departments they can be assigned to.

Hint: Don't use LIMIT; use a WHERE clause.

```
SELECT
 e.* , d.*
FROM
 employees e
CROSS JOIN
 departments d
WHERE e.emp_no < 10011
ORDER BY e.emp_no,d.dept_no;
```

→ use emp\_no to limit the number of employees (the 11th person not included)

## 74. Aggregate functions and JOIN

Example: find the average salaries of men and women in the company

```
7 • SELECT
8 e.gender, AVG(s.salary) AS average_salary
9 FROM
10 employees e
11 JOIN
12 salaries s ON e.emp_no=s.emp_no
13 GROUP BY gender;
14
```

| gender | average_salary |
|--------|----------------|
| M      | 63755.9134     |
| F      | 63769.1222     |

Possible problem: delivering first values is simply how MySQL behaves when it is asked to extract a single value from a column that groups data

The screenshot shows the MySQL Workbench interface. On the left, a query editor window displays the following SQL code:

```

1 • SELECT
2 e.emp_no, e.gender, AVG(s.salary) AS average_salary
3 FROM
4 employees e
5 JOIN
6 salaries s ON e.emp_no = s.emp_no
7 GROUP BY gender;

```

The result grid shows the output of the query:

| emp_no | gender | average_salary |
|--------|--------|----------------|
| 10001  | M      | 63755.9134     |
| 10002  | F      | 63769.1222     |

On the right, two table definitions are shown:

- employees** table:
 

```

emp_no INT
birth_date DATE
first_name VARCHAR(14)
last_name VARCHAR(16)
gender ENUM('M','F')
hire_date DATE

```
- salaries** table:
 

```

emp_no INT (FK)
salary INT
from_date DATE
to_date DATE

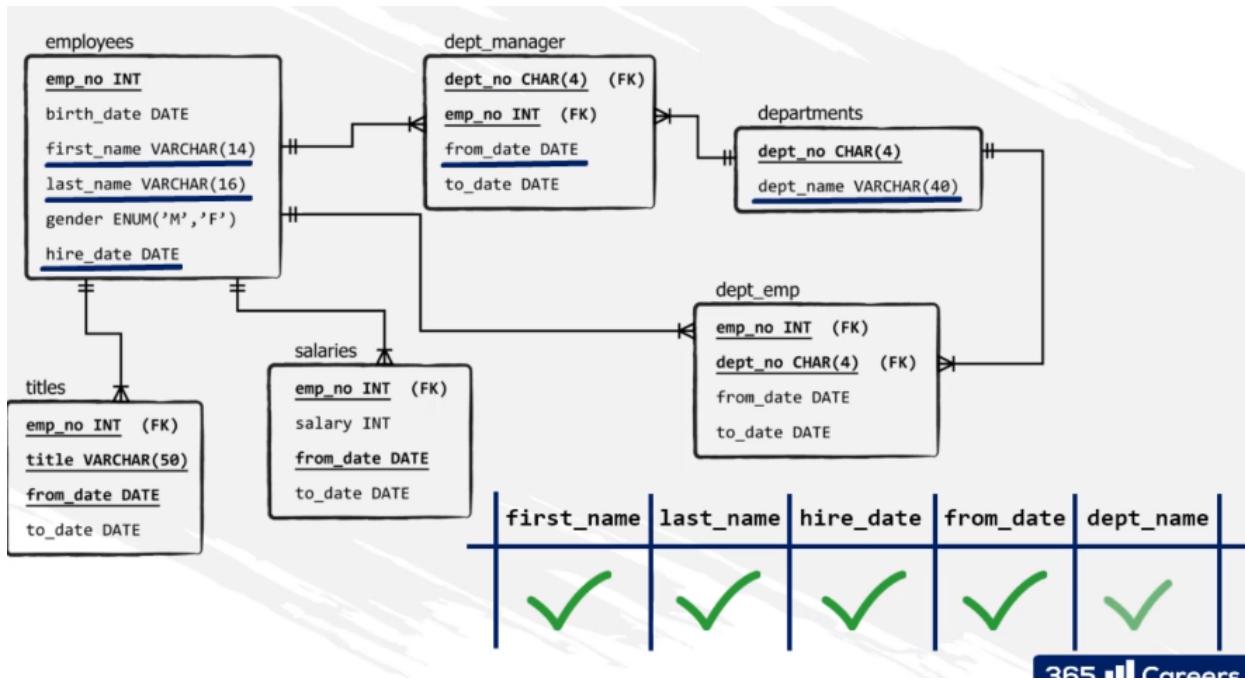
```

There is no real meaning to put emp\_no as the first column

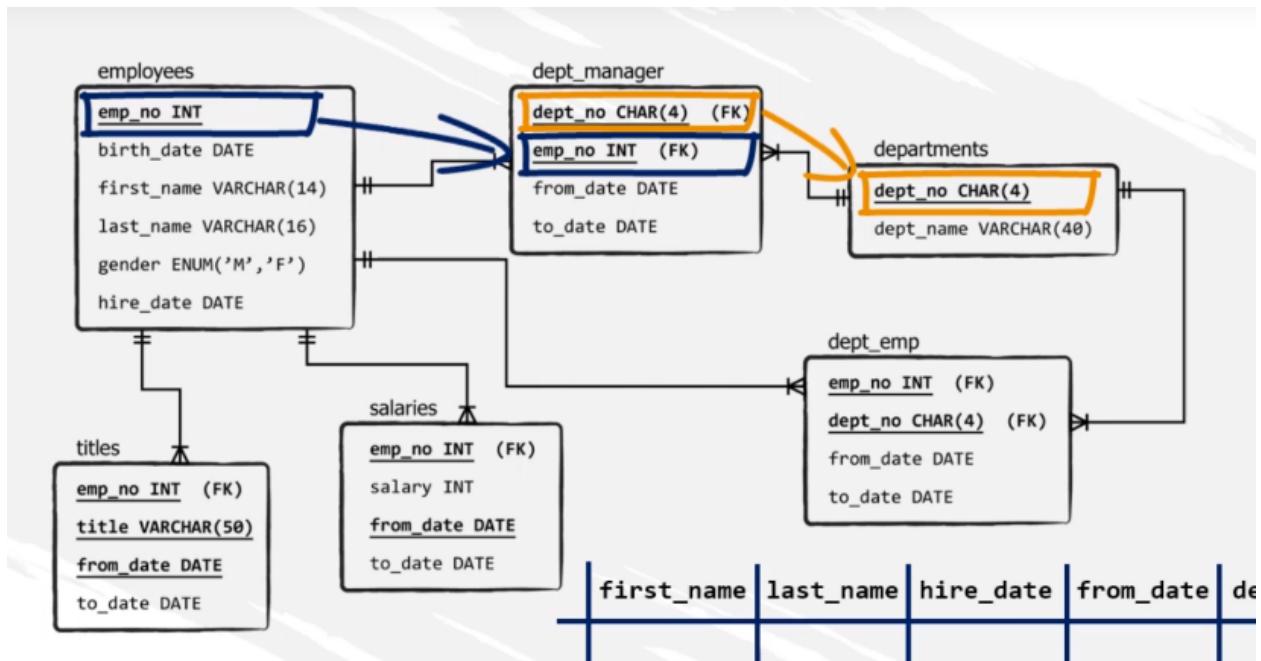
## 75. JOIN more than two tables

**When creating a query that joins multiple tables, you must back it with a strong intuition and crystal-clear idea of how you would like the tables to be connected**

Example: join three tables



- Employees and dept\_manager can be JOINed by emp\_no
- Dept\_number and departments can be JOINed by dept\_no



SELECT

```
e.first_name,
e.last_name,
```

```

e.hire_date,
m.from_date,
d.dept_name
FROM
 employees e
 JOIN
 dept_manager m ON e.emp_no = m.emp_no # this is the table with two matching
columns
 JOIN
 departments d ON m.dept_no = d.dept_no
;

```

|   | first_name | last_name    | hire_date  | from_date  | dept_name       |
|---|------------|--------------|------------|------------|-----------------|
| ▶ | Margareta  | Markovitch   | 1985-01-01 | 1985-01-01 | Marketing       |
|   | Vishwani   | Minakawa     | 1986-04-12 | 1991-10-01 | Marketing       |
|   | Ebru       | Alpin        | 1985-01-01 | 1985-01-01 | Finance         |
|   | Isamu      | Legleitner   | 1985-01-14 | 1989-12-17 | Finance         |
|   | Shirish    | Ossenbruggen | 1985-01-01 | 1985-01-01 | Human Resources |
|   | Karsten    | Sigstam      | 1985-08-04 | 1992-03-21 | Human Resources |
|   | Krassimir  | Wegerle      | 1985-01-01 | 1985-01-01 | Production      |
|   | Rosine     | Cools        | 1985-11-22 | 1988-09-09 | Production      |

There is no change when you switch departments and employees

```
1 • SELECT
2 e.first_name,
3 e.last_name,
4 e.hire_date,
5 m.from_date,
6 d.dept_name
7 FROM
8 departments d
9 JOIN
10 dept_manager m ON d.dept_no = m.dept_no
11 JOIN
12 employees e ON m.emp_no = e.emp_no
13 ;
14
15
```

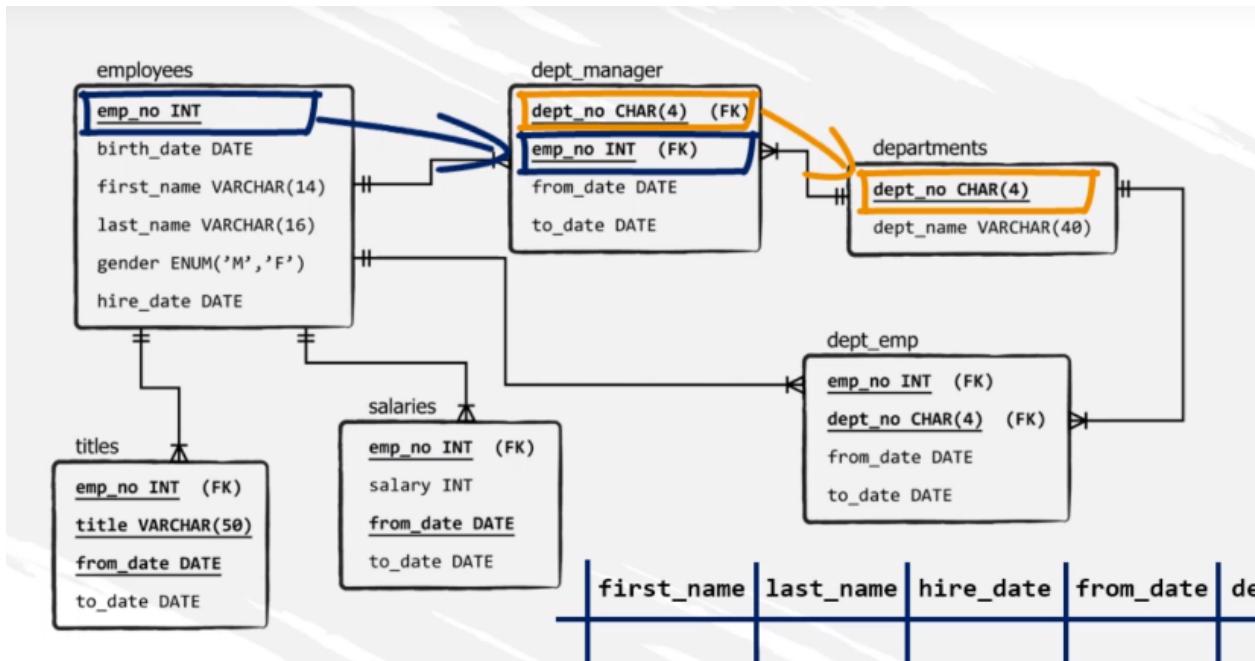
| first_name | last_name    | hire_date  | from_date  | dept_name       |
|------------|--------------|------------|------------|-----------------|
| Leon       | DasSarma     | 1986-10-21 | 1992-04-25 | Development     |
| Ebri       | Aloin        | 1985-01-01 | 1985-01-01 | Finance         |
| Isamu      | Leleitner    | 1985-01-14 | 1989-12-17 | Finance         |
| Shirish    | Ossenbruegen | 1985-01-01 | 1985-01-01 | Human Resources |
|            |              |            | 1992-03-21 | Human Resources |
|            |              |            | 1985-01-01 | Marketing       |
| Vishwani   | Minakawa     | 1986-04-12 | 1991-10-01 | Marketing       |
| Krassimir  | Weoerle      | 1985-01-01 | 1985-01-01 | Production      |

Changing the JOIN type to left/right join can produce different results.

```
• SELECT
 e.first_name,
 e.last_name,
 e.hire_date,
 m.from_date,
 d.dept_name
 FROM
 departments d
 JOIN
 dept_manager m ON d.dept_no = m.dept_no
 RIGHT JOIN
 employees e ON m.emp_no = e.emp_no
;
```

**Exercise:**

Select all managers' first and last name, hire date, job title, start date, and department name.



Manager is also a condition!!!

Key1:

```

SELECT
 e.first_name,
 e.last_name,
 e.hire_date,
 t.title,
 t.from_date,
 d.dept_name
FROM
 titles t
 JOIN
 employees e ON t.emp_no = e.emp_no
 JOIN
 dept_manager dm ON e.emp_no = dm.emp_no
 JOIN
 departments d ON dm.dept_no = d.dept_no
WHERE t.title = "Manager"
ORDER BY e.emp_no
;

```

Key2:

```
SELECT
```

```

e.first_name,
e.last_name,
e.hire_date,
t.title,
m.from_date,
d.dept_name

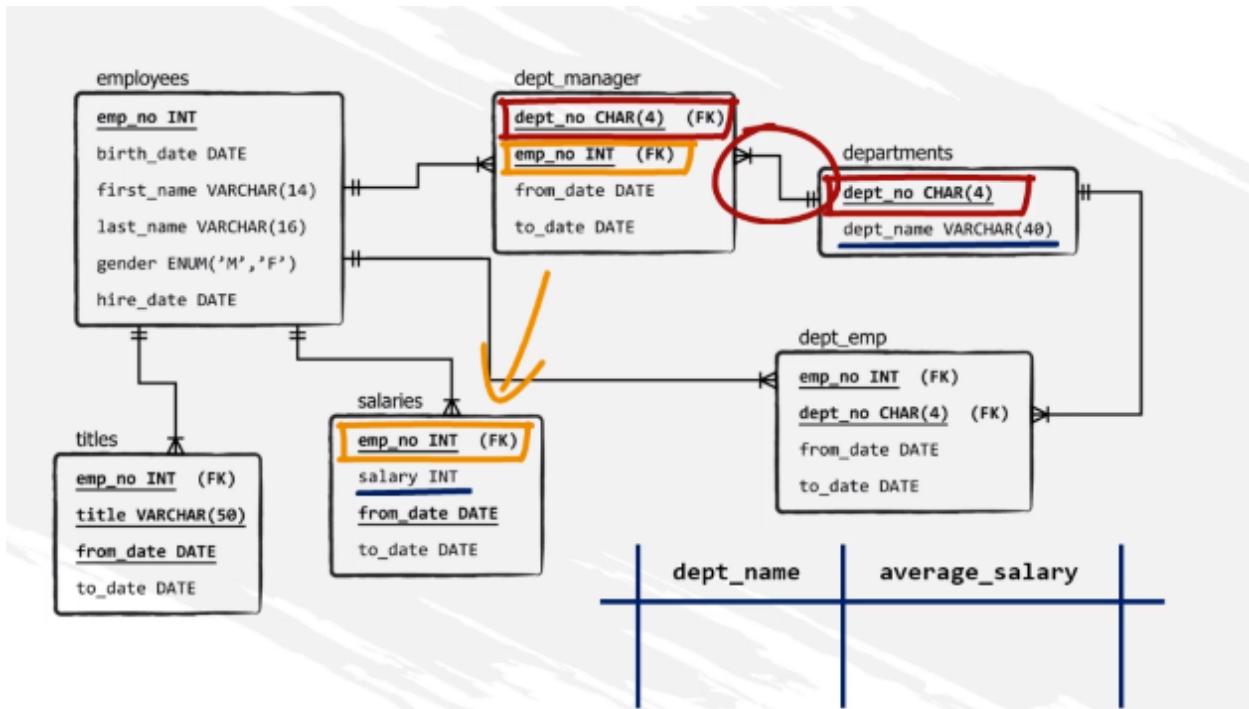
FROM
employees e
JOIN
dept_manager m ON e.emp_no = m.emp_no
JOIN
departments d ON m.dept_no = d.dept_no
JOIN
titles t ON e.emp_no = t.emp_no
AND m.from_date = t.from_date
ORDER BY e.emp_no;

```

## 76. Tips and tricks of JOINs

- One should look for key columns, which are common between tables involved in the analysis and are necessary to solve the task
- These columns do not need to be primary or foreign keys

For example, in order to join dept\_name and average\_salary, there is no need to use a fourth table like employees. The red and orange show the sharing matching columns



```

SELECT
 d.dept_name, AVG(salary)
FROM
 departments d
 JOIN
 dept_manager dm ON d.dept_no = dm.dept_no
 JOIN
 salaries s ON dm.emp_no = s.emp_no
GROUP BY d.dept_name
;

```

|   | dept_name          | AVG(salary) |
|---|--------------------|-------------|
| ▶ | Customer Service   | 54959.6724  |
|   | Development        | 59658.1176  |
|   | Finance            | 70815.8889  |
|   | Human Resources    | 58286.0556  |
|   | Marketing          | 88371.6857  |
|   | Production         | 56233.4000  |
|   | Quality Management | 67130.9355  |
|   | Research           | 77535.1818  |
|   | Sales              | 85738.7647  |

(1) IMPORT: GROUP BY must be added. Otherwise only the first row will be shown.

```

7 • SELECT
8 d.dept_name, AVG(salary)
9 FROM
10 departments d
11 JOIN
12 dept_manager dm ON d.dept_no = dm.dept_no
13 JOIN
14 salaries s ON dm.emp_no = s.emp_no
15 #GROUP BY d.dept_name
16 ;
17

```

| Result Grid      |             |
|------------------|-------------|
| dept_name        | AVG(salary) |
| Customer Service | 66924.2706  |

(2) We can also list the results according to dept\_no by using ORDER BY:

```

7 • SELECT
8 d.dept_name, AVG(salary)
9 FROM
10 departments d
11 JOIN
12 dept_manager dm ON d.dept_no = dm.dept_no
13 JOIN
14 salaries s ON dm.emp_no = s.emp_no
15 GROUP BY d.dept_name
16 ORDER BY d.dept_no
17 ;

```

| Result Grid        |             |
|--------------------|-------------|
| dept_name          | AVG(salary) |
| Marketing          | 88371.6857  |
| Finance            | 70815.8889  |
| Human Resources    | 58286.0556  |
| Production         | 56233.4000  |
| Development        | 59658.1176  |
| Quality Management | 67130.9355  |
| Sales              | 85738.7647  |
| Research           | 77535.1818  |
| Customer Service   | 54959.6724  |

(3) adding alias and sorting the results by numbers

```
7 • SELECT
8 d.dept_name, AVG(salary) AS average_salary
9 FROM
10 departments d
11 JOIN
12 dept_manager dm ON d.dept_no = dm.dept_no
13 JOIN
14 salaries s ON dm.emp_no = s.emp_no
15 GROUP BY d.dept_name
16 ORDER BY AVG(salary) DESC #or just average_salary
17 ;
```

The screenshot shows a database query result grid. At the top, there are buttons for 'Result Grid' (selected), 'Filter Rows', 'Export', and 'Wrap Cell Content'. The result grid has two columns: 'dept\_name' and 'average\_salary'. The data rows are: Marketing (88371.6857), Sales (85738.7647), Research (77535.1818), Finance (70815.8889), Quality Management (67130.9355), Development (59658.1176), Human Resources (58286.0556), Production (56233.4000), and Customer Service (54959.6724). The 'Sales' row is highlighted.

| dept_name          | average_salary |
|--------------------|----------------|
| Marketing          | 88371.6857     |
| Sales              | 85738.7647     |
| Research           | 77535.1818     |
| Finance            | 70815.8889     |
| Quality Management | 67130.9355     |
| Development        | 59658.1176     |
| Human Resources    | 58286.0556     |
| Production         | 56233.4000     |
| Customer Service   | 54959.6724     |

(4) Setting up value conditions for salary

```
SELECT
 d.dept_name, AVG(salary) AS average_salary
FROM
 departments d
 JOIN
 dept_manager dm ON d.dept_no = dm.dept_no
 JOIN
 salaries s ON dm.emp_no = s.emp_no
GROUP BY d.dept_name
HAVING average_salary > 60000
ORDER BY AVG(salary) DESC #or just average_salary
;
```

```

7 • SELECT
8 d.dept_name, AVG(salary) AS average_salary
9 FROM
10 departments d
11 JOIN
12 dept_manager dm ON d.dept_no = dm.dept_no
13 JOIN
14 salaries s ON dm.emp_no = s.emp_no
15 GROUP BY d.dept_name
16 HAVING average_salary > 60000
17 ORDER BY AVG(salary) DESC #or just average_salary
18 ;

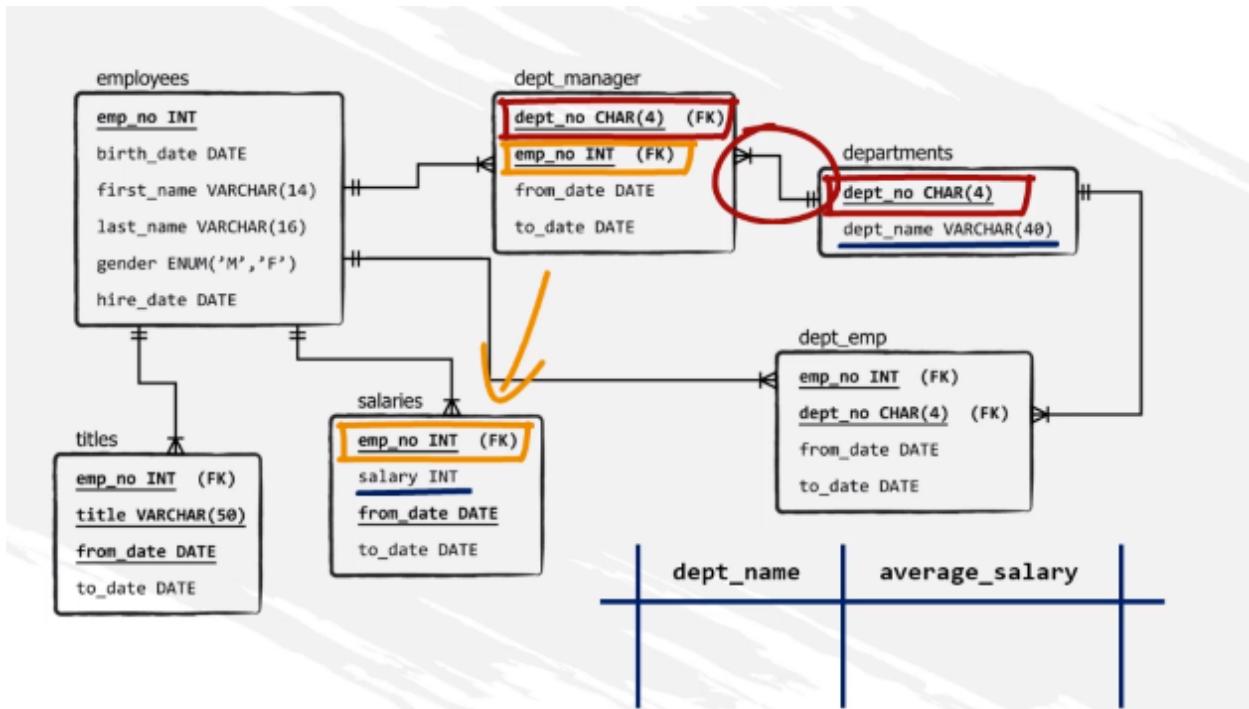
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

| dept_name          | average_salary |
|--------------------|----------------|
| Marketing          | 88371.6857     |
| Sales              | 85738.7647     |
| Research           | 77535.1818     |
| Finance            | 70815.8889     |
| Quality Management | 67130.9355     |

**Exercise:**

**How many male and how many female managers do we have in the ‘employees’ database?**



```

SELECT e.gender, COUNT(dm.emp_no)
FROM
 employees e
 JOIN
 dept_manager dm ON e.emp_no = dm.emp_no
GROUP BY e.gender;

```

## 77. UNION and UNION ALL

We make a duplicate of the first 20 rows of the employees table; then we add a new instance

```

DROP TABLE IF EXISTS employees_dup;
CREATE TABLE employees_dup (
 emp_no int(11),
 birth_date date,
 first_name varchar(14),
 last_name varchar(16),
 gender enum("M","F"),
 hire_date date
);

```

```

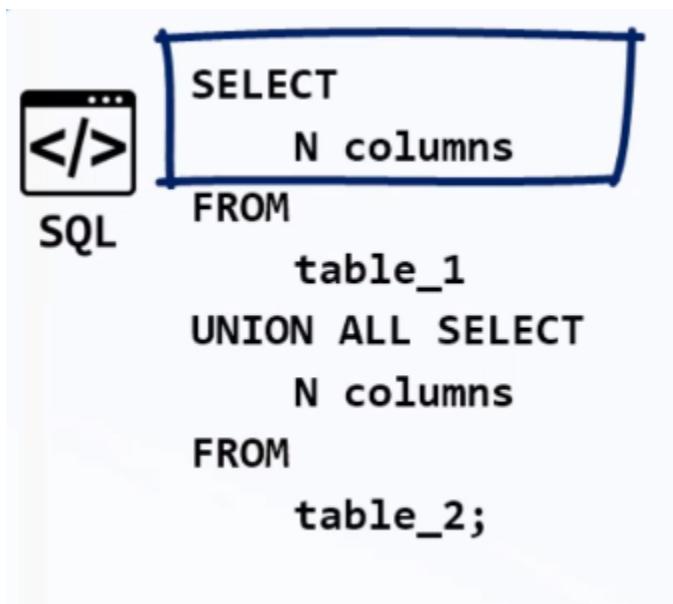
INSERT INTO employees_dup
SELECT
 e.*
```

```
FROM employees e
LIMIT 20;
```

```
INSERT INTO employees_dup VALUES
("10001","1953-09-02","Georgi","Facello","M","1986-06-26")
```

#### UNION ALL

- Used to combine a few SELECT statements in a single statement
- It is a tool to UNIFY tables



**Conditions: the same number of columns from each table; these columns should have the same name, same order, and related data types. If there are columns missing, you need to create NULL ones (NULL AS)**

```
SELECT
 e.emp_no,
 e.first_name,
 e.last_name,
 NULL AS dept_no,
 NULL AS from_date
FROM
 employees_dup e
WHERE
 e.emp_no = 10001
UNION ALL
SELECT
 NULL AS emp_no,
```

```

NULL AS first_name,
NULL AS last_name,
m.dept_no,
m.from_date
FROM
 dept_manager m;

```

|   | emp_no | first_name | last_name | dept_no | from_date  |
|---|--------|------------|-----------|---------|------------|
| ▶ | 10001  | Georgi     | Facello   | NULL    | NULL       |
|   | 10001  | Georgi     | Facello   | NULL    | NULL       |
|   | NULL   | NULL       | NULL      | d001    | 1985-01-01 |
|   | NULL   | NULL       | NULL      | d001    | 1991-10-01 |
|   | NULL   | NULL       | NULL      | d002    | 1985-01-01 |
|   | NULL   | NULL       | NULL      | d002    | 1989-12-17 |
|   | NULL   | NULL       | NULL      | d003    | 1985-01-01 |
|   | NULL   | NULL       | NULL      | d003    | 1992-03-21 |
|   | NULL   | NULL       | NULL      | d004    | 1985-01-01 |

The result is merge of the two tables

### UNION vs. UNION ALL

- When uniting two identically organized tables
  - UNION displays only distinct values in the output (no repetition)
  - UNION ALL retrieves duplicates as well (including possible repetition)
- Therefore UNION uses more resources (computing power and storage space) → better results; UNION ALL takes less → faster

Exercise:

Go forward to the solution and execute the query. What do you think is the meaning of the minus sign before subset A in the last row (ORDER BY -a.emp\_no DESC)?

SELECT

\*

FROM

(SELECT

e.emp\_no,

e.first\_name,

```
e.last_name,
NULL AS dept_no,
NULL AS from_date

FROM

employees e

WHERE

last_name = 'Denis' UNION SELECT

NULL AS emp_no,
NULL AS first_name,
NULL AS last_name,
dm.dept_no,
dm.from_date

FROM

dept_manager dm) as a

ORDER BY -a.emp_no DESC;
```

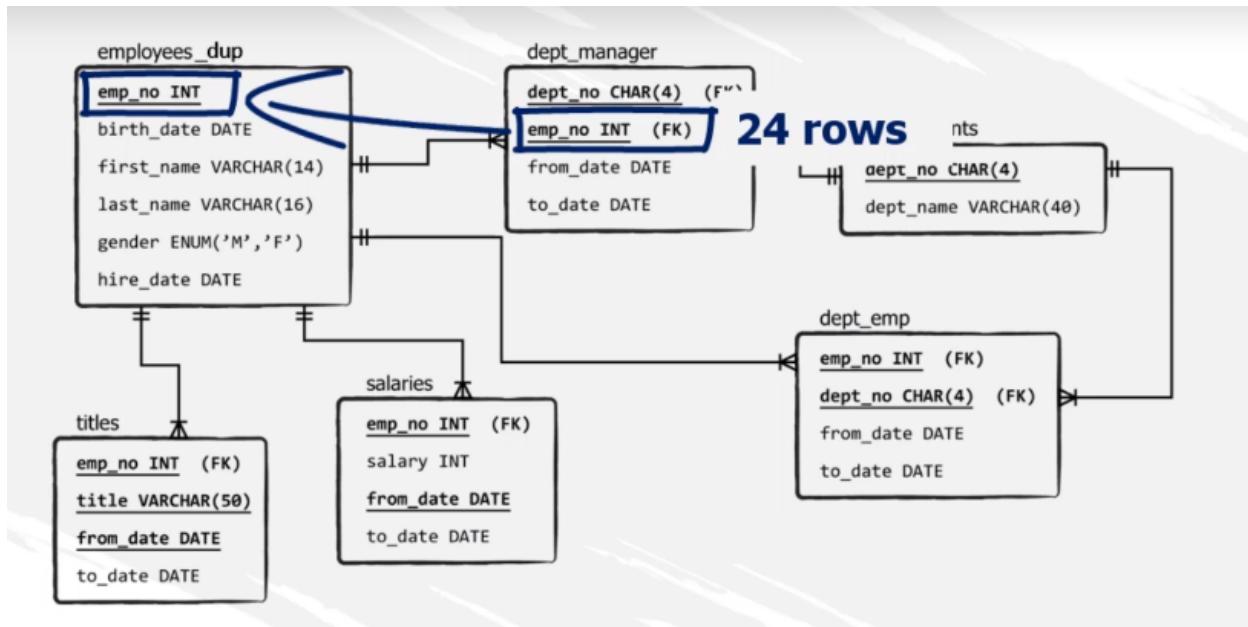
I think minus means ordering while putting all NULLs in the end???

## Section 14. SQL subqueries

### 78. SQL subqueries with IN nested inside where

Sub-queries: a query within another query (inner/nested query = inner select)  
If they are part of another query, called an outer query (= outer select)

**Example: the same 24 people in dept\_manager needs to be found in employees**



```

SELECT
 e.first_name, e.last_name
FROM
 employees e
WHERE
 e.emp_no IN (SELECT
 dm.emp_no
 FROM
 dept_manager dm);

```

```

7 # employee numbers that can be found in the "Department Manager" t
8 • SELECT
9 e.first_name, e.last_name
10 FROM
11 employees e
12 WHERE
13 e.emp_no IN (SELECT
14 dm.emp_no
15 FROM
16 dept_manager dm); T

```

outer query

```

SELECT
 e.emp_no IN (SELECT
 dm.emp_no
 FROM
 dept_manager dm); I

```

subquery (inner query)

- The SQL engine starts by running the inner query
- Then it uses its returned output, which is intermediate, to execute the outer query

A subquery may return a single value (scalar), a single row, a single column, or an entire table

**Exercise:** Extract the information about all department managers who were hired between the 1st of January 1990 and the 1st of January 1995.

Hire\_date is not start\_date. The former can only be found in the employees table.

```

dm.*
FROM
 dept_manager dm
WHERE
 dm.emp_no IN(SELECT
 e.emp_no
 FROM
 employees e
 WHERE
 e.hire_date BETWEEN "1990-01-01" AND "1995-01-01");

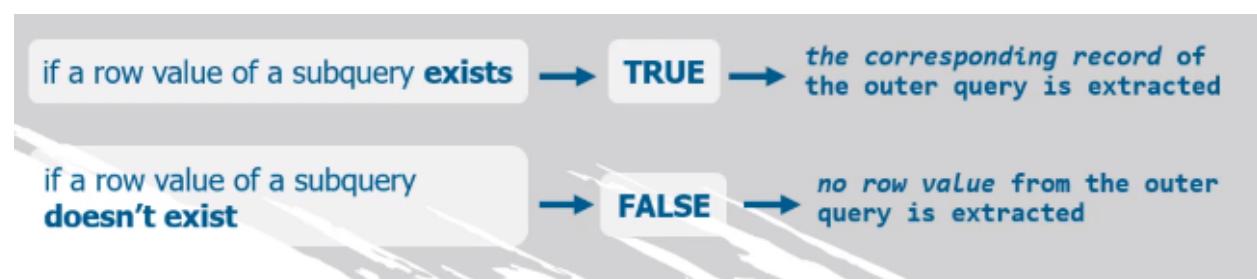
```

## 79. Subqueries with EXIST-NOT EXISTS Nested inside WHERE

**EXISTS**

Checks whether certain row values are found within a subquery

- This check is conducted row by row
- It returns a boolean value



Example: return anyone's first and last names if found in the manager table

```

SELECT
 e.first_name, e.last_name
FROM
 employees e
WHERE
 EXISTS(SELECT
 *
 FROM
 dept_manager dm
 WHERE
 dm.emp_no = e.emp_no
);

```

A table is within the WHERE EXISTS

The differences between EXISTS and IN

| EXISTS                                            | IN                                |
|---------------------------------------------------|-----------------------------------|
| Tests row values for existence                    | Searches among values             |
| Quicker in retrieving <u>large</u> amount of data | Faster in <u>smaller</u> datasets |

Tips:

- (1) Put ORDER BY outside the subquery for clarity

- **SELECT**

```

e.first_name, e.last_name
FROM
 employees e
WHERE
 EXISTS(SELECT
 *
 FROM
 dept_manager dm
 WHERE
 dm.emp_no = e.emp_no
)

```

**ORDER BY emp\_no;**

- (2) Some nested queries can be rewritten using JOINS (esp. Inner queries using ), which is more efficient

- **subqueries:**

- allow for better structuring of the outer query
  - thus, each inner query can be thought of in isolation
  - hence the name of SQL - Structured Query Language!
- in some situations, the use of subqueries is much more intuitive compared to the use of complex joins and unions
- many users prefer subqueries simply because they offer enhanced code readability

**Exercise:**

Select the entire information for all employees whose job title is “Assistant Engineer”.

Hint: To solve this exercise, use the 'employees' table.

Entire information in employees →

```
SELECT
 e.*
FROM
 employees e
WHERE
 EXISTS(SELECT
 *
 FROM
 titles t
 WHERE
 t.emp_no = e.emp_no AND title = "Assistant Engineer"
)
ORDER BY emp_no;
```

## In this lecture:



challenging  
task



exercise

**Task: assign employee number 110022 as a manager to all employees from 10001 to 10020, and employee number 110039 as a manager to all employees from 10021 to 10040.**

### employees

```
emp_no INT
birth_date DATE
first_name VARCHAR(14)
last_name VARCHAR(16)
gender ENUM('M','F')
hire_date DATE
```

### dept\_manager

```
dept_no CHAR(4) (FK)
emp_no INT (FK)
from_date DATE
to_date DATE
```

Subset A: 10001 to 10020

Subset B: 10021 to 10040

SA  
UNION  
SB

Step1: add subquery to SA based on the statement of 110022 as a manager

```
1 • SELECT
2 emp_no
3 FROM
4 dept_manager
5 WHERE
6 emp_no = 110022;
```

|   | employee_ID | dept_code | manager_ID |
|---|-------------|-----------|------------|
| I | < 10020     |           |            |
|   | < 10020     |           |            |

```

SELECT
e.emp_no AS employee_ID,
MIN(de.emp_no) AS department_code,

(SELECT
 emp_no
FROM
 dept_manager
WHERE
 emp_no = 110022) AS manager_ID

FROM
employees e
JOIN
dept_emp de ON e.emp_no = de.emp_no
WHERE e.emp_no <= 10020
GROUP BY e.emp_no
ORDER BY e.emp_no;
;

```

**Thought process: determine what columns to display + determine what source tables (or any JOINs/EXISTS/INs)**

Step2: wrap up the Subset A

```

SELECT
A.*
FROM
(SELECT
 e.emp_no AS employee_ID,
 MIN(de.emp_no) AS department_code,

(SELECT
 emp_no
FROM
 dept_manager
WHERE
 emp_no = 110022) AS manager_ID

FROM
employees e
JOIN
dept_emp de ON e.emp_no = de.emp_no

```

```
WHERE e.emp_no <= 10020
GROUP BY e.emp_no
ORDER BY e.emp_no) AS A;
```

Step3: UNION + three changes in the original A code

```
SELECT
 A.*
FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.emp_no) AS department_code,

 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110022) AS manager_ID

 FROM
 employees e
 JOIN
 dept_emp de ON e.emp_no = de.emp_no
 WHERE e.emp_no <= 10020
 GROUP BY e.emp_no
 ORDER BY e.emp_no) AS A
UNION
SELECT
 B.*
FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.emp_no) AS department_code,

 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110039) AS manager_ID

 FROM
```

```
employees e
JOIN
dept_emp de ON e.emp_no = de.emp_no
WHERE e.emp_no > 10020
GROUP BY e.emp_no
ORDER BY e.emp_no
LIMIT 20) AS B;
```

**Exercise 1:**

Starting your code with “DROP TABLE”, create a table called “emp\_manager” (emp\_no – integer of 11, not null; dept\_no – CHAR of 4, null; manager\_no – integer of 11, not null).

```
DROP TABLE IF EXISTS emp_manager;
CREATE TABLE emp_manager(
 emp_no INT(11) NOT NULL,
 dept_no CHAR(4) NULL,
 manager_no INT(11) NOT NULL
);
```

**Exercise 2:**

Fill emp\_manager with data about employees, the number of the department they are working in, and their managers.

Your query skeleton must be:

**Insert INTO emp\_manager SELECT**

**U.\***

**FROM**

**(A)**

**UNION (B) UNION (C) UNION (D) AS U;**

A and B should be the same subsets used in the last lecture (SQL Subqueries Nested in SELECT and FROM). In other words, assign employee number 110022 as a manager to all employees from 10001 to 10020 (this must be subset A), and employee number 110039 as a manager to all employees from 10021 to 10040 (this must be subset B).

**Use the structure of subset A to create subset C, where you must assign employee number 110039 as a manager to employee 110022.**

**Following the same logic, create subset D. Here you must do the opposite - assign employee 110022 as a manager to employee 110039.**

**Your output must contain 42 rows.**

**Good luck!**

```
INSERT INTO emp_manager
SELECT
 u.*
FROM
 (SELECT
 a.*
 FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.dept_no) AS department_code,
 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110022) AS manager_ID
 FROM
 employees e
 JOIN dept_emp de ON e.emp_no = de.emp_no
 WHERE
 e.emp_no <= 10020
 GROUP BY e.emp_no
 ORDER BY e.emp_no) AS a UNION SELECT
 b.*
 FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.dept_no) AS department_code,
 (SELECT
 emp_no
 FROM
 dept_manager
```

```

 WHERE
 emp_no = 110039) AS manager_ID
 FROM
 employees e
 JOIN dept_emp de ON e.emp_no = de.emp_no
 WHERE
 e.emp_no > 10020
 GROUP BY e.emp_no
 ORDER BY e.emp_no
 LIMIT 20) AS b UNION SELECT
 c.*
 FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.dept_no) AS department_code,
 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110039) AS manager_ID
 FROM
 employees e
 JOIN dept_emp de ON e.emp_no = de.emp_no
 WHERE
 e.emp_no = 110022
 GROUP BY e.emp_no) AS c UNION SELECT
 d.*
 FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.dept_no) AS department_code,
 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110022) AS manager_ID
 FROM
 employees e
 JOIN dept_emp de ON e.emp_no = de.emp_no
 WHERE
 e.emp_no = 110039
 GROUP BY e.emp_no) AS d) as u;

```

**Each block is like this:**

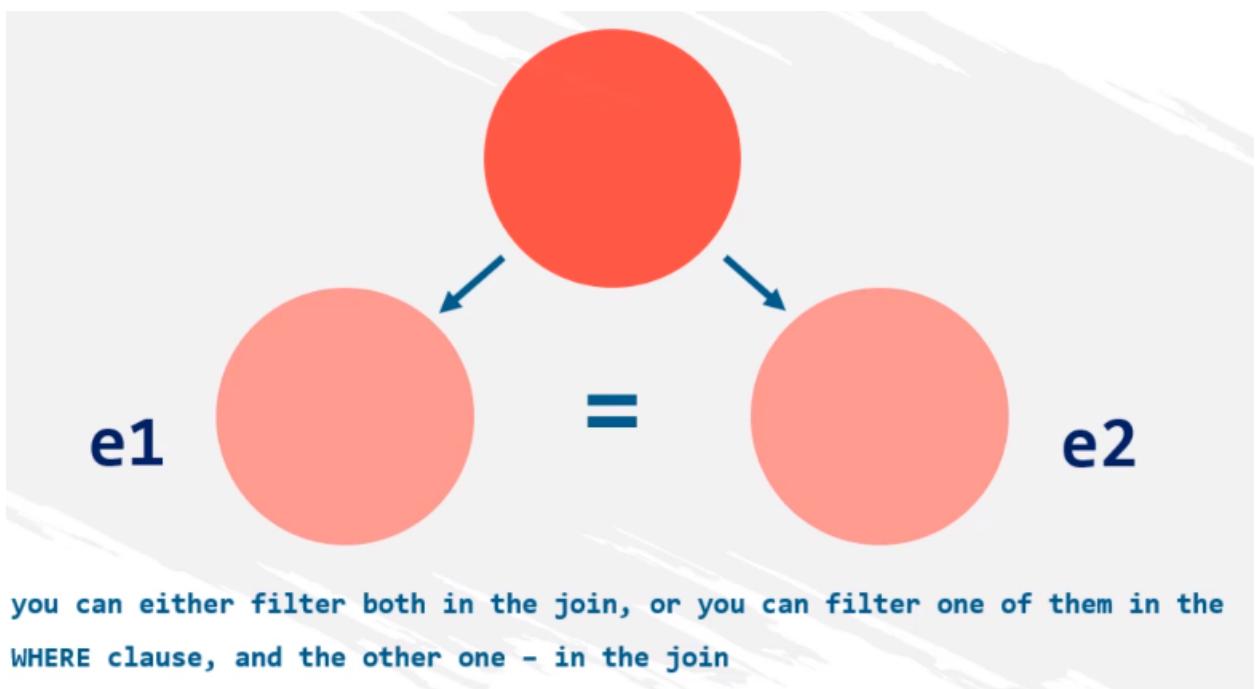
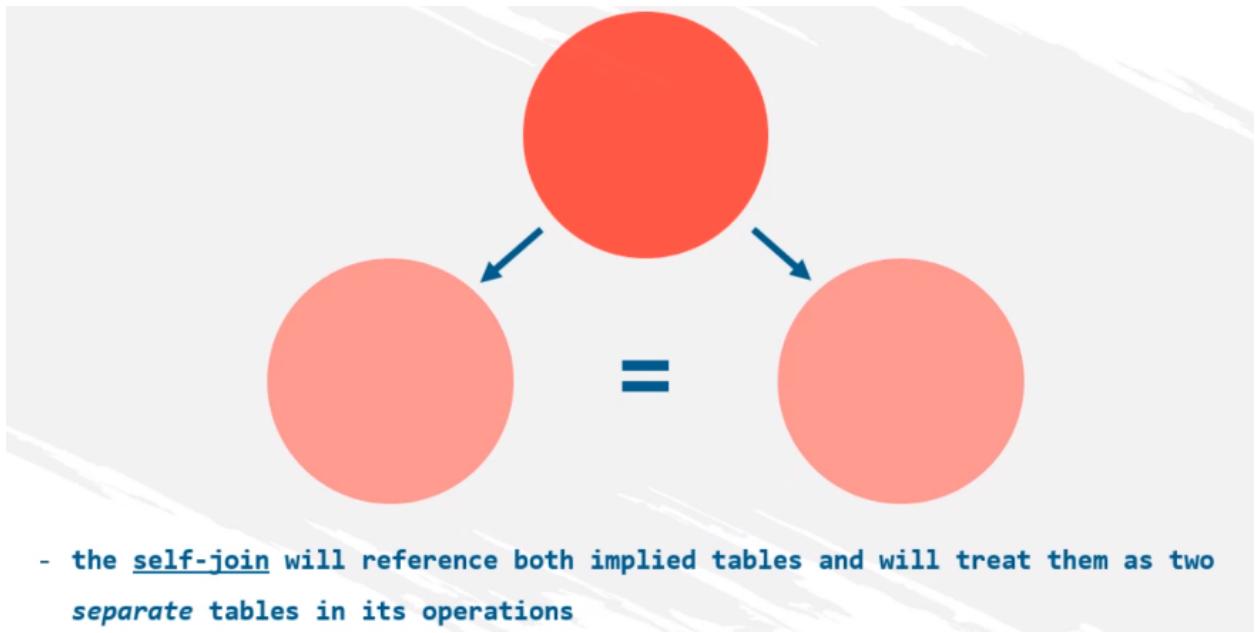
```
SELECT
 A.*
FROM
 (SELECT
 e.emp_no AS employee_ID,
 MIN(de.emp_no) AS department_code,
 (SELECT
 emp_no
 FROM
 dept_manager
 WHERE
 emp_no = 110022) AS manager_ID
 FROM
 employees e
 JOIN
 dept_emp de ON e.emp_no = de.emp_no
 WHERE e.emp_no <= 10020
 GROUP BY e.emp_no
 ORDER BY e.emp_no) AS A
```

## Section 15. SQL Self Join

### 81. Self-join

Applied when a table must join itself

- If you'd like to combine certain rows of a table with other rows of the same table, you need a self-join



Two cases: filtering both or filtering one

Example: From the emp\_manager table, extract the record data only of those employees who are managers as well.

```
SELECT
 e1.*
```

```

FROM
 emp_manager e1
 JOIN
 emp_manager e2 ON e1.emp_no = e2.manger_no;

```

The screenshot shows the SQL Server Management Studio interface with two tables, **e1** and **e2**, displayed in separate panes. Both tables have columns: **emp\_no**, **dept\_no**, and **manager\_no**. The data in both tables is identical, showing various employee numbers, department numbers, and manager numbers. A blue box highlights the first row of table **e1** (emp\_no 10017, dept\_no d001, manager\_no 110022). A red box highlights the last row of table **e1** (emp\_no 110039, dept\_no d001, manager\_no 110022). A blue arrow points from the highlighted row in **e1** to the corresponding row in the result grid below. The result grid shows the joined data with the same two rows highlighted. The SQL query in the query editor is:

```

SELECT
 e1.*
FROM
 emp_manager e1
 JOIN
 emp_manager e2 ON e1.emp_no = e2.manger_no;

```

The status bar at the bottom indicates the execution time and number of rows returned.

How do we delete repeated data?

(1)

**SELECT DISTINCT**

e1.\*

FROM

emp\_manager e1

JOIN

emp\_manager e2 ON e1.emp\_no = e2.manger\_no;

(2)

**SELECT**

e1.\*

FROM

emp\_manager e1

JOIN

emp\_manager e2 ON e1.emp\_no = e2.manger\_no

```

WHERE
 e2.emp_no IN(SELECT
 manger_no
 FROM
 emp_manager
);

```

## Section 16. SQL Views

### 82. View

View is a virtual table whose contents are obtained from an existing table or tables called base tables.

The retrieval happens through an SQL statement, incorporating into the view.

The view does not contain any real data; the data is physically stored in the base table; the view simply shows the data contained in the base table.

Example: emp+manager

```

SELECT
 emp_no, from_date, to_date, COUNT(emp_no) AS Num
FROM
 dept_emp
GROUP BY emp_no
HAVING Num >1;

```

The screenshot shows a database interface with a results grid and an action output window. The results grid contains the following data:

| emp_no | dept_no | from_date  | to_date    |
|--------|---------|------------|------------|
| 10001  | d005    | 1986-06-26 | 9999-01-01 |
| 10002  | d007    | 1986-08-03 | 9999-01-01 |
| 10003  | d004    | 1995-12-03 | 9999-01-01 |
| 10004  | d004    | 1986-12-01 | 9999-01-01 |
| 10005  | d003    | 1989-09-12 | 9999-01-01 |
| 10006  | d005    | 1990-08-05 | 9999-01-01 |
| 10007  | d008    | 1989-02-10 | 9999-01-01 |
| 10008  | d005    | 1989-03-11 | 2000-07-31 |
| 10009  | d006    | 1985-02-18 | 9999-01-01 |
| 10010  | d004    | 1996-11-24 | 2000-06-26 |
| 10011  | d006    | 2000-06-26 | 9999-01-01 |

Action Output:

1 17:07:51 SELECT \* FROM dept\_emp

Message: 331603 row(s) returned

The output's number is very high because some employees' history was repeated. Our goal here is to visualize only the period encompassing the last contract of each employee.

The screenshot shows a SQL editor window with a toolbar at the top. The main area contains a template for creating a view:

```
CREATE VIEW view_name AS
SELECT
 column_1, column_2, ..., column_n
FROM
 table_name;
```

Below the template, there is a status bar with some text and numbers.

```
CREATE VIEW v_dept_emp_latest_date AS
SELECT
 emp_no, MAX(from_date) AS from_date, MAX(to_date) AS to_date
FROM
 dept_emp
GROUP BY emp_no;
```

After execution, this view can be found (1) on the schema:

The screenshot shows a database schema browser. On the left, there is a tree view of the schema:

- titles
- Views
  - current\_dept\_emp
  - dept\_emp\_latest\_da
  - v\_dept\_emp\_late
- Stored Procedures
- Functions

Below the tree, there are two database objects:

- sales
- sys

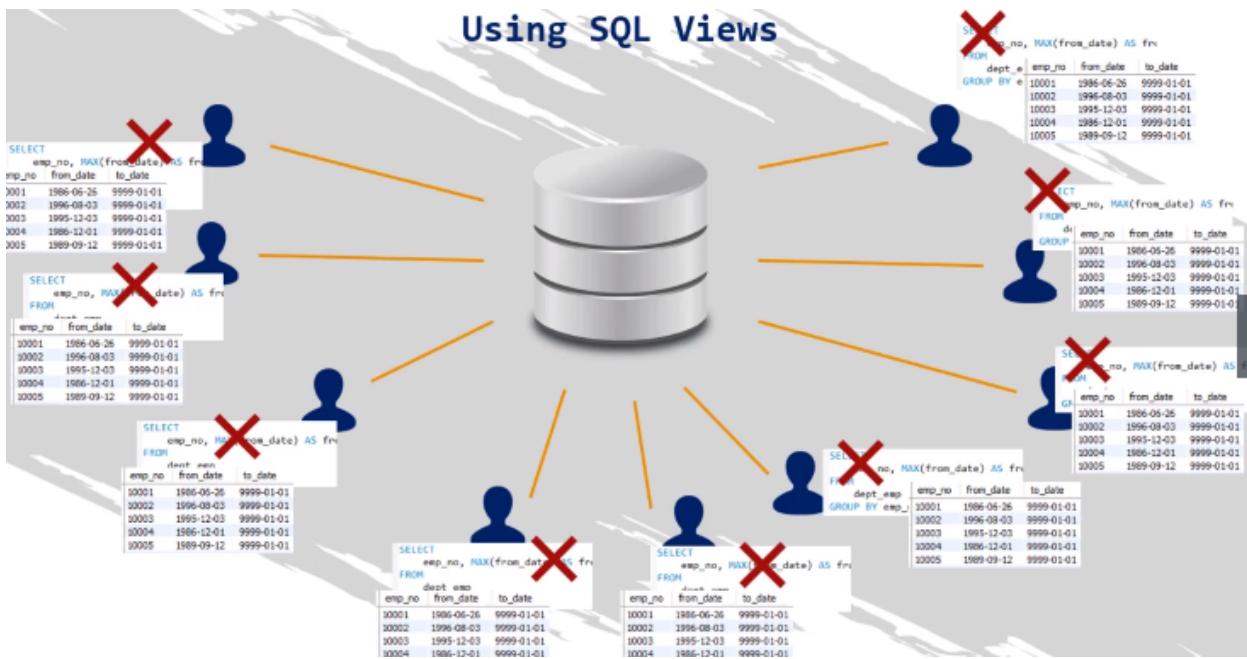
On the right, there is a result grid titled "Result Grid" showing the data from the view:

|   | emp_no | from_date  | to_date    |
|---|--------|------------|------------|
| ▶ | 10001  | 1986-06-26 | 9999-01-01 |
|   | 10002  | 1996-08-03 | 9999-01-01 |
|   | 10003  | 1995-12-03 | 9999-01-01 |
|   | 10004  | 1986-12-01 | 9999-01-01 |
|   | 10005  | 1989-09-12 | 9999-01-01 |

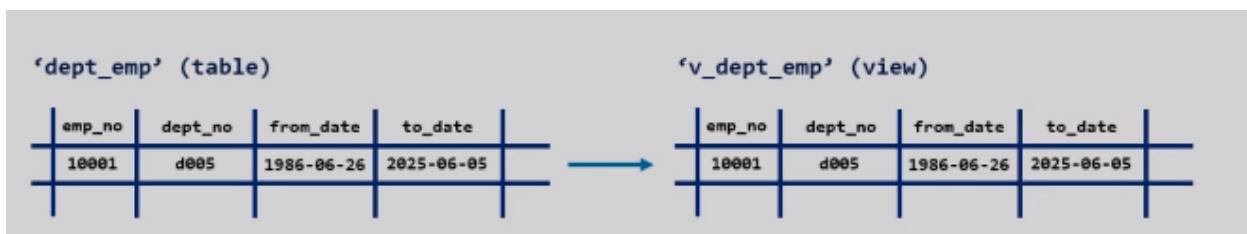
(2) Or use this SELECT line:

```
SELECT * FROM employees.v_dept_emp_latest_date;
```

## Using SQL Views



A view acts as a shortcut for writing the same SELECT statement every time a new user wants to check the table. Therefore it saves a lot of coding time and storage; any change in the table can be synced in the view.



### Exercise:

Create a view that will extract the average salary of all managers registered in the database. Round this value to the nearest cent.

If you have worked correctly, after executing the view from the “Schemas” section in Workbench, you should obtain the value of 66924.27.

```
CREATE VIEW avg_salary_manager AS
 SELECT
 AVG(s.salary)
 FROM
 salaries s
 JOIN
 dept_manager dm ON s.emp_no=dm.emp_no;
```

(Then use the schema-views on the left)

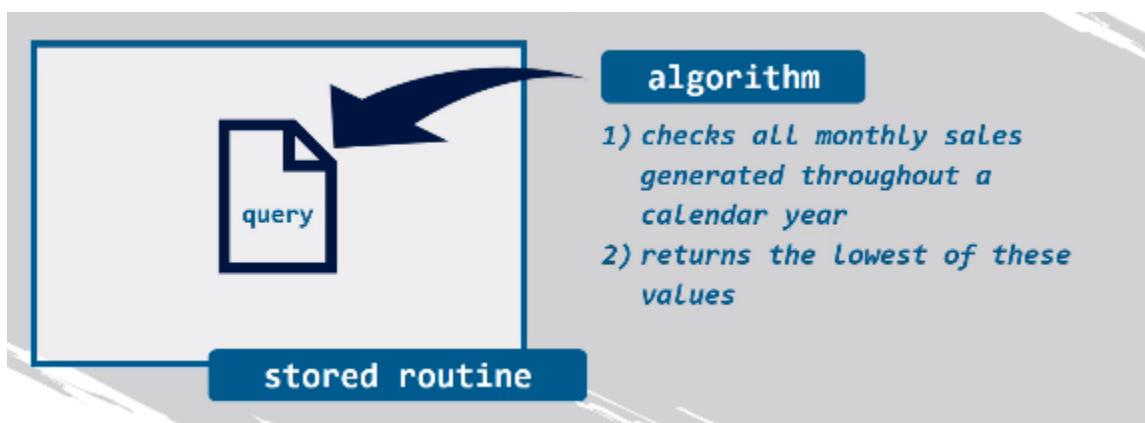
## Section 17. Stored routines

### 83. Introduction

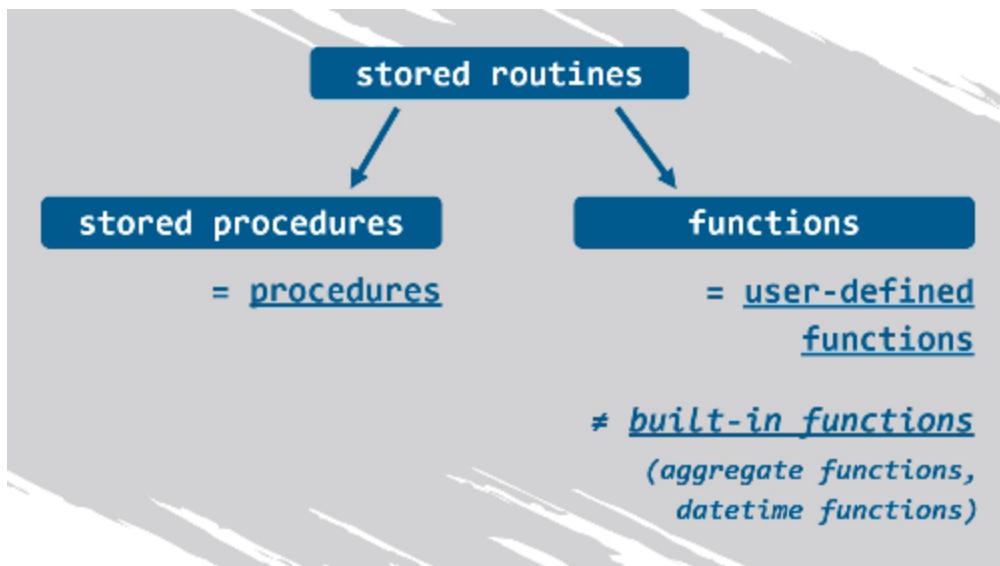
Stored routine: an SQL statement, or a set of SQL statements, that can be stored on the database server

Whenever a user wants to run the query, they can call, reference, or invoke the routine.

Example:

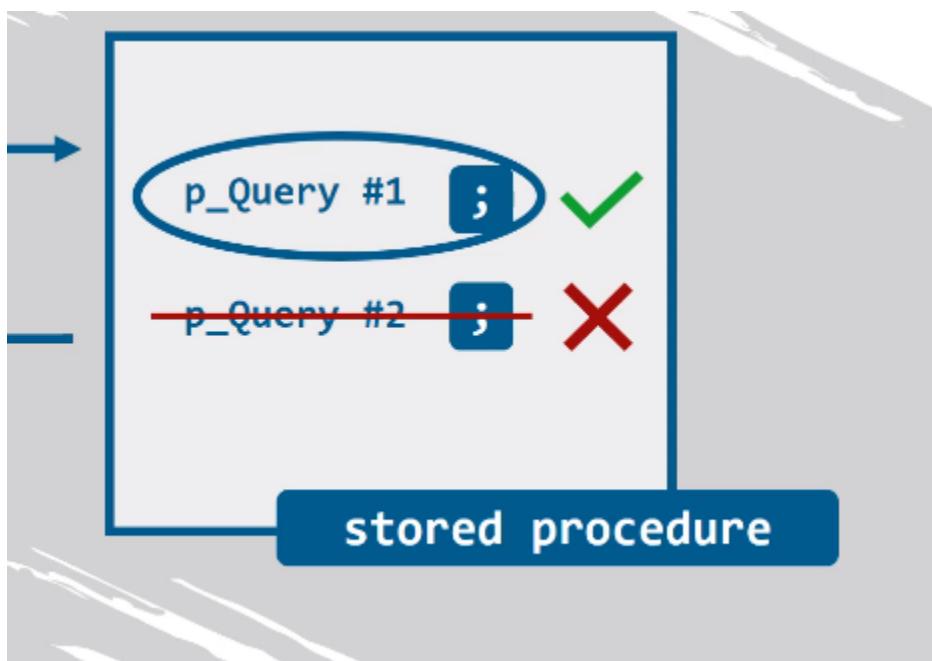


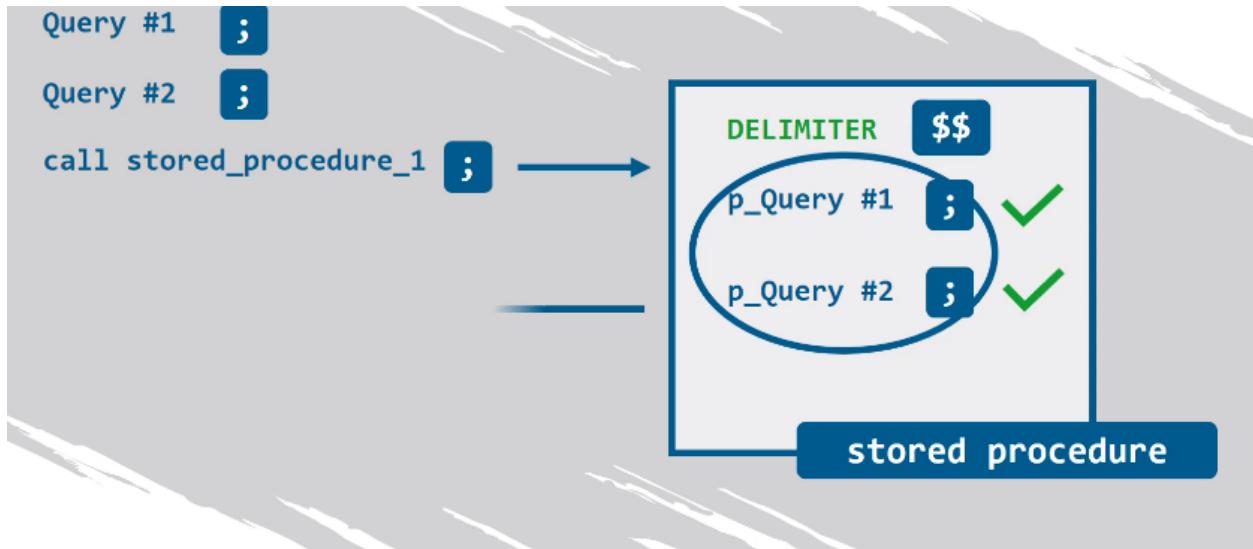
This statement is used by tons of users, so making it a routine is necessary.



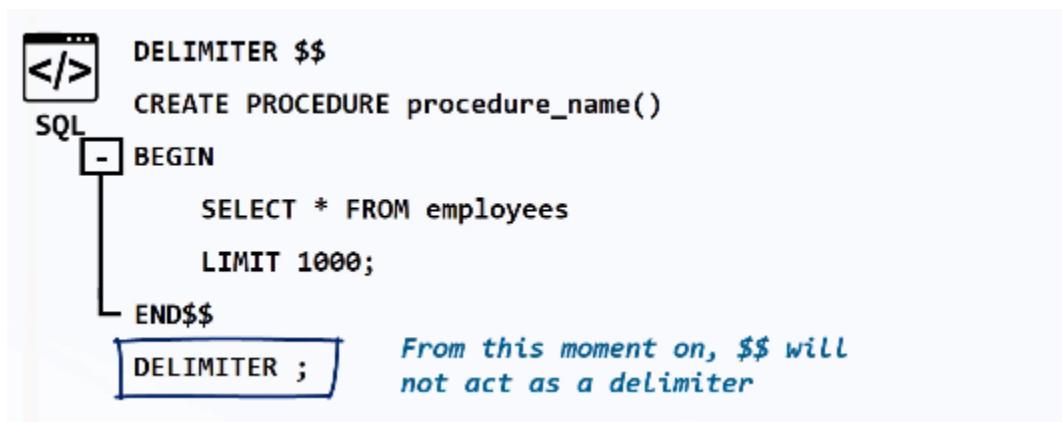
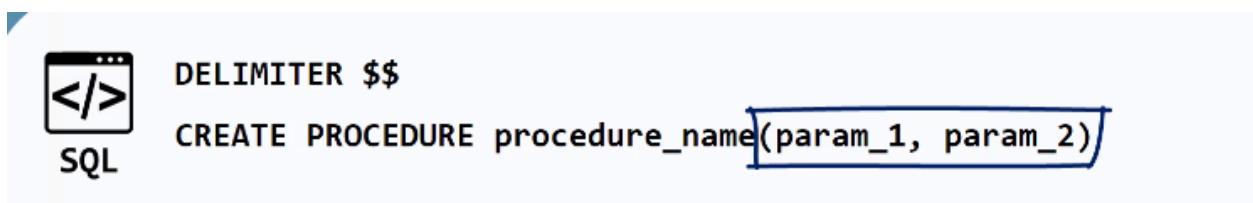
#### 84. MySQL syntax for stored procedures

**Use \$\$ or // as a delimiter, so the queries before can be stored.**





Syntax of stored procedure



Example: select the first 1000 tokens from employees

USE employees;

DROP PROCEDURE IF EXISTS select\_employees;

```

DELIMITER $$

CREATE PROCEDURE select_employees()
BEGIN
 SELECT * FROM employees
 LIMIT 1000;
END$$

DELIMITER ;

```

85. Check the procedure you create

**There are two ways:**

(1) Call method

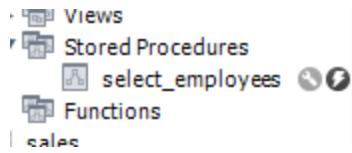
The screenshot shows the MySQL Workbench interface. In the top query editor window, the command `CALL database_name.procedure_name();` is entered. In the bottom result grid, the output of the procedure is displayed as a table with columns: emp\_no, birth\_date, first\_name, last\_name, gender, and hire\_date. The data consists of five rows of employee information.

|   | emp_no | birth_date | first_name | last_name | gender | hire_date  |
|---|--------|------------|------------|-----------|--------|------------|
| ▶ | 10001  | 1953-09-02 | Georgi     | Facello   | M      | 1986-06-26 |
|   | 10002  | 1964-06-02 | Bezalel    | Simmel    | F      | 1985-11-21 |
|   | 10003  | 1959-12-03 | Parto      | Bamford   | M      | 1986-08-28 |
|   | 10004  | 1954-05-01 | Chirstian  | Koblick   | M      | 1986-12-01 |
|   | 10005  | 1955-01-21 | Kyoichi    | Maliniak  | M      | 1989-09-12 |

Since we have “`USE database_name`” at the beginning, there is a default database. So there is no need to assign the name in the call.

```
call select_employees();
```

(2) Press the lightning button in the schema



By clicking the wrench button, we can view and edit the procedure code.

A screenshot of the MySQL Workbench 'Routine Editor' window. The title bar says 'exercise\*' and 'select\_employees - Routine'. The 'Name:' field contains 'select\_employees'. Below it, the 'DDL:' section shows the SQL code for the stored procedure:

```
1 • CREATE DEFINER='root'@'localhost' PROCEDURE `select_employees`()
2 BEGIN
3 SELECT * FROM employees
4 LIMIT 1000;
5 END
```

**Exercise: Create a procedure that will provide the average salary of all employees.**

**Then, call the procedure.**

```
DROP PROCEDURE IF EXISTS average_salary;
```

```
DELIMITER $$
CREATE PROCEDURE average_salary()
BEGIN
 SELECT AVG(salary) FROM salaries;
END$$
```

```
DELIMITER ;
```

```
call average_salary();
```

Some other ways to call the procedure...

```
CALL avg_salary;
```

```
CALL avg_salary();

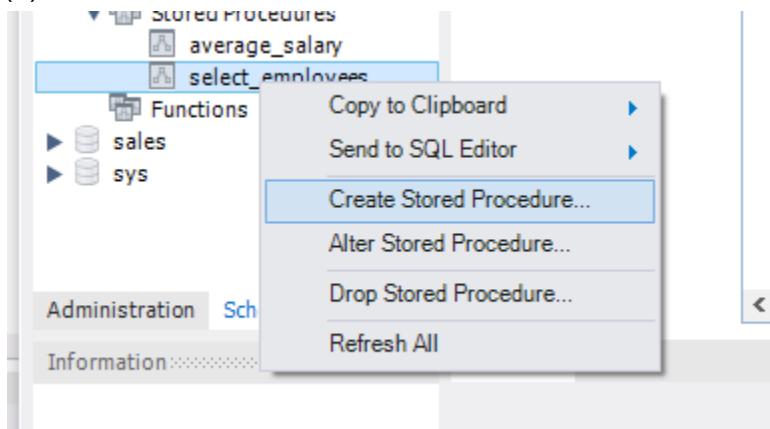
CALL employees.avg_salary;

CALL employees.avg_salary();
```

86. Another way to create a procedure

Click on create a procedure on the schema

(1)



(2)

ercise\* new\_procedure - Routine X

Name:  The name of the routine is parsed automatically from the DDL statement. The DDL is parsed automatically while you type.

DDL:

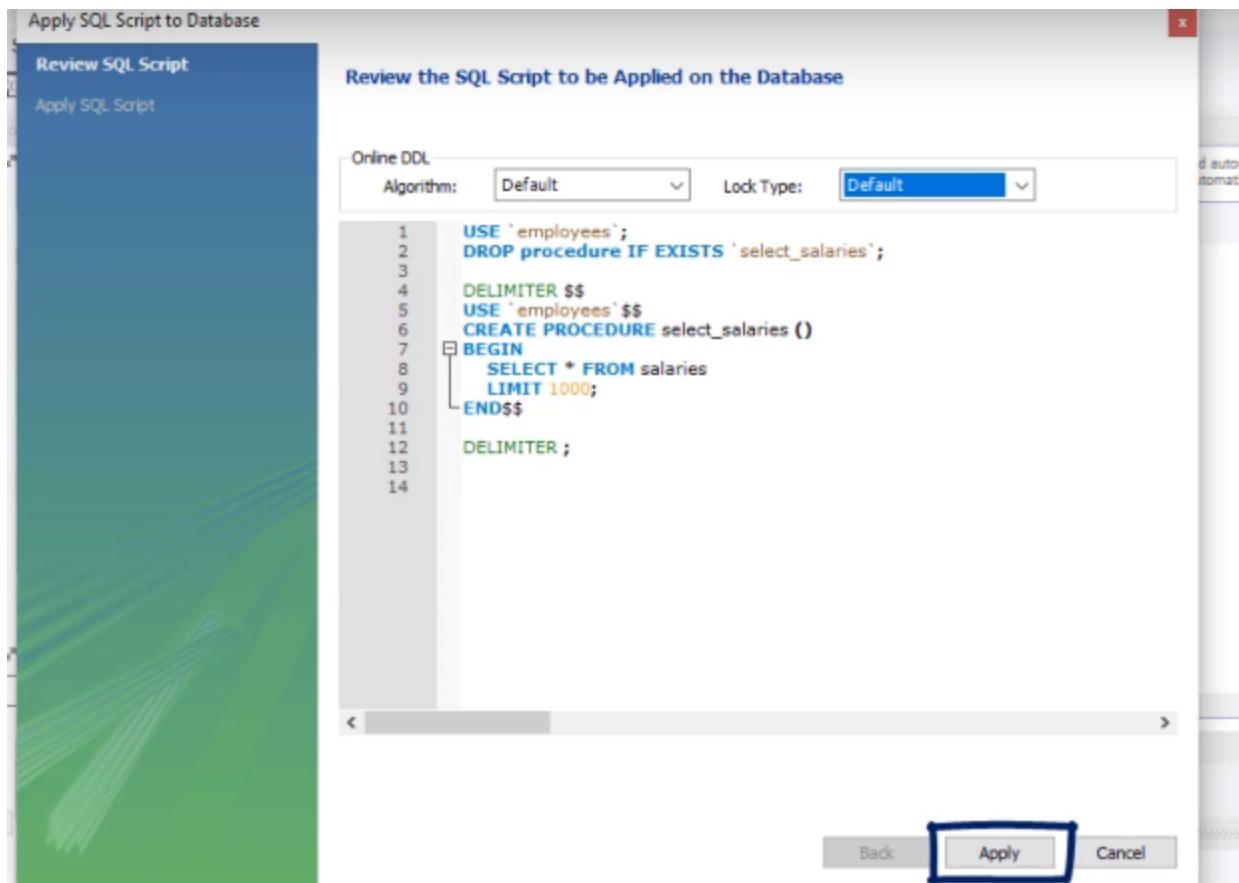
```
1 • CREATE PROCEDURE select_salaries ()
2 BEGIN
3 SELECT * FROM salaries
4 LIMIT 1000;
5 END
6
```

< >

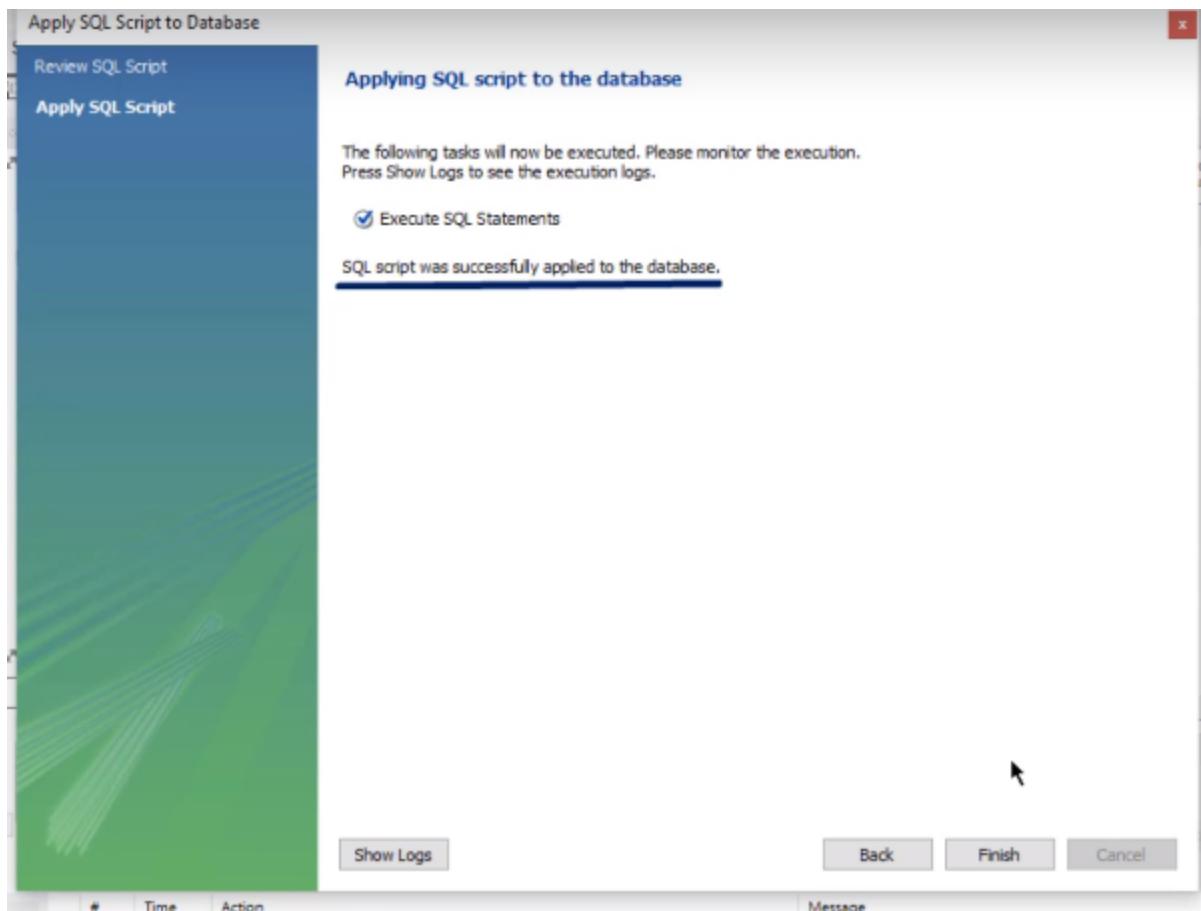
routine

Apply Revert

(3)



(4)



(5) run the procedure (on the left)

MANAGEMENT

- Server Status
- Client Connections
- Users and Privileges
- Status and System Variables
- Data Export
- Data Import/Restore

INSTANCE

- Startup / Shutdown
- Server Logs
- Options File

PERFORMANCE

- Dashboard
- Performance Reports
- Performance Schema Setup

SCHEMAS

- employees
- Tables
- Views
- Stored Procedures
- select\_employees
- select\_salaries
- Functions

Information

Procedure: select\_salaries

Result Grid

| emp_no | salary | from_date  | to_date    |
|--------|--------|------------|------------|
| 10001  | 60117  | 1986-06-26 | 1987-06-26 |
| 10001  | 62102  | 1987-06-26 | 1988-06-25 |
| 10001  | 66074  | 1988-06-25 | 1989-06-25 |
| 10001  | 66596  | 1989-06-25 | 1990-06-25 |
| 10001  | 66961  | 1990-06-25 | 1991-06-25 |

Action Output

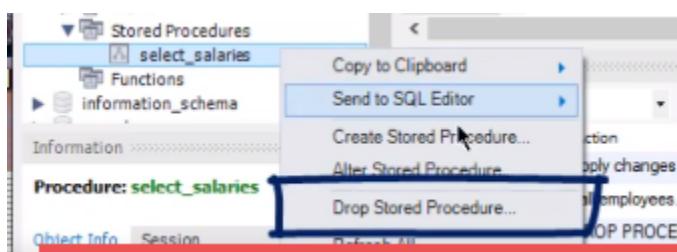
| # | Time     | Action                            | Message              |
|---|----------|-----------------------------------|----------------------|
| 5 | 14:58:13 | call employees.select_employees() | 1000 row(s) returned |
| 6 | 15:09:55 | Apply changes to select_salaries  | Changes applied      |

## (6) drop a procedure

A.

**DROP PROCEDURE procedure\_name;**

B. right click on the procedure in the scheme



### 86. Stored procedures with an input parameter

A stored routine can perform a calculation that transforms an input value in an output value

Stored procedures can take an input value and then use it in the query, or queries, written in the body of the procedure

- This value is represented by the IN parameter
- After that calculation is ready, the result will be returned

(1) Syntax:

```

DELIMITER $$

CREATE PROCEDURE procedure_name(in parameter)
BEGIN
 SELECT...;
END$$

DELIMITER ;

```

---

```

DELIMITER $$ name data_type
• USE employees $$ name data_type
• CREATE PROCEDURE emp_salary(IN p_emp_no INTEGER)
 BEGIN
 SELECT
 e.first_name, e.last_name, s.salary, s.from_date, s.to_date
 FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no

```

'ER \$\$

Parameter contains both name and data type.

USE employees;

DROP PROCEDURE IF EXISTS emp\_salary;

DELIMITER \$\$

USE employees \$\$

CREATE PROCEDURE emp\_salary(**IN p\_emp\_no INTEGER**)

BEGIN

```

 SELECT
 e.first_name, e.last_name, s.salary, s.from_date, s.to_date
 FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no
WHERE
 e.emp_no = p_emp_no;

```

END\$\$

DELIMITER ;

Run the procedure on the left side and test with 11300

The screenshot shows the MySQL Workbench interface. On the left, the 'Schemas' tree is open, showing the 'employees' schema with its tables (departments, dept\_emp, dept\_manager, employees, salaries, titles), views, stored procedures (average\_salary, emp\_salary, select\_employees, select\_salaries), and functions. In the main pane, a query editor window is open with the following SQL code:

```
1 • call employees.emp_salary(11300);
2
```

A modal dialog titled 'Call stored procedure employees.emp\_salary' is displayed. It contains a message: 'Enter values for parameters of your procedure and click <Execute> to create an SQL editor and run the call:' Below this, there is a parameter input field for 'p\_emp\_no' with the value '11300' and a type indicator '(IN) INTEGER'. At the bottom right of the dialog are 'Execute' and 'Cancel' buttons. A result grid table is visible at the bottom of the dialog, showing the output of the procedure call.

(2) Procedures with one input parameter can be used with aggregate functions too

Example: calculate the average salary instead of several salaries  
→ change the procedure name and column name as AVG()

```
USE employees;
DROP PROCEDURE IF EXISTS emp_salary;
```

```
DELIMITER $$
USE employees $$
CREATE PROCEDURE emp_avg_salary(IN p_emp_no INTEGER)
BEGIN
 SELECT
 e.first_name, e.last_name, AVG(s.salary), s.from_date, s.to_date
 FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no
 WHERE
 e.emp_no = p_emp_no;
END$$
```

```
DELIMITER ;

call emp_avg_salary(11300);
```

## 87. Stored procedures with an output parameter

We can add two or more parameters

```
 SQL
DELIMITER $$
CREATE PROCEDURE procedure_name(in parameter, out parameter)
BEGIN
 SELECT * FROM employees
 LIMIT 1000;
END$$
DELIMITER ;
```

Out\_parameter: it will represent the variable containing the output value of the operation executed by the query of the stored procedure

Example: add one more parameter to the example in #86

```

1 • USE employees;
2 • DROP procedure IF EXISTS emp_avg_salary_out;
3
4 DELIMITER $$*
5 • CREATE PROCEDURE emp_avg_salary_out(in p_emp_no INTEGER, out p_avg_salary DECIMAL(10,2))
6 BEGIN
7 SELECT
8 AVG(s.salary)
9 INTO p_avg_salary
10 FROM
11 employees e
12 JOIN
13 salaries s ON e.emp_no = s.emp_no
14 WHERE
15 e.emp_no = p_emp_no;

```

USE employees;  
DROP PROCEDURE IF EXISTS emp\_avg\_salary\_out;

```

DELIMITER $$*
USE employees $$*
CREATE PROCEDURE emp_avg_salary_out(IN p_emp_no INTEGER, OUT p_avg_salary
DECIMAL(10,2))
BEGIN
 SELECT
 AVG(s.salary)
 INTO p_avg_salary
 FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no
 WHERE
 e.emp_no = p_emp_no;
END$$

```

DELIMITER ;

- **SELECT... INTO make the outside parameter not required!**

The screenshot shows a MySQL Workbench interface. The top window is titled 'exercise' and contains the following SQL code:

```
1 • set @p_avg_salary = 0;
2 • call employees.emp_avg_salary_out(11300, @p_avg_salary);
3 • select @p_avg_salary;
4
```

The bottom window is titled 'Result Grid' and displays the output of the last query:

|   | @p_avg_salary |
|---|---------------|
| ▶ | 48193.80      |

**Exercise:** Create a procedure called 'emp\_info' that uses as parameters the first and the last name of an individual, and returns their employee number.

```
USE employees;
DROP PROCEDURE IF EXISTS emp_info;

DELIMITER $$

USE employees $$

CREATE PROCEDURE emp_info(IN p_first_name CHAR(255), in p_last_name CHAR(255),
OUT p_emp_no INTEGER)
BEGIN
 SELECT
 e.emp_no
 INTO p_emp_no
 FROM
 employees e
```

```

 WHERE
 e.first_name = p_first_name AND e.last_name = p_last_name;
END$$

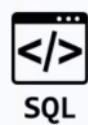
DELIMITER ;

```

## 88. Variables

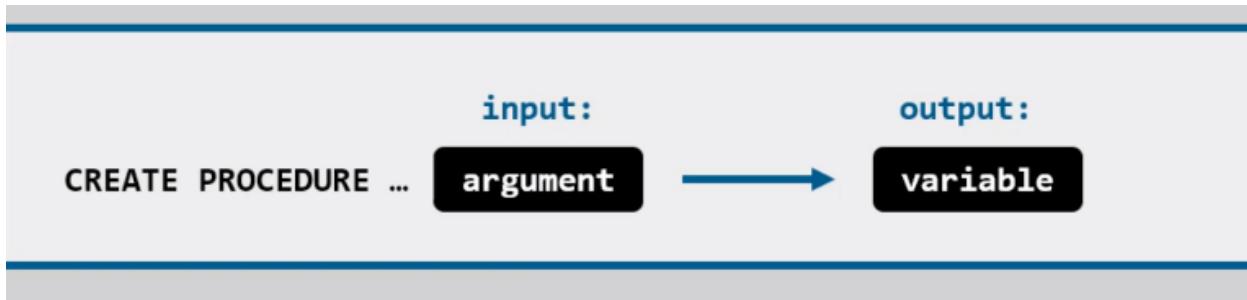
When you are defining a program, such as a stored procedure for instance, you can say you are using “parameters”

- “Parameters” are a more abstract term



**DELIMITER \$\$**  
**CREATE PROCEDURE procedure\_name (in parameter , out parameter )**

Once the structure has been solidified, then it will be applied to the database. The input you insert is typically referred to as the “argument”, while the output value is stored in a “variable”



**x=arg, y=variable**

Example:

- (1) Create a variable

```
SET @v_avg_salary = 0
```

- (2) Extract a value that will be assigned to the newly created variable (call the procedure)  
`CALL employees.emp_avg_salary_out(11300,@v_avg_salary);`

- (3) Ask the software to display the procedure

```
SELECT @v_avg_salary;
```

```
8 • SET @v_avg_salary = 0;
9
10 • CALL employees.emp_avg_salary_out(11300,@v_avg_salary);
11
12 • SELECT @v_avg_salary;
```

| Result Grid   |              |
|---------------|--------------|
|               | Filter Rows: |
| @v_avg_salary | 48193.80     |

IN-OUT parameters

input = output

CREATE PROCEDURE ... OUT parameter ←

**Exercise:** Create a variable, called 'v\_emp\_no', where you will store the output of the procedure you created in the last exercise.

Call the same procedure, inserting the values 'Aruna' and 'Journel' as a first and last name respectively.

Finally, select the obtained output.

```

SET @v_emp_no = 0;

CALL emp_info('Aruna', 'Journel', @v_emp_no);

SELECT @v_emp_no;

```

## 89. User-defined functions in MySQL

The code structure is similar to stored procedures.

```

DELIMITER $$

CREATE FUNCTION function_name(parameter data_type) RETURNS data_type
DECLARE variable_name data_type
BEGIN
 SELECT ...
 RETURN variable_name
END$$
DELIMITER ;

```

**Annotations:**

- An arrow points from the word "parameter" in the `CREATE FUNCTION` line to the text: *here you have no OUT parameters to define between the parentheses after the object's name*.
- An arrow points from the word "variable\_name" in the `DECLARE` line to the text: *all parameters are IN, and since this is well known, you need not explicitly indicate it with the word, 'IN'*.

Although there is no OUT parameters, there is a return value. The return value is obtained after running the query contained in the body of the function.

Example: write a user-defined function for average salary

```

DELIMITER $$

CREATE FUNCTION f_emp_avg_salary(p_emp_no INTEGER) RETURNS DECIMAL(10,2) #
this is just a returned value
DETERMINISTIC # must add this to avoid 1418 error; or add “DETERMINISTIC; NO SQL,
READS SQL DATA” (one of the words would be fine)

```

BEGIN

```

DECLARE v_avg_salary DECIMAL(10,2); # this is a step to DECLARE a returned value as
a variable

```

```

SELECT
 AVG(s.salary)
INTO v_avg_salary FROM
 employees e

```

```

JOIN
 salaries s ON e.emp_no=s.emp_no
WHERE e.emp_no = p_emp_no; #link the parameter with the table
RETURN v_avg_salary; # return the variable at the end
END $$

DELIMITER ;

```

Then use SELECT to call the function...

```
SELECT f_emp_avg_salary(11300)
```

More about the 1418 repair words:

- DETERMINISTIC – it states that the function will always return identical result given the same input
- NO SQL – means that the code in our function does not contain SQL (rarely the case)
- READS SQL DATA – this is usually when a simple SELECT statement is present

#### **Exercise:**

**Create a function called ‘emp\_info’ that takes for parameters the first and last name of an employee, and returns the salary from the newest contract of that employee.**

**Hint: In the BEGIN-END block of this program, you need to declare and use two variables – v\_max\_from\_date that will be of the DATE type, and v\_salary, that will be of the DECIMAL (10,2) type.**

**Finally, select this function.**

```

DELIMITER $$

CREATE FUNCTION emp_info(p_first_name varchar(255), p_last_name varchar(255))
RETURNS decimal(10,2)

DETERMINISTIC NO SQL READS SQL DATA

BEGIN
 DECLARE v_max_from_date date;
 DECLARE v_salary decimal(10,2);

```

```

SELECT
 MAX(from_date)

INTO v_max_from_date FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no

WHERE
 e.first_name = p_first_name
 AND e.last_name = p_last_name;

SELECT
 s.salary

INTO v_salary FROM
 employees e
 JOIN
 salaries s ON e.emp_no = s.emp_no

WHERE
 e.first_name = p_first_name
 AND e.last_name = p_last_name
 AND s.from_date = v_max_from_date;

RETURN v_salary;

END$$
DELIMITER ;

```

```
SELECT EMP_INFO('Aruna', 'Journe');
```

| Result Grid |                             |  |  | Filter Rows: |
|-------------|-----------------------------|--|--|--------------|
|             | EMP_INFO('Aruna', 'Journe') |  |  |              |
| ▶           | 45709.00                    |  |  |              |

### Technical differences between procedures and functions

| Stored procedure                                      | User-defined function                              |
|-------------------------------------------------------|----------------------------------------------------|
| Does not return a value                               | Returns a value                                    |
| CALL procedure                                        | SELECT procedure                                   |
| Cannot be embedded in a SELECT (in the example below) | Can be embedded in a SELECT (in the example below) |

### Conceptual differences

| Stored procedure                       | User-defined function          |
|----------------------------------------|--------------------------------|
| Can Have multiple OUT parameters       | Can ONLY return a single value |
| INSERT, UPDATE, and DELETE are allowed | No INSERT, UPDATE, or DELETE   |

- If you need to obtain more than one value as a result of a calculation, you are better off using a procedure
- **If you need to have just one value to be returned, then you can use a function**

### Function can be combined with a traditional SELECT statement

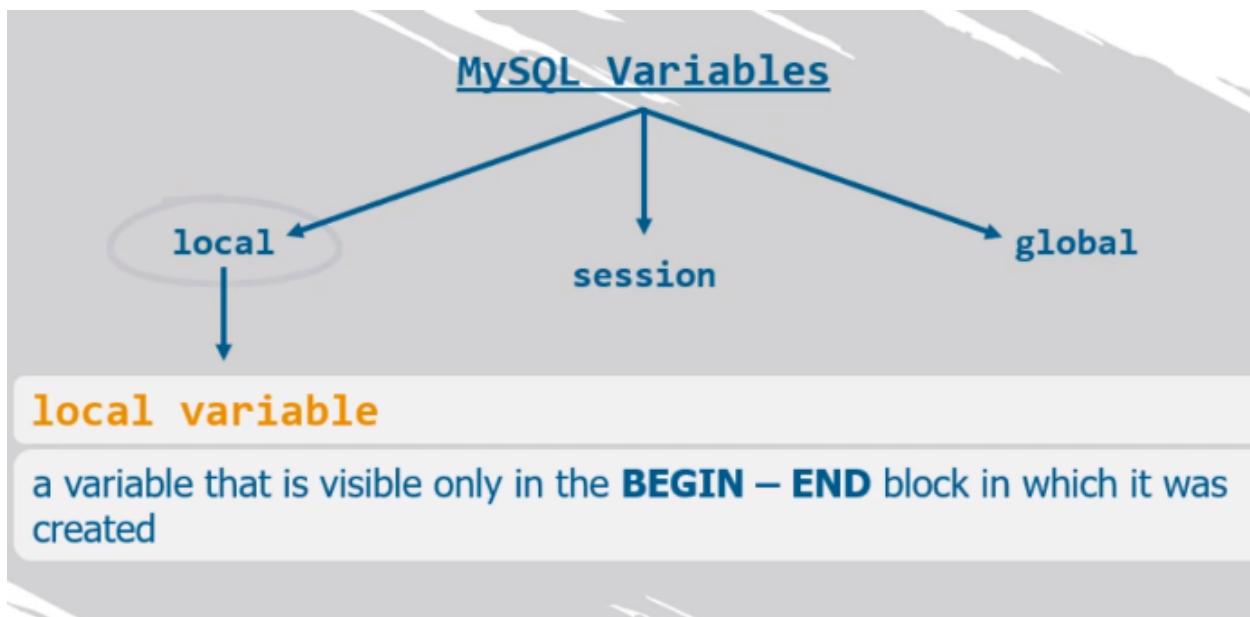
```
SET @v_emp_no = 11300;
SELECT
 emp_no,
 first_name,
 last_name,
 f_emp_avg_salary(@v_emp_no) AS avg_salary
FROM
 employees
WHERE
 emp_no = @v_emp_no;
```

| Result Grid   Filter Rows:     Exp |        |            |           |            |
|------------------------------------|--------|------------|-----------|------------|
|                                    | emp_no | first_name | last_name | avg_salary |
| ▶                                  | 10789  | Aruna      | JourneI   | 41991.20   |

## Section 18. Advanced SQL

## 92. Types of MySQL variables - local variables

- (1) Scope: the region of a computer program where a phenomenon, such as variable, is considered valid. Also known as visibility.
- (2) Three types of variables:



Example: `f_emp_avg_salary` is a local variable

The keyword to define a local variable is **DECLARE**.

When we SELECT this variable OUTSIDE “END\$\$”, there is an error. This means that the scope of this variable is only within BEGIN... END...

| # | Action                                                                    | Message                                                       | Duration / Fetch |
|---|---------------------------------------------------------------------------|---------------------------------------------------------------|------------------|
| 2 | CREATE FUNCTION f_emp_avg_salary (p_emp_no integer) RETURNS decimal(10,2) | 0 rows affected.                                              | 0.000 sec.       |
| 3 | SELECT v_avg_salary                                                       | Error Code: 1054 Unknown column 'v_avg_salary' in field list' |                  |

Similarly, we create a v\_avg\_salary\_2 in an embedded BEGIN...END structure within the original one. There will be an error if we mention this variable outside this embedded structure.

The screenshot shows the MySQL Workbench interface. In the main editor window, a function definition is written in SQL:

```
7 DECLARE v_avg_salary DECIMAL(10,2);
8
9 BEGIN
10 DECLARE v_avg_salary_2 DECIMAL(10,2);
11 END;
12
13 SELECT
14 AVG(s.salary)
15 INTO v_avg_salary_2 FROM
16 employees e
17 JOIN
18 salaries s ON e.emp_no = s.emp_no
19 WHERE
20 e.emp_no = p_emp_no;
21
22 RETURN v_avg_salary_2;
23 END$$
24
25 DELIMITER ;
26
27 • SELECT v_avg_salary;
```

In the 'Output' pane, the results of the function creation are displayed:

| # | Time     | Action                                                                    | Message                                               |
|---|----------|---------------------------------------------------------------------------|-------------------------------------------------------|
| 6 | 12:32:58 | DROP FUNCTION IF EXISTS f_emp_avg_salary                                  | 0 row(s) affected                                     |
| 7 | 12:33:02 | CREATE FUNCTION f_emp_avg_salary (p_emp_no integer) RETURNS decimal(10,2) | Error Code: 1327. Undeclared variable: v_avg_salary_2 |

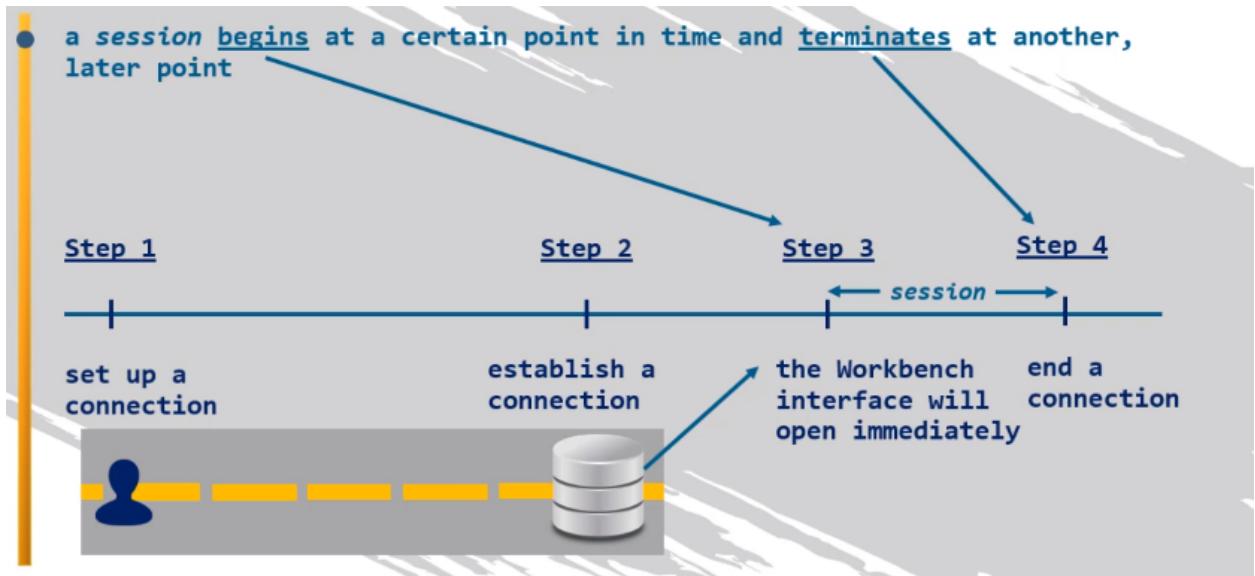
A red arrow points from the text 'Undeclared variable: v\_avg\_salary\_2' in the error message to the variable declaration 'v\_avg\_salary\_2' in the code.

### 93. Session variables

Session: a series of information exchange interactions, or a dialogue, between a computer and a user → e.g., a dialogue between the MySQL server and a client application like MySQL Workbench

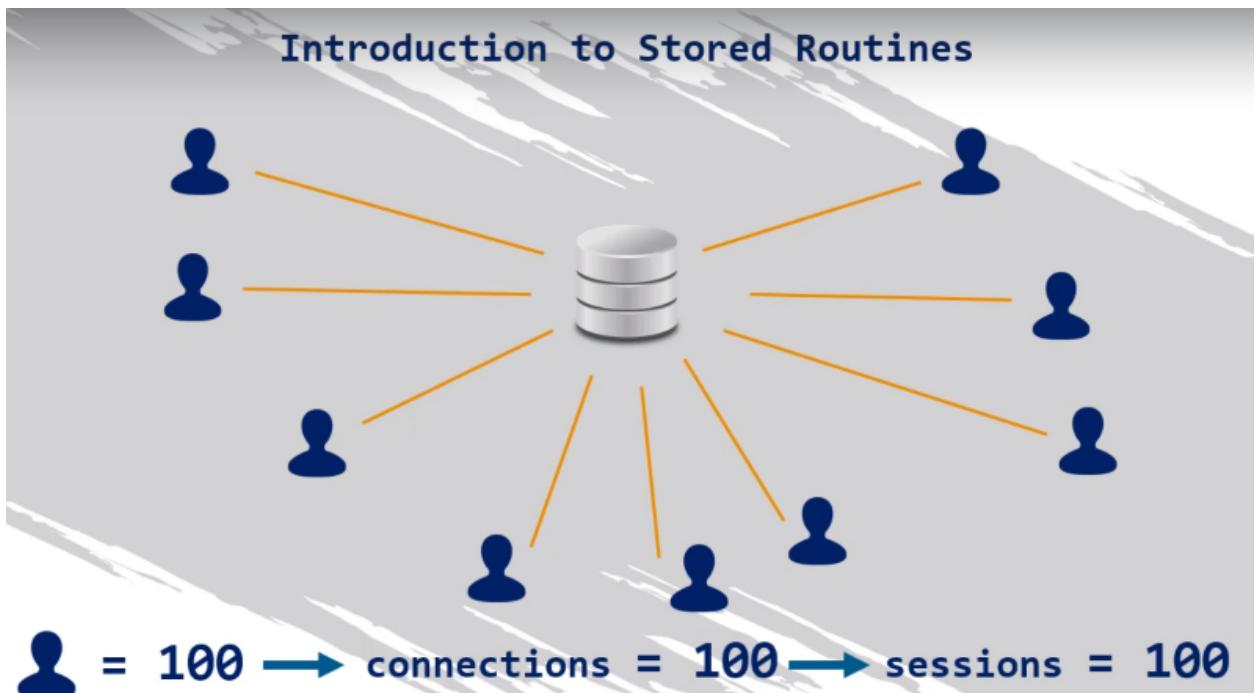
A session begins at a certain point in time and terminates at another later point

For example, a user logs into a SQL server:



There are certain SQL objects that are valid for a specific session only

**Session variable:** a variable that exists only for the session in which you are operating



Session variable must use “SET” and “@”

exercise\* x SQL File 3\*

```
1 #CREATE DATABASE IF NOT E
2
3 #DROP TABLE sales;
4 #DROP TABLE customers;
5 #DROP TABLE items;
6
7 #call emp_avg_salary_out(
8
9 • SET @s_var1 = 3;
10
11 • SELECT @s_var1
12
```

This variable can be called in another SQL tab in the same connection

The screenshot shows the MySQL Workbench interface. At the top, there are two tabs: 'exercise\*' and 'SQL File 3\*'. Below the tabs is a toolbar with various icons. A status bar at the bottom indicates '1 • SELECT @s\_var1'. The main area contains a result grid with one row. The grid has two columns: the first column is empty, and the second column is labeled '@s\_var1' with the value '3'.

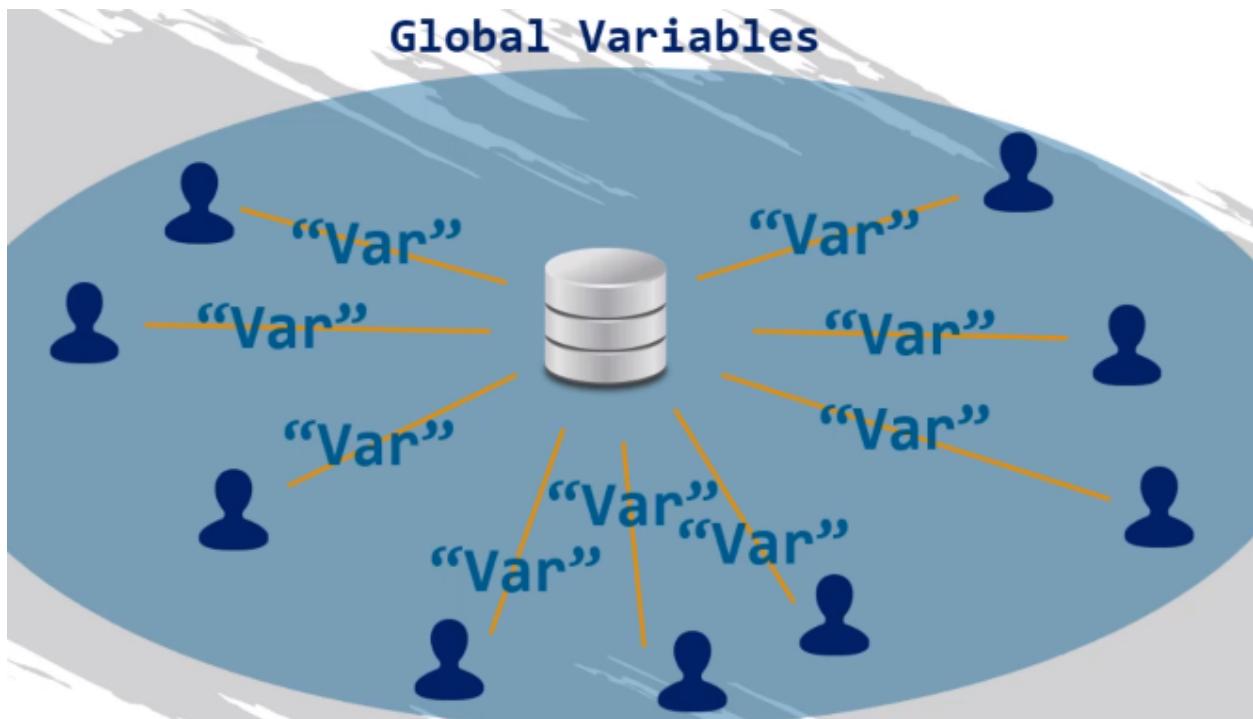
|   | @s_var1 |
|---|---------|
| ▶ | 3       |

Session can be defined here:

Local instance MySQL80

#### 94. Global variables

Global variable applies to all connections related to a specific server



Syntax:

```

</> SQL SET GLOBAL var_name = value;
</> SQL SET @@global.var_name = value;

```

You cannot set just any variable as global:

- A specific group of predefined variables in MySQL is suitable for this job. They are called system variables

`.max_connections()`

- Indicates the maximum number of connections to a server that can be established at a certain point in time

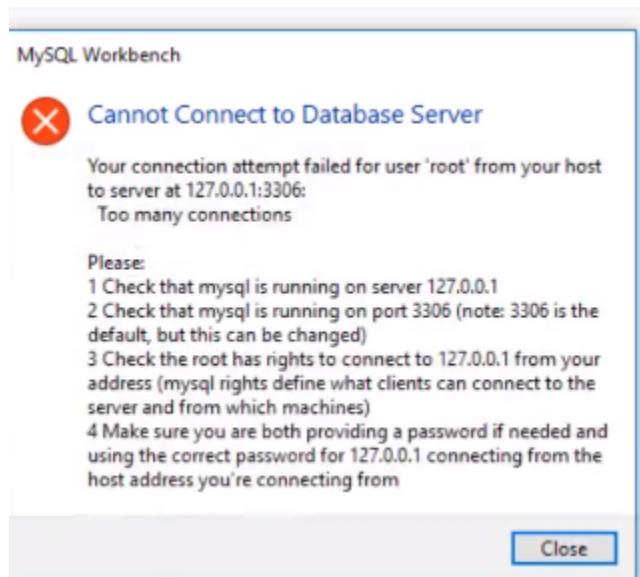
`.max_join_size()`

- Sets the maximum memory space allocated for the joins created by a certain connection

Example:

We can reset the connection# from 1000 to 1. When we log into another server at the home page, there is a warning:

```
SET GLOBAL max_connections = 1000;
SET @@global.max_connections =1;
```



Question 1:

**A person with what profession do you think will be in control of setting global variables in a database?**

a data scientist

a data analyst

a database administrator

any SQL specialist

## 95. User-defined vs System variables

- variables in MySQL can be characterized according to the way they have been created

**user-defined**

variables that can be *set by the user manually*

**system**

variables that are *pre-defined on our system - the MySQL server*

- variables in MySQL can be characterized according to the way they have been created

|              | local | session | global |
|--------------|-------|---------|--------|
| user-defined | ✓     | ✓*      | ✗      |
| system       | ✗     | ✓*      | ✓      |

- both user-defined and system variables can be set as session variables there are limitations to this rule!

### Session variables:

- (1) Some of the system variables can be defined as global only, e.g., .max\_connections()
- (2) Some can be defined as session variable

```
6
7 • SET SESSION sql_mode='STRICT_TRANS_TABLES,NO_ZERO_DATE,NO_AUTO_CREATE_USER,NO_ENGINE_SUBSTITUTION';
8
```

T

A user can define a local or a session variable.

Some system variable can be set as session variables or global variables

```
1 • SET SESSION max_connections = 1000;
2
3 • SET @@global.max_connections = ;
4
5
6
7 • SET GLOBAL sql_mode='STRICT_TRANS_TABLES,NO_ZERO_DATE,NO_AUTO_CREATE_USER,NO_ENGINE_SUBSTITUTION';
8
```

can be set as **global only**

could be set either as a **session or a global** variable

Question 1:

Which of the following variables is not a system variable in MySQL?

- a variable setting the date format
- a variable setting the maximum size of the joins used on a server
- a variable defining the maximum size of the files exchanged on a certain server
- a variable setting the minimum salary paid to employees registered in the database

## 96. MySQL Triggers

A trigger is a MySQL object that can “trigger” a specific action or calculation ‘before’ or ‘after’ an **INSERT, UPDATE, or DELETE** statement has been executed. For instance, a trigger can be activated before a new record is inserted into a table, or after a record has been updated.

(1) Before-insert trigger

DELIMITER \$\$

```
CREATE TRIGGER before_salaries_insert
BEFORE INSERT ON salaries
FOR EACH ROW
BEGIN
 IF NEW.salary < 0 THEN
 SET NEW.salary = 0;
 END IF;
END $$

DELIMITER ;
```

In general, NEW refers to a row that has just been inserted or updated. In our case, after we insert a new record, “NEW dot salary” will refer to the value that will be inserted in the “Salary” column of the “Salaries” table.

After inserting a negative salary row into the table, the table now fixes the value as 0. **The “before\_salaries\_insert” trigger was activated automatically.** It corrected the value of minus 92,891 we tried to insert.

```

26 • INSERT INTO salaries VALUES ('10001', -92891, '2010-06-22', '9999-01-01');
27
28 • SELECT
29 *
30 FROM
31 salaries
32 WHERE
33 emp_no = '10001';
34
35

```

|  | emp_no | salary | from_date  | to_date    |
|--|--------|--------|------------|------------|
|  | 10001  | 66596  | 1989-06-25 | 1990-06-25 |
|  | 10001  | 66961  | 1990-06-25 | 1991-06-25 |
|  | 10001  | 71046  | 1991-06-25 | 1992-06-24 |
|  | 10001  | 74333  | 1992-06-24 | 1993-06-24 |
|  | 10001  | 75286  | 1993-06-24 | 1994-06-24 |
|  | 10001  | 75994  | 1994-06-24 | 1995-06-24 |
|  | 10001  | 76884  | 1995-06-24 | 1996-06-23 |
|  | 10001  | 80013  | 1996-06-23 | 1997-06-23 |
|  | 10001  | 81025  | 1997-06-23 | 1998-06-23 |
|  | 10001  | 81097  | 1998-06-23 | 1999-06-23 |
|  | 10001  | 84917  | 1999-06-23 | 2000-06-22 |
|  | 10001  | 85112  | 2000-06-22 | 2001-06-22 |
|  | 10001  | 85097  | 2001-06-22 | 2002-06-22 |
|  | 10001  | 88958  | 2002-06-22 | 9999-01-01 |
|  | 10001  | 0      | 2010-06-22 | 9999-01-01 |
|  | NULL   | NULL   | NULL       | NULL       |

## (2) Before-update trigger

DELIMITER \$\$

```

CREATE TRIGGER trig_upd_salary
BEFORE UPDATE ON salaries
FOR EACH ROW
BEGIN
 IF NEW.salary < 0 THEN
 SET NEW.salary = OLD.salary;
 END IF;
END $$
```

DELIMITER ;

We add a negative salary value to see whether this works:

### UPDATE salaries

```
SET
 salary = - 50000
WHERE
 emp_no = '10001'
 AND from_date = '2010-06-22';
```

```
30 • SELECT
31 *
32 FROM
33 salaries
34 WHERE
35 emp_no = '10001'
36 AND from_date = '2010-06-22';
37
38
```

| emp_no | salary | from_date  | to_date    |
|--------|--------|------------|------------|
| 10001  | 0      | 2010-06-22 | 9999-01-01 |

System functions: System functions can also be called built-in functions. Often applied in practice, they provide data related to the moment of the execution of a certain query.

```
SELECT SYSDATE();
SELECT DATE_FORMAT(SYSDATE(), '%y-%m-%d') as today;
```

```
DELIMITER $$
```

```
CREATE TRIGGER trig_ins_dept_mng
AFTER INSERT ON dept_manager
FOR EACH ROW
BEGIN
 DECLARE v_curr_salary int;
```

```

SELECT
 MAX(salary)
INTO v_curr_salary FROM
 salaries
WHERE
 emp_no = NEW.emp_no;

IF v_curr_salary IS NOT NULL THEN
 UPDATE salaries
 SET
 to_date = SYSDATE()
WHERE
 emp_no = NEW.emp_no and to_date = NEW.to_date;

 INSERT INTO salaries
 VALUES (NEW.emp_no, v_curr_salary + 20000, NEW.from_date,
NEW.to_date);
END IF;
END $$

DELIMITER ;

INSERT INTO dept_manager VALUES ('111534', 'd009', date_format(sysdate(), "%y-%m-%d"),
'9999-01-01');

SELECT
*
FROM
 dept_manager
WHERE
 emp_no = 111534;

```

Conceptually, this was an ‘after’ trigger that automatically added \$20,000 to the salary of the employee who was just promoted as a manager. Moreover, it set the start date of her new contract to be the day on which you executed the insert statement.

#### **Exercise:**

**Create a trigger that checks if the hire date of an employee is higher than the current date. If true, set this date to be the current date. Format the output appropriately (YY-MM-DD).**

```

USE employees;
COMMIT;

DELIMITER $$

CREATE TRIGGER salary_update
BEFORE UPDATE ON employees_dup
FOR EACH ROW
BEGIN
 IF NEW.hire_date > date_format(sysdate(), '%y-%m-%d') THEN #formating the date
value into YY-MM-DD
 SET NEW.hire_date = date_format(sysdate(), '%y-%m-%d');
 END IF;
END $$

DELIMITER ;

```

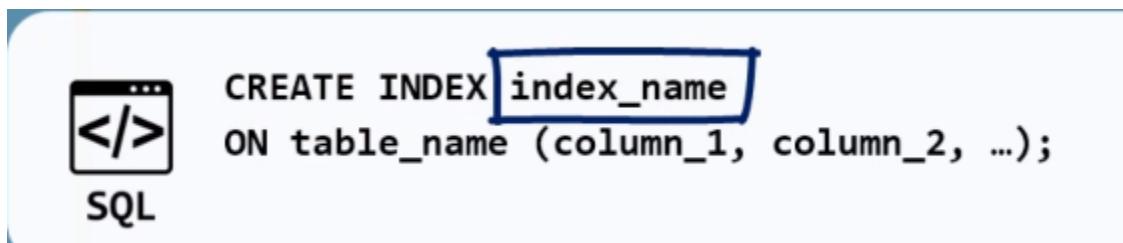
## 97. MySQL Indexes

The index of a table functions like the index of a book

Data is taken from a column of the table and is sorted in a certain order in a distance place, called an index.

Your dataset will typically contains 100,000

(1) Syntax:



Example: after add the index creation line, the processing speed decreases a lot

SELECT

\*

```

FROM
 employees
WHERE
 hire_date > "2000-01-01";

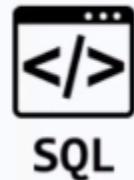
CREATE INDEX i_hire_date ON employees(hire_date);

SELECT
 *
FROM
 employees
WHERE
 hire_date > "2000-01-01";

```

|    |          |                                                                  |                                                        |                       |
|----|----------|------------------------------------------------------------------|--------------------------------------------------------|-----------------------|
| 24 | 11:18:42 | SELECT * FROM employees WHERE hire_date > "2000-01-01" LIMIT ... | 14 row(s) returned                                     | 0.140 sec / 0.000 sec |
| 25 | 11:18:42 | CREATE INDEX i_hire_date ON employees(hire_date)                 | 0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0 | 1.282 sec             |
| 26 | 11:19:01 | SELECT * FROM employees WHERE hire_date > "2000-01-01" LIMIT ... | 14 row(s) returned                                     | 0.000 sec / 0.000 sec |

## (2) Composite index: applying on multiple columns



```

CREATE INDEX index_name
ON table_name (column 1, column 2, ...);

```

```

SELECT
 *
FROM
 employees
WHERE
 first_name = "Georgi" AND last_name = "Facello";

CREATE INDEX i_composite ON employees(first_name,last_name);

SELECT
 *
FROM
 employees
WHERE
 first_name = "Georgi" AND last_name = "Facello";

```

```

✓ 27 11:31:23 SELECT * FROM employees WHERE first_name = "Georgi" AND last_n... 2 row(s) returned 0.141 sec / 0.000 sec
✓ 28 11:32:00 CREATE INDEX i_composite ON employees(first_name,last_name) 0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0 1.859 sec
✓ 29 11:32:08 SELECT * FROM employees WHERE first_name = "Georgi" AND last_n... 2 row(s) returned 0.000 sec / 0.000 sec

```

### (3) Ways of view indexes

#### A - in the schemas

**Indexes in Table**

| Visible                             | Key         | Type  | Uni... | Columns               |
|-------------------------------------|-------------|-------|--------|-----------------------|
| <input checked="" type="checkbox"/> | PRIMARY     | BTREE | YES    | emp_no                |
| <input checked="" type="checkbox"/> | i_hire_date | BTREE | NO     | hire_date             |
| <input checked="" type="checkbox"/> | i_composite | BTREE | NO     | first_name, last_name |

**Index Details**

Key Name: **i\_hire\_date**  
 Index Type: **BTREE** Packed:  
 Allows NULL: Unique: **NO**  
 Cardinality: **5578**  
 Comment:  
 User Comment:

**Columns in table**

| Column     | Type          | Nullable | Indexes     |
|------------|---------------|----------|-------------|
| emp_no     | int           | NO       | PRIMARY     |
| birth_date | date          | NO       |             |
| first_name | varchar(14)   | NO       | i_composite |
| last_name  | varchar(16)   | NO       | i_composite |
| gender     | enum('M','F') | NO       |             |
| hire_date  | date          | NO       | i_hire_date |

#### B - SHOW INDEX FROM database FROM table

```
SHOW INDEX FROM employees FROM employees;
```

| Table     | Non_unique | Key_name    | Seq_in_index | Column_name | Collation | Cardinality | Sub_part | Packed |
|-----------|------------|-------------|--------------|-------------|-----------|-------------|----------|--------|
| employees | 0          | PRIMARY     | 1            | emp_no      | A         | 299645      | NULL     | NULL   |
| employees | 1          | i_hire_date | 1            | hire_date   | A         | 5578        | NULL     | NULL   |
| employees | 1          | i_composite | 1            | first_name  | A         | 1265        | NULL     | NULL   |
| employees | 1          | i_composite | 2            | last_name   | A         | 279496      | NULL     | NULL   |

### (4) Using INDEX

|                       |                                                                         |
|-----------------------|-------------------------------------------------------------------------|
| <u>small datasets</u> | the costs of having an index might be higher than the benefits          |
| <u>large datasets</u> | a well-optimized index can make a positive impact on the search process |

Exercise1: drop i\_hire\_date

```
ALTER TABLE employees
DROP INDEX i_hire_date;
```

#### Exercise2:

Select all records from the ‘salaries’ table of people whose salary is higher than \$89,000 per annum.

Then, create an index on the ‘salary’ column of that table, and check if it has sped up the search of the same SELECT statement.

```
SELECT
*
FROM
 salaries
WHERE salary > 89000;

CREATE INDEX high_salary ON salaries(salary);
```

|             |                                                      |                                                        |                       |
|-------------|------------------------------------------------------|--------------------------------------------------------|-----------------------|
| 33 11:40:40 | SELECT * FROM salaries WHERE salary > 89000 LIMIT... | 1000 row(s) returned                                   | 0.000 sec / 0.000 sec |
| 34 11:41:23 | CREATE INDEX high_salary ON salaries(salary)         | 0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0 | 2.844 sec             |
| 35 11:41:33 | SELECT * FROM salaries WHERE salary > 89000 LIMIT... | 1000 row(s) returned                                   | 0.047 sec / 0.000 sec |

#### 98. CASE statement

CASE is a type of condition clause:

```

CASE
 WHEN... THEN...
 ELSE...
END AS...

```

**Example:re-represent the genders in full words**

```

SELECT
 emp_no,
 first_name,
 last_name,
 CASE
 WHEN gender = "M" THEN "Male"
 ELSE "Female"
 END AS gender
FROM
employees;

```

| emp_no | first_name | last_name | gender |
|--------|------------|-----------|--------|
| 10001  | Georgi     | Facello   | Male   |
| 10002  | Bezalel    | Simmel    | Female |
| 10003  | Parto      | Bamford   | Male   |
| 10004  | Chirstian  | Koblick   | Male   |
| 10005  | Kyoichi    | Maliniak  | Male   |
| 10006  | Anneke     | Preusig   | Female |
| 10007  | Tatyana    | Deckow    | Female |

**Another way is to put gender after the CASE and not use =**

```

SELECT
 emp_no,
 first_name,
 last_name,
 CASE gender
 WHEN "M" THEN "Male"
 ELSE "Female"
 END AS gender
FROM
employees;

```

**A third way is to use a reduced IF structure:**

```

SELECT
 emp_no,

```

```

first_name,
last_name,
IF(gender = "M","Male", "Female") AS gender
FROM
 employees;

```

### The difference between CASE and IF:

**IF VS CASE**

you can have just one conditional expression

we can have multiple conditional expressions

```

SELECT
 dm.emp_no,
 e.first_name,
 e.last_name,
 MAX(s.salary) - MIN(s.salary) AS salary_difference,
 CASE
 WHEN MAX(s.salary) - MIN(s.salary) > 30000 THEN 'Salary was raised by more than $30,000'
 WHEN MAX(s.salary) - MIN(s.salary) BETWEEN 20000 AND 30000 THEN
 'Salary was raised by more than $20,000 but less than $30,000'
 ELSE 'Salary was raised by less than $20,000'
 END AS salary_increase
FROM
 dept_manager dm
 JOIN
 employees e ON e.emp_no = dm.emp_no
 JOIN
 salaries s ON s.emp_no = dm.emp_no
GROUP BY s.emp_no;

```

```

SELECT
 dm.emp_no,
 e.first_name,
 e.last_name,
 MAX(s.salary) - MIN(s.salary) AS salary_difference,
 CASE
 WHEN MAX(s.salary) - MIN(s.salary) > 30000 THEN 'Salary was raised by more than $30,000'
 WHEN MAX(s.salary) - MIN(s.salary) BETWEEN 20000 AND 30000 THEN
 'Salary was raised by more than $20,000 but less than $30,000'
 ELSE 'Salary was raised by less than $20,000'
 END AS salary_increase
FROM
 dept_manager dm
 JOIN
 employees e ON e.emp_no = dm.emp_no
 JOIN
 salaries s ON s.emp_no = dm.emp_no
GROUP BY s.emp_no;

```

There can be two WHENs for CASE

### Exercise 1:

Similar to the exercises done in the lecture, obtain a result set containing the employee number, first name, and last name of all employees with a number higher than 109990. Create a fourth column in the query, indicating whether this employee is also a manager, according to the data provided in the dept\_manager table, or a regular employee.

The key of this problem is to find out what two tables to overlap. Using dept\_manager can make the boolean easier by saying any employees not shown in this table is an employee.

```

SELECT
 e.emp_no,
 e.first_name,
 e.last_name,
 CASE
 WHEN dm.emp_no IS NOT NULL THEN "Yes"
 ELSE "No"
 END AS if_manager
FROM employees e
 LEFT JOIN
 dept_manager dm ON e.emp_no = dm.emp_no
WHERE e.emp_no > 109990;

```

### **Exercise2:**

**Extract a dataset containing the following information about the managers: employee number, first name, and last name. Add two columns at the end – one showing the difference between the maximum and minimum salary of that employee, and another one saying whether this salary raise was higher than \$30,000 or NOT.**

**If possible, provide more than one solution.**

(1) Use CASE

```

SELECT
 e.emp_no,
 e.first_name,
 e.last_name,
 MAX(s.salary) - MIN(s.salary) AS salary_difference
 #CASE
 # WHEN MAX(s.salary)-MIN(s.salary) > 30000 THEN "Yes"
 # ELSE "No"
 #END AS higher_than_30k
FROM
 salaries s
JOIN
 employees e ON s.emp_no = e.emp_no
JOIN
 dept_manager dm ON e.emp_no = dm.emp_no
GROUP BY s.emp_no #must include, otherwise only the first line
;
```

(2) Use IF

SELECT

```
dm.emp_no,
e.first_name,
e.last_name,
MAX(s.salary) - MIN(s.salary) AS salary_difference,
```

```
IF(MAX(s.salary) - MIN(s.salary) > 30000, 'Salary was raised by more than $30,000',
'Salary was NOT raised by more than $30,000') AS salary_increase
```

FROM

```
dept_manager dm
```

JOIN

```
employees e ON e.emp_no = dm.emp_no
```

JOIN

```
salaries s ON s.emp_no = dm.emp_no
```

```
GROUP BY s.emp_no;
```

**Exercise3:** Extract the employee number, first name, and last name of the first 100 employees, and add a fourth column, called “current\_employee” saying “Is still employed” if the employee is still working in the company, or “Not an employee anymore” if they aren’t.

**Hint:** You’ll need to use data from both the ‘employees’ and the ‘dept\_emp’ table to solve this exercise.

Still employed means the to\_date is bigger than the current time, sysdate()

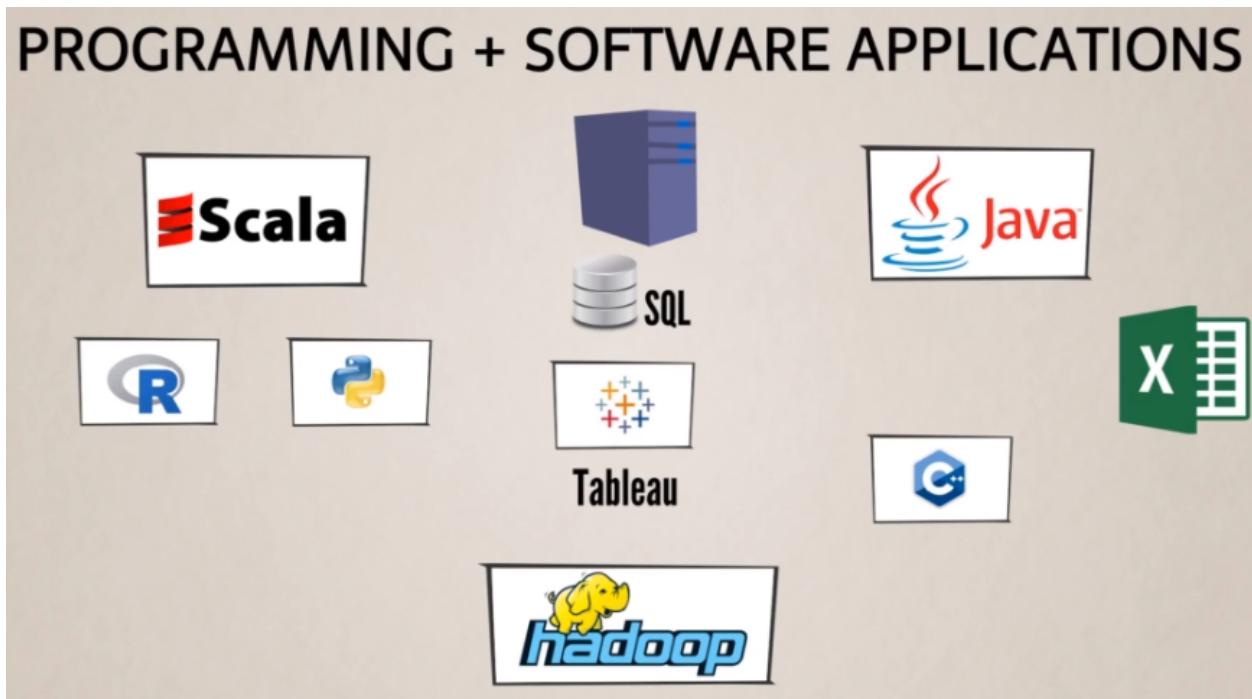
```

SELECT
 e.emp_no,
 e.first_name,
 e.last_name,
 IF(MAX(de.to_date) > sysdate(), "Is still employed","Not an employee any more") AS
current_employee
FROM
 employees e
JOIN
 dept_emp de ON e.emp_no = de.emp_no
GROUP BY de.emp_no
LIMIT 100
;

```

## 98. Combining SQL and Tableau

- Advantage of using them together
- A company usually uses different tools for their analytics



Reasons?

- ✓ optimize the capacity of the chosen software
- ✓ practical reasons
- ✓ cost reduction
- ✓ data security
- ✓ historical reasons
- ✓ a tool which deals with all domains together does not exist

- Kinds of integrations
  - A. You can use customized SQL in Tableau

## STORED PROCEDURES

**SQL**  
better performance

**Tableau**  
good performance

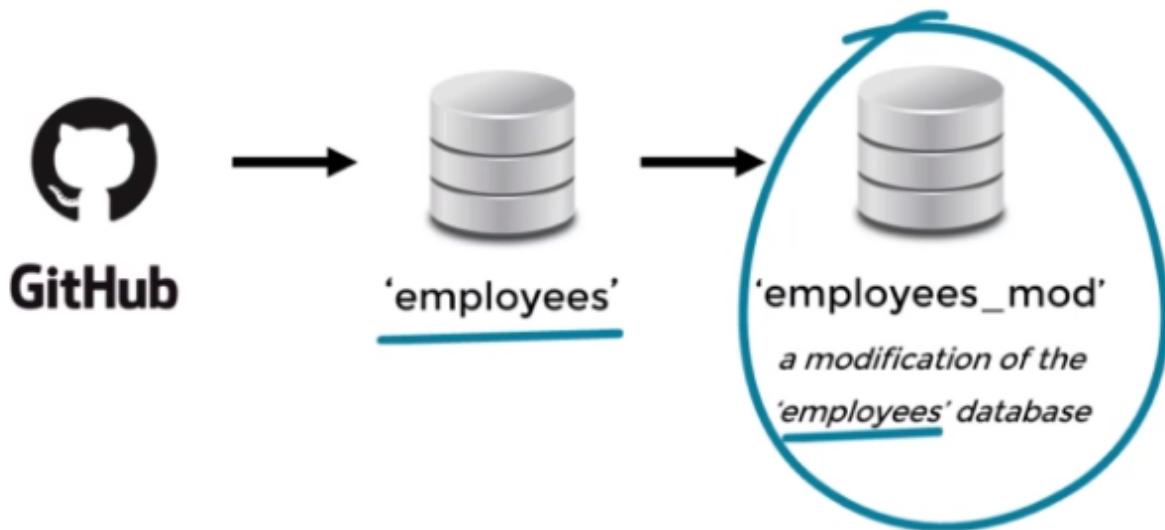
the Tableau in-memory engine is also suitable for similar types of calculations when only using a smaller dataset

- B. SQL is better than tableau to pre-process the data  
For example, transpose, convert data types

Problem structure:

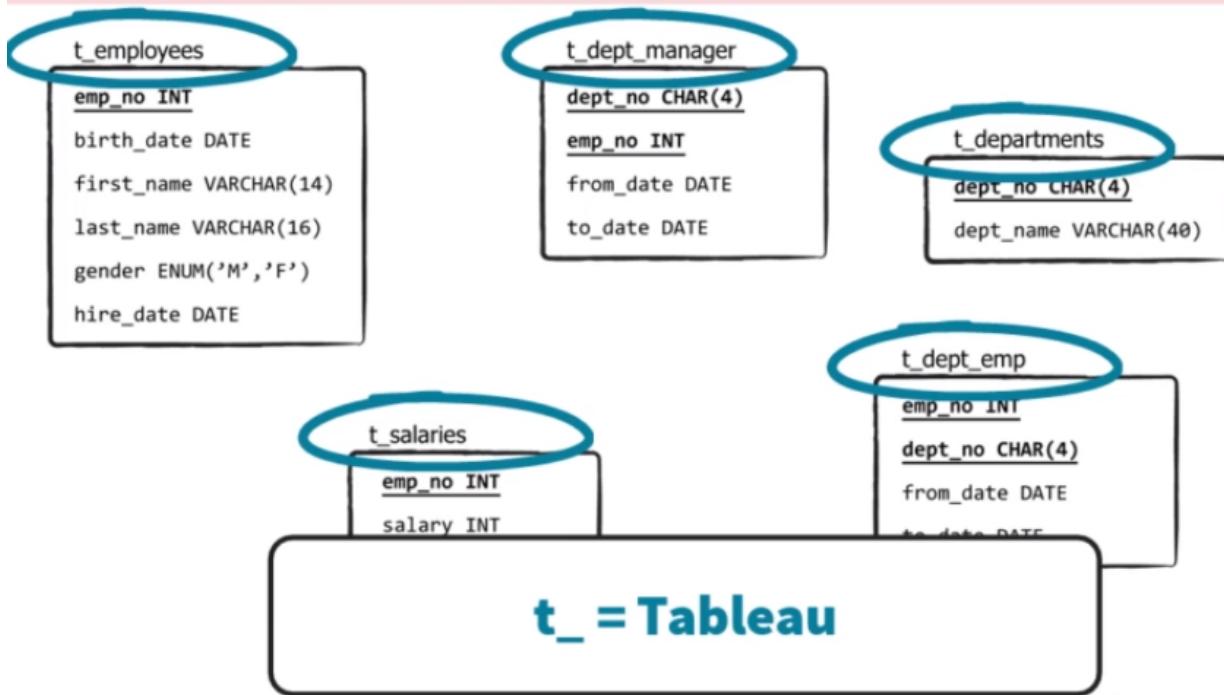
- Receive a business task
- Use SQL to execute a query retrieving a relevant dataset from the database
- Export the newly obtained data in a csv for tableau
- Visualization in tableau

Database description:



Structure:

## Database: employees\_mod

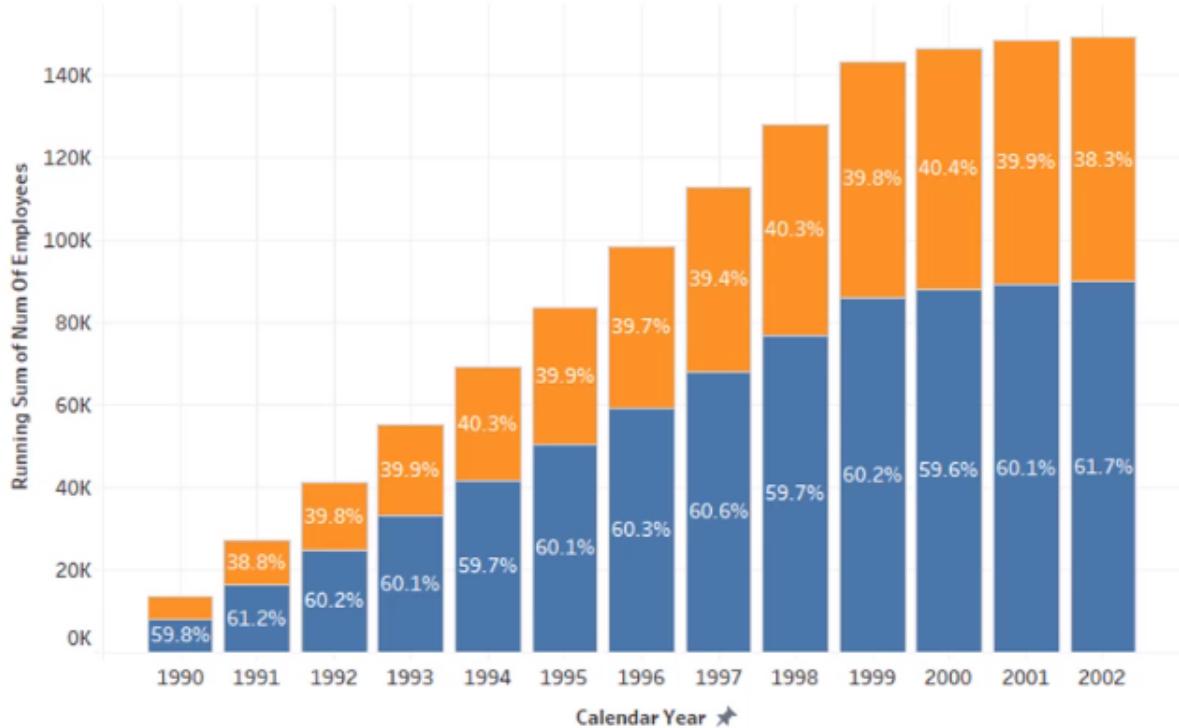


### 99. Task 1

- SQL slice

Create a visualization that provides a breakdown between the male and female employees working in the company each year, starting from 1990

Chart 1



The data needs to be retrieved from the database slice like this

|  | calendar_year | gender | num_of_employees |
|--|---------------|--------|------------------|
|  | 1990          | M      | 8128             |
|  | 1990          | F      | 5469             |
|  | 1991          | M      | 8286             |
|  | 1991          | F      | 5254             |
|  | 1992          | M      | 8472             |
|  | 1992          | F      | 5592             |
|  | 1993          | M      | 8473             |
|  | 1993          | F      | 5614             |
|  | 1994          | M      | 8457             |
|  | 1994          | F      | 5713             |
|  | 1995          | M      | 8613             |
|  | 1995          | F      | 5727             |
|  | 1996          | M      | 8811             |
|  | 1996          | F      | 5812             |
|  | 1997          | M      | 8924             |
|  | 1997          | F      | 5792             |
|  | 1998          | M      | 8925             |
|  | 1998          | F      | 6028             |
|  | 1999          | M      | 9195             |
|  | 1999          | F      | 6072             |
|  | 2000          | M      | 1894             |
|  | 2000          | F      | 1284             |
|  | 2001          | M      | 1153             |
|  | 2001          | F      | 765              |
|  | 2002          | M      | 668              |
|  | 2002          | F      | 414              |

```

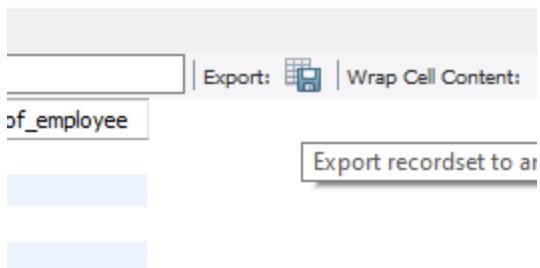
SELECT
 YEAR(d.from_date) AS Calendar_year,
 e.gender,
 COUNT(e.emp_no) as num_of_employee
FROM
 t_employees e
JOIN
 t_dept_emp d ON d.emp_no = e.emp_no
GROUP BY calendar_year, e.gender
HAVING calendar_year >= 1990;

```

→ remember to use having to modify the conditions; also it's about the start date

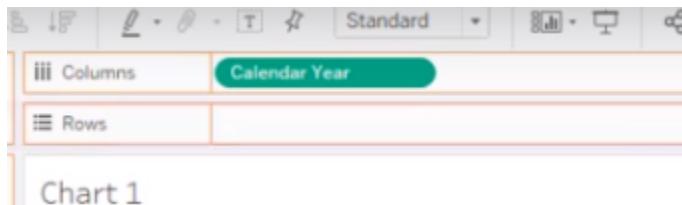
- Transfer the data to Tableau

Click on save here:



Column: X axis

Row: Y axis



## 100. Task 2

Question: compare the number of male managers with female managers from different departments for each year, starting from 1990.

Steps: think about the charts and then SQL

Need to use area chart this time.

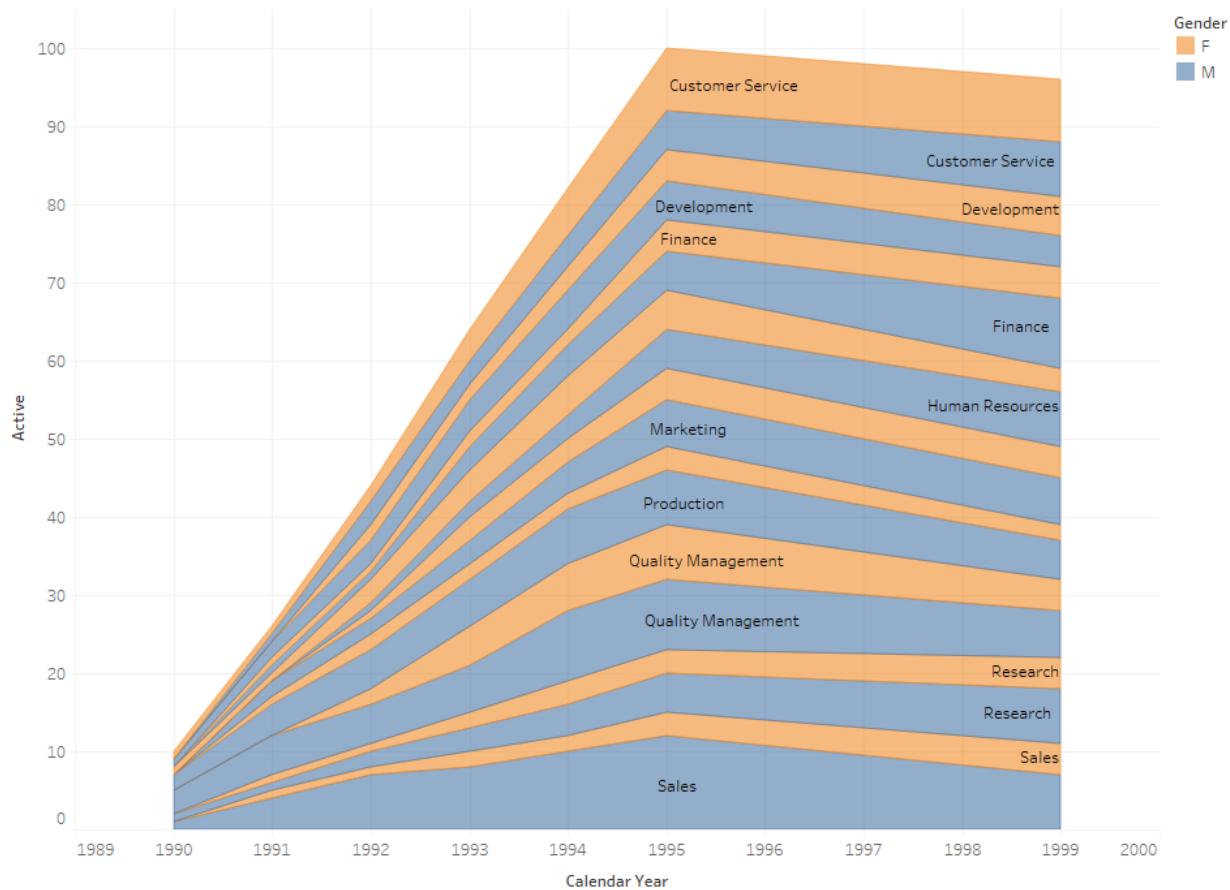
| dept_name | gender | emp_no | from_date  | to_date    | calendar_year | active |
|-----------|--------|--------|------------|------------|---------------|--------|
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1990          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1991          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1992          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1993          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1994          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1995          | 1      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1996          | 1      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1997          | 1      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1998          | 1      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 1999          | 0      |
| Marketino | M      | 110022 | 1995-12-30 | 1998-12-29 | 2000          | 0      |
| Marketino | M      | 110039 | 1997-04-09 | 9999-01-01 | 1990          | 0      |
| Marketino | M      | 110039 | 1997-04-09 | 9999-01-01 | 1991          | 0      |

→ during some employee's work, it is possible that they were NOT manager at that time, so we need a conditional active column for this variable

```

SELECT
 d.dept_name,
 ee.gender,
 dm.emp_no,
 dm.from_date,
 dm.to_date,
 e.calendar_year,
 CASE
 WHEN e.calendar_year >= YEAR(dm.from_date)
 AND e.calendar_year <= YEAR(dm.to_date)
 THEN 1
 ELSE 0
 END AS active
FROM
 (SELECT
 YEAR(hire_date) as calendar_year
 FROM
 t_employees
 GROUP BY calendar_year
) e
 CROSS JOIN
 t_dept_manager dm
 JOIN
 t_employees ee ON dm.emp_no = ee.emp_no
 JOIN
 t_departments d
 ON
 dm.dept_no = d.dept_no
ORDER BY dm.emp_no AND e.calendar_year;

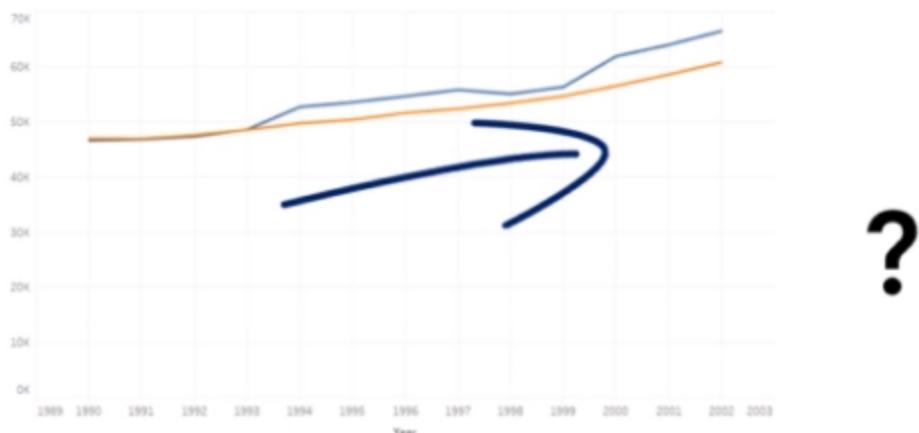
```



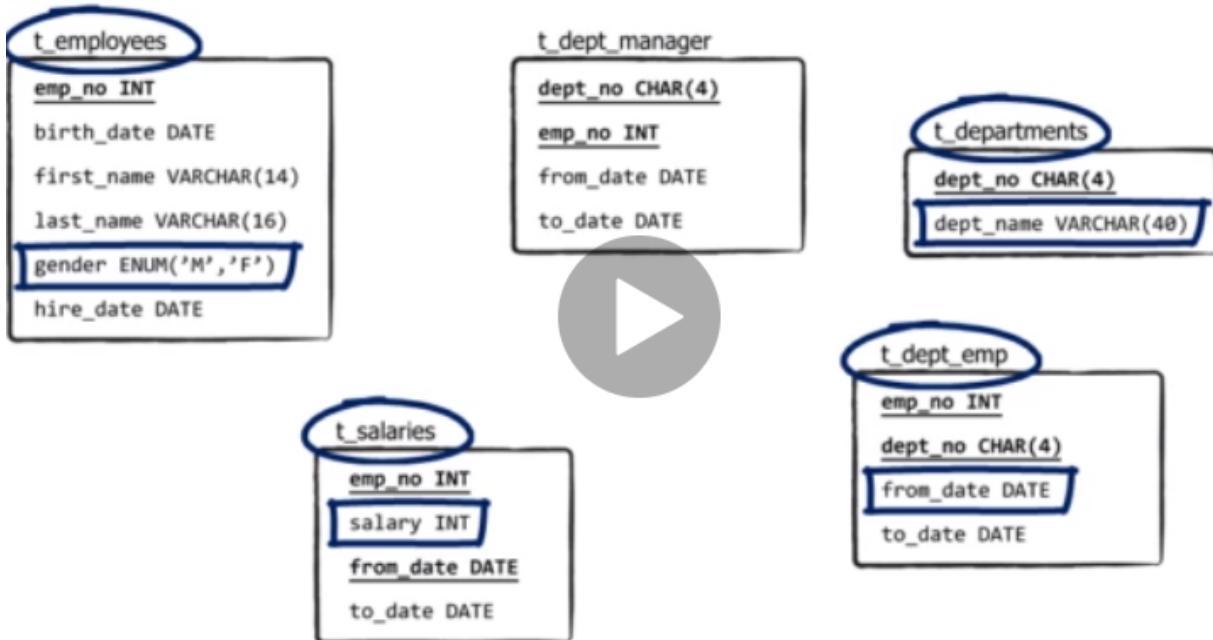
### Task 3:

Compare the average salary of female and male employees in the entire company until 2002, and add a filter allowing you to see that per each department

Using line chart



|   | gender    | dept_name | salary | calendar_year |
|---|-----------|-----------|--------|---------------|
| M | Marketina | 58895.85  | 1990   |               |
| M | Marketina | 59232.75  | 1991   |               |
| M | Marketina | 59743.08  | 1992   |               |
| M | Marketina | 60436.85  | 1993   |               |
| M | Marketina | 64547.55  | 1994   |               |
| M | Marketina | 65377.05  | 1995   |               |
| M | Marketina | 66467.56  | 1996   |               |
| M | Marketina | 67253.18  | 1997   |               |
| M | Marketina | 66332.51  | 1998   |               |
| M | Marketina | 67594.58  | 1999   |               |
| M | Marketina | 73248.01  | 2000   |               |
| M | Marketina | 75364.26  | 2001   |               |
| M | Marketina | 77525.24  | 2002   |               |
| F | Marketina | 57358.31  | 1990   |               |
| F | Marketina | 57670.20  | 1991   |               |



```

SELECT
 te.gender,
 ROUND(AVG(ts.salary),1),
 td.dept_name,
 YEAR(tde.from_date) as calendar_year
FROM
 t_employees te
 JOIN
 t_salaries ts ON te.emp_no = ts.emp_no
 JOIN
 t_dept_emp tde ON ts.emp_no = tde.emp_no
 JOIN
 t_departments td ON td.dept_no = tde.dept_no

GROUP BY td.dept_name, te.gender, calendar_year
HAVING calendar_year <= 2002
ORDER BY tde.dept_no;

```

Task 4:

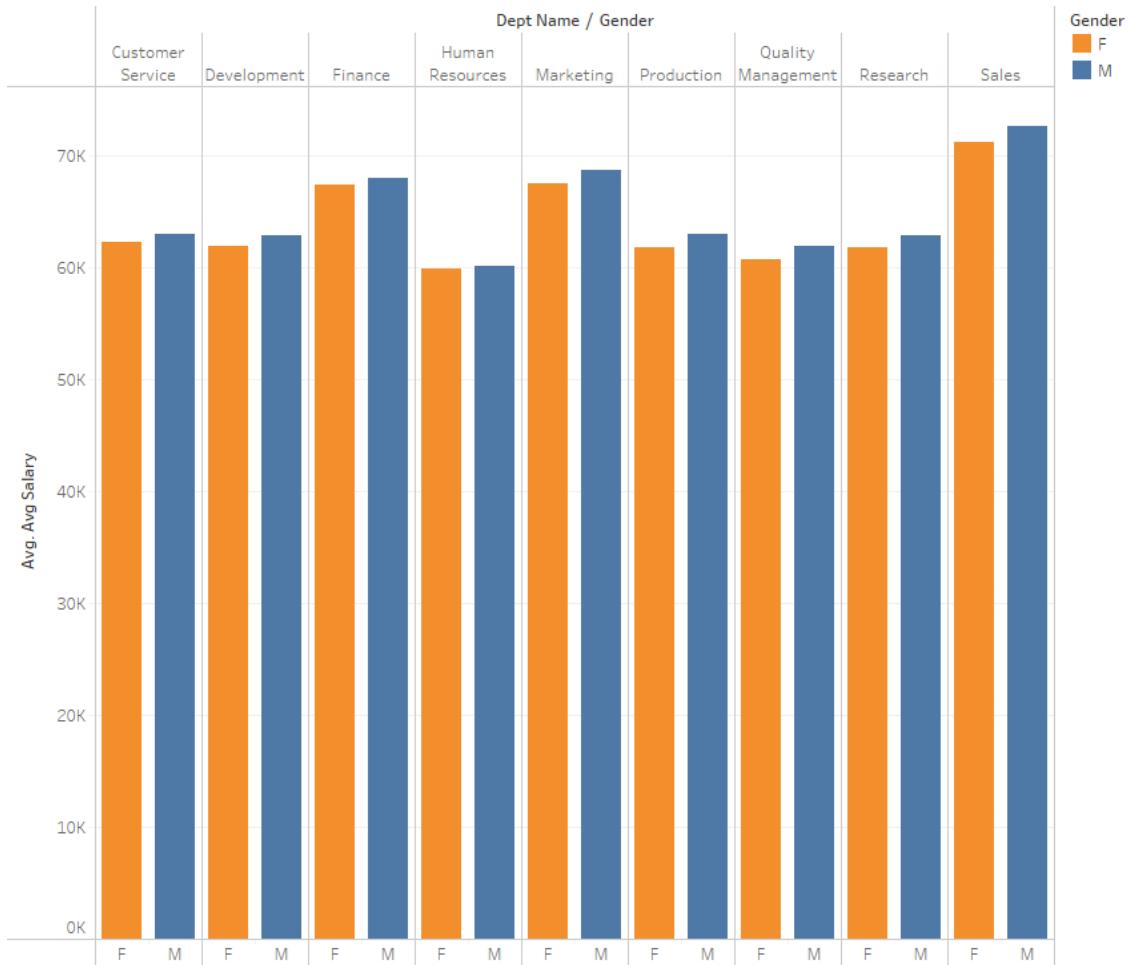
Create an SQL stored procedure that will allow you to obtain the average male and female salary per department within a certain salary range. Let this range be defined by two values the user can insert when calling the procedure (using double-bar charts).

```
DROP PROCEDURE IF EXISTS filter_salary;

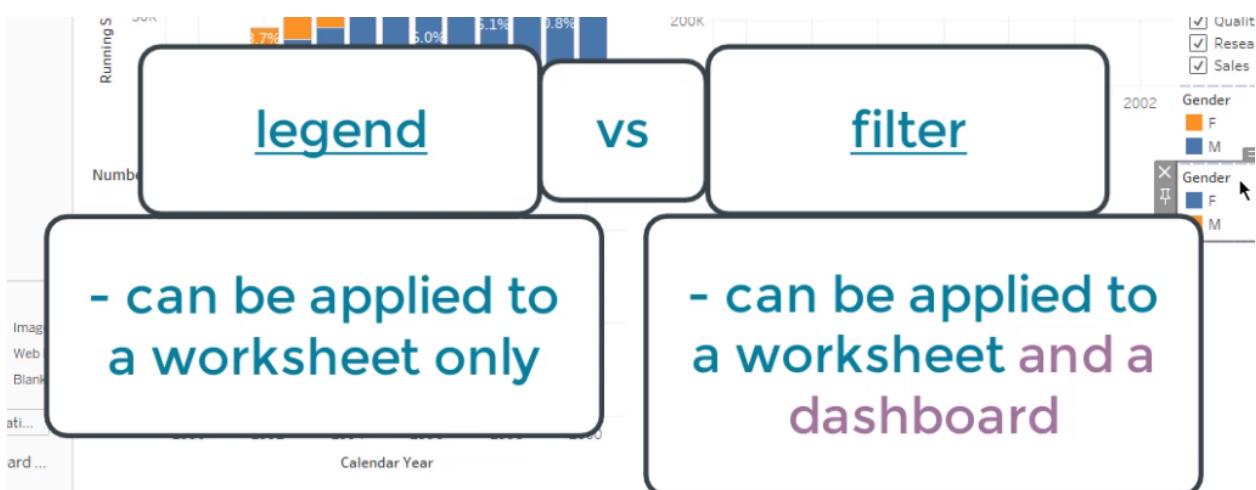
DELIMITER $$
CREATE PROCEDURE filter_salary (IN min FLOAT, IN max FLOAT)
BEGIN
SELECT
 e.gender, d.dept_name, AVG(s.salary) as avg_salary
FROM
 t_salaries s
 JOIN
 t_employees e ON s.emp_no = e.emp_no
 JOIN
 t_dept_emp de ON de.emp_no = e.emp_no
 JOIN
 t_departments d ON d.dept_no = de.dept_no
 WHERE s.salary BETWEEN min AND max
GROUP BY d.dept_no, e.gender;
END$$

DELIMITER ;

CALL filter_salary(50000, 90000);
```



## Total dashboard



## → using Apply to Worksheets

The screenshot shows two Tableau dashboards. The top dashboard features a treemap visualization of departmental hierarchy over time. The bottom dashboard features a bar chart comparing department names across different gender categories.

**Treemap Dashboard:**

- Y-axis: department (Customer Service, Quality Management, Research, Sales).
- X-axis: Calendar Year (1992, 1994, 1996, 1998).
- Legend: Gender (F, M).

**Bar Chart Dashboard:**

- Y-axis: department (Human Resou., Marke.., Produ.., Quality Mana.., Rese..).
- X-axis: Dept Name / Gender (Finance, Human Resou., Market, Production, Quality Management, Research).

A context menu is open on the treemap visualization, specifically on the 'Customer Service' node. The menu options include:

- All Using Related Data Sources
- All Using This Data Source
- Selected Worksheets...
- Only This Worksheet** (selected)

The main menu of the context menu is titled "Apply to Worksheets" and includes the following items:

- Edit Filter...
- Apply to Worksheets** (selected)
- Format Filter and Set Controls...
- Customize
- Show Title
- Edit Title...

Below these are several filter options, each with a radio button or checkbox:

- Single Value (list)
- Single Value (dropdown)
- Single Value (slider)
- Multiple Values (list)**
- Multiple Values (dropdown)
- Multiple Values (custom list)
- Wildcard Match

Further down the menu are:

- Only Relevant Values
- All Values in Database** (selected)
- Include Values**
- Exclude Values
- Add Show/Hide Button

At the bottom of the menu are floating and height-related options:

- Floating
- Fix Height
- Edit Height...

Finally, there are dashboard-level options:

- Select Container: Vertical
- Deselect
- Remove from Dashboard
- Rename Dashboard Item...

